

PHASE 1 DevOps Infrastructure Project

Wednesday, 17 September 2025 6:43 am

Goal: To understand how to build and deploy modern cloud-native infrastructure

Key Components:

- VM setup with Vmware
- Kubernetes (K3s cluster setup)
- CNI plugins (Calico)
- Dockerfile + app manifests (sampapp, DB, Redis, etc)
- GitHub Project structure
- Network Policies + DNS Troubleshooting
- Persistent Volumes + StorageClass
- Cronjob for postgres backups
- Ingress Setup via Helm
- CI/CD: GitHub > Jenkins > ArgoCD (GitOps)
- Ansible for node provisioning
- Istio ServiceMesh
- Prometheus + Grafana + Kiali
- Middleware, MessageBroker: RabbitMQ

Designed and deployed Kubernetes-based DevOps lab from scratch with Calico CNI, Helm-based Ingress, and GitOps workflow via Jenkins and ArgoCD

- **WHAT I DID**
 - o **VM set up (Vmware)**
 - o **SSH set up between my host machine and the 2 VMs**
 - o **Installed Docker on both VMs**
 - o **Installed K3s on both VMs and joined worker node to master node cluster**
 - o **Set up cluster access from local (kube config)**
 - o **Created simple feedback form app**
 - o **Created a Dockerfile for app**
 - o **Built and Pushed image to DockerHub**
 - o **Created pod yaml manifests**
 - o **Pushed yamls to GitHub**
 - o **Deployed YAMLs to cluster**
 - o **Replaced Flannel with Calico**
 - o **Set up some test Netpols**
 - o **Created PVC for Postgres**
 - o **Cronjob for backups for Postgres**
 - o **GitHub Resync (as some edits were made to cluster directly)**
 - o **Ingress**
 - o **CI/CD with Jenkins & ArgoCD**
 - o **Node Provisioning with Ansible**
 - o **Istio ServiceMesh**
 - o **Prometheus + Grafana + Kiali for observability**
 - o **RabbitMQ**

Each section will show the steps I carried out, and if any issues were faced, a short description of the ISSUE is included, with a reference to which Troubleshooting Log page to refer to for more detailed explanations (the issue number within the reference Troubleshooting Log page will be the same as in

the section). But do note that the issues may not be in chronological order - pls follow Reference Log # AND Issue #.

Troubleshooting Log 1 - Miscellaneous Topics Issues

Troubleshooting Log 2 - Postgres

Troubleshooting Log 3 - Ingress

Troubleshooting Log 4 - CI/CD w Jenkins and ArgoCD

1 VM and ssh setup

Sunday, 10 August 2025 8:08 pm

Technology:

Vmware Workstation Pro 17

VM Creation

1. Click on "Create a New Virtual Machine"
2. Select Installer disc image file (iso)
 - Click on Browse and navigate to the directory I saved the ISO Image:
➤ C:\Users\mikia\Downloads\ubuntu-22.04.5-live-server-amd64.iso
3. Name VM (master1 and worker1)
4. Select "Store virtual disk as a single file"
5. Customize Hardware > New CD/DVD (SATA) change to Auto detect
6. Click Finish and let OS boot up

System Bootup

1. Click on first option in GRUB menu (Try or install Ubuntu Server)
2. Keep clicking Done for all the options offered
3. Name the server as your name, nodename, your username and pw
4. Select Install OpenSSH server
5. Wait for configurations to complete then click Reboot Now
6. When screen shows "failed to mount /cdrom", press enter
7. Login to server and shutdown now
8. Click on VM settings and edit CD/DVD (SATA) to Use physical drive (autodetect)
9. Power on VM again

SSH Setup

1. On host machine, generate the ssh id_rsa private and public keys:

```
$ ssh-keygen -t rsa -b 4096 -C "WorkerNode"
```

-t = key type

Rsa is RSA key pair... alt include ed25519, ecdsa etc

-b = bits

For RSA, common values are 2048 or 4096

4096 is stronger encryption but a little slower than 2048

For ed25519, don't need -b because it's fixed strength

-C = comment only

We will use same key for both master1 and worker1 nodes.

2. Copy contents of the id_rsa.pub key and copy over to /home/miki/.ssh/authorized_keys in VM

```
$ printf '%s\n' 'ssh-rsa
AAAAAB3NzaC1yc2EAAAQABAAQCx3A54D3j0mVcDNajgY/U3RI1aGN20vPm8/VVLsStUOfjX
y1Kt0t4z0WV+RLjAmKUM4DH9/VKzOlue/XYTJQS35pUGBC/IcEfCHTTN3INdMcYctps5OSyIEhIN/A5
ramlxTMy25qH5c860lfcwoxnd52ksmly+efmtMaizUKFhiH28T3MTKm+lT3mqo8XvTjjYnX9XRujF5/
FLSm4ChCpbBu8vEUa+tVjQT6MmAkRypdcqd0g3a79xr3wUq6pzLLM/VMsMkvt3c8lw9QV6rmbvEY
n7UNQt4hgF3r+Exz++qkX00Por/S7971QCTYi/EY3Xdy8MVnoXLiacDabWwurtcn8JQ6DGZDd2HBI4
oeLfK/e258oJ6BWVT/YDMdcv7ApRGjLc7d7OT1Se74hnPoEec1lh7eeOsVmVaKzh0o2rTMfBoffCmp
odikhB+bTirWS0v3JwgDill6wXO8ZqNpAkREW/jqGh0tC7lb2qZg55LGU/z2g0vRi7TJuFP9O0AOlhz+q
i6DOPVNz9FTUU40c/qzTvbyvLhKEFMI$qnXOZ9Fullfpp6CxfZYkWL4OKXS1uLE9jrbPxrM6d4oy3549
Vvv3Anq4+IVMQUH6oVAIO9lztVjd3B+QMpRMssAmnEN9CpVE+HXBE2+'
```

```
7JO8P3gwNdNbK6Z5Nfsj/ZfUUqu5qERw== ubuntu-key' > /home/miki/.ssh/authorized_keys
```

3. Ensure key is a one liner -- :

```
$ sed -i 's/\r$//' /home/miki/.ssh/authorized_keys
```

4. Change ownership of .ssh dir and contents :

```
$ chown -R miki:miki /home/miki/.ssh/
```

5. Change permissions

```
$ chmod 600 /home/miki/.ssh/authorized_keys  
$ chmod 700 /home/miki/.ssh/
```

6. Restart ssh service

```
$ systemctl restart ssh
```

7. Verify effective sshd config:

```
$ sudo sshd -T | grep -E "pubkeyauthentication|authorizedkeysfile|passwordauthentication"
```

Expected Output:

```
pubkeyauthentication yes  
passwordauthentication yes  
authorizedkeysfile .ssh/authorized_keys .ssh/authorized_keys2
```

8. Remove # from AuthorizedKeysFile line in sshd_config

```
# sudo grep -i authorized_keys /etc/ssh/sshd_config  
[sudo] password for miki:  
# Expect .ssh/authorized_keys2 to be disregarded by default in future.  
#AuthorizedKeysFile .ssh/authorized_keys .ssh/authorized_keys2
```

9. To access with nodenames instead of IP

Make a `~/.ssh/config` file in local machine (can use Git Bash)

```
# Master node  
Host master1  
    HostName 192.168.81.128  
    User miki  
    IdentityFile ~/.ssh/id_rsa  
  
# Worker node 1  
Host worker1  
    HostName 192.168.81.129  
    User miki  
    IdentityFile ~/.ssh/id_rsa
```

Change permissions of file to 600. Now can ssh in to master1 and worker1

To debug ssh:

```
ssh -vvv miki@192.168.81.128 / master1
```

Had to regenerate ssh key bc git bash/windows was creating some string for my passphrase that I couldn't ssh using.

```
$ mv ~/.ssh/id_rsa ~/.ssh/id_rsa.bak  
  
$ mv ~/.ssh/id_rsa.pub ~/.ssh/id_rsa.pub.bak  
  
$ ssh-keygen -t rsa -b 4096 -C "localhost" -N "" -f ~/.ssh/id_rsa  
Generating public/private rsa key pair.  
Your identification has been saved in /c/Users/mikia/.ssh/id_rsa  
Your public key has been saved in /c/Users/mikia/.ssh/id_rsa.pub  
The key fingerprint is:  
SHA256:gfuPxY3HbvkoLWqfQHhWiQiTr8VHtTCIF5QFAR17+l0 localhost  
The key's randomart image is:  
+---[RSA 4096]----+  
| +**O+.. |  
| oo=+.+ o |  
| ++.+ + |  
| +=.o |  
| o+S E |  
| . * o = |  
| + =+. |  
| .=o++. |  
| .oo++o.. |  
+---[SHA256]----+
```

Copy to node

```
$ ssh-copy-id -i ~/.ssh/id_rsa.pub miki@192.168.81.128  
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/c/Users/mikia/.ssh/id_rsa.pub"  
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are  
already installed  
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to  
install the new keys  
miki@192.168.81.128's password:
```

Number of key(s) added: 1

Now try logging into the machine, with: "ssh -i /c/Users/mikia/.ssh/id_rsa 'miki@192.168.81.128'"

and check to make sure that only the key(s) you wanted were added.

2 Install Docker

Sunday, 10 August 2025 9:11 pm

Technology:

Docker

<https://docs.docker.com/engine/install/ubuntu/#install-using-the-repository>

1. Set up Docker's apt repo

```
# Add Docker's official GPG key:  
sudo apt-get update  
sudo apt-get install ca-certificates curl  
sudo install -m 0755 -d /etc/apt/keyrings  
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc  
sudo chmod a+r /etc/apt/keyrings/docker.asc
```

Add the repository to Apt sources:

```
echo \  
  "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc]  
https://download.docker.com/linux/ubuntu \  
  $(./etc/os-release && echo "${UBUNTU_CODENAME:-$VERSION_CODENAME}") stable" | \  
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null  
sudo apt-get update
```

2. Install the Docker packages

```
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-  
plugin
```

3. Verify installation success

```
sudo docker run hello-world
```

3 Install K3s

Sunday, 10 August 2025 9:14 pm

Technology:

K3s

On Control plane node:

1. Update your system

```
# sudo apt update && sudo apt upgrade -y
```

2. Install curl if not already installed

```
# sudo apt install -y curl
```

3. Install k3s

```
# curl -sfL https://get.k3s.io | sh -
```

4. Wait for k3s to be fully installed and running

```
# Check the status
```

```
# sudo systemctl status k3s
```

5. Verify installation

```
# sudo k3s kubectl get nodes
```

6. Export kubeconfig for easy access

```
# sudo k3s kubectl config view --raw > ~/.kube/config  
# export KUBECONFIG=~/.kube/config
```

7. Make a user copy of kubeconfig so I can access cluster from username miki

```
mkdir -p ~/.kube  
sudo cp /etc/rancher/k3s/k3s.yaml ~/.kube/config  
sudo chown $USER:$(id -gn) ~/.kube/config  
chmod 600 ~/.kube/config  
unset KUBECONFIG  
export KUBECONFIG=$HOME/.kube/config  
kubectl get no
```

8. Check cluster -- master node shld be there

```
# kubectl get no
```

8. Retrieve join token from master node

```
# cat /var/lib/rancher/k3s/server/node-token
```

On Worker Node:

1. Join worker node to cluster

```
# curl -sfL https://get.k3s.io | K3S_URL=https://<MASTER_IP>:6443 K3S_TOKEN=<NODE_TOKEN> sh -
```

2. Check service

```
# systemctl status k3s-agent
```

Back on Master Node:

1. Verify both nodes are connected

```
# k get no
```

4 Simple App creation

Monday, 11 August 2025 8:48 pm

Step 1: Install Python Tools (from here is on Vscode terminal)

Run in windows terminal (in SampApp dir):

```
# python -m venv venv  
# .\venv\Scripts\activate
```

** If you get a security policy error, run:
\$ Set-ExecutionPolicy -ExecutionPolicy RemoteSigned -Scope Process
then try again.

```
# pip install flask
```

Step 2: Project Folder Structure

Create these files:

```
SampApp/  
    ├── k8s-manifests/  
    |    └── (contents will come later)  
    |  
    ├── static/  
    |    └── style.css  
    ├── templates  
    |    ├── index.html  
    |    └── success.html  
    ├── feedback.db  
    ├── requirements.txt  
    └── app.py
```

static/style.css

```
body {  
    font-family: Arial, sans-serif;  
    background-color: #f4f4f4;  
}  
.container {  
    width: 40%;  
    margin: auto;  
    background: white;  
    padding: 2em;  
    margin-top: 5%;  
    box-shadow: 0 0 10px rgba(0,0,0,0.1);  
}  
input, textarea {  
    width: 100%;  
    padding: 0.8em;  
    margin-top: 1em;  
}
```

```
button {
    padding: 0.8em 2em;
    margin-top: 1em;
}
```

templates/index.html

```
<html>
<head>
    <title>Feedback Form</title>
    <link rel="stylesheet" href="{{ url_for('static', filename='style.css') }}">
</head>
<body>
    <div class="container">
        <h2>Feedback Form</h2>
        <form method="POST">
            <input type="text" name="name" placeholder="Your Name" required>
            <input type="email" name="email" placeholder="Your Email" required>
            <textarea name="message" placeholder="Your Feedback" required>
        </form>
    </div>
</body>
</html>
```

templates/success.html

```
<!DOCTYPE html>
<html>
<head>
    <title>Success</title>
</head>
<body>
    <h2>Thanks for your feedback, {{ name }}!</h2>
</body>
</html>
```

requirements.txt

```
Flask
```

app.py

```
from flask import Flask, render_template, request
import sqlite3
import os
app = Flask(__name__)
DB_PATH = 'feedback.db'
def init_db():
    if not os.path.exists(DB_PATH):
        conn = sqlite3.connect(DB_PATH)
        c = conn.cursor()
        c.execute('''CREATE TABLE feedback
                    (id INTEGER PRIMARY KEY AUTOINCREMENT,
                     name TEXT NOT NULL,
                     email TEXT NOT NULL,
                     message TEXT NOT NULL)'''')
        conn.commit()
        conn.close()
```

```
@app.route('/', methods=['GET', 'POST'])
def index():
    if request.method == 'POST':
        name = request.form['name']
        email = request.form['email']
        message = request.form['message']
        conn = sqlite3.connect(DB_PATH)
        c = conn.cursor()
        c.execute("INSERT INTO feedback (name, email, message) VALUES
        (?, ?, ?)", (name, email, message))
        conn.commit()
        conn.close()
        return render_template('success.html', name=name)
    return render_template('index.html')
if __name__ == '__main__':
    init_db()
    app.run(host='0.0.0.0', port=5000)
```

5 Dockerfile

Wednesday, 3 September 2025 8:55 pm

Build Dockerfile in app's root dir

```
FROM python:3.10-slim
# Set env var
ENV PYTHONDONTWRITEBYTECODE=1
ENV PYTHONUNBUFFERED=1
# Set the working directory in the container
WORKDIR /app
# Copy requirements first (for caching efficiency)
COPY requirements.txt .
# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt
# Copy the rest of the application
COPY . .
# Expose port (must match app.py)
EXPOSE 5000
# Run the Flask app
CMD ["python", "app.py"]
```

6 Push image to Docker Hub

Sunday, 7 September 2025 1:26 pm

Download Docker desktop and make account on Dockerhub and desktop (I use Github acc for both).

Check that docker desktop is installed, do a docker info to check that there is output.

Open docker desktop.

On vscode Terminal, go into main app dir.

1. Build image

*docker desktop mac and windows uses BuildKit by default (buildx is part of it)

*Image and repo name shld be lower case

```
# docker buildx build --load -t username/repo_name:tag root_dir
```

```
docker buildx build --load -t mikiangelica/sampapp:v1 .
```

Check that image is created:

```
# docker images
```

2. Push image

Make sure to login with:

```
docker login
```

Push image:

```
docker push mikiangelica/sampapp:v1
```

*image name shld match repo name basically

Can check on docker hub UI if image is there

7 Pod yaml manifests

Sunday, 7 September 2025 1:41 pm

Create the different pod manifests (app, postgres, redis, nginx, busybox)

- **Sampapp** (main app)
- **Postgres** (db)
- **Redis** (cache)
- **Nginx** (reverse proxy for webapp AKA receives incoming HTTP requests and forwards to actual app)
- **Busybox** (for testing)

```
k8s-manifests/
├── sampapp.yaml
├── postgres.yaml
├── redis.yaml
├── nginx.yaml
└── busybox.yaml
```

Sampapp

```
apiVersion: v1
kind: Service
metadata:
  name: sampapp-service
spec:
  selector:
    app: sampapp
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 5000
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sampapp
  labels:
    app: sampapp
spec:
  replicas: 2
  selector:
    matchLabels:
      app: sampapp
  template:
    metadata:
      labels:
        app: sampapp
    spec:
      containers:
        - name: sampapp
          image: mikiangelica/sampapp:v1
```

```
ports:  
  - containerPort: 5000
```

Postgres

```
apiVersion: v1  
kind: Service  
metadata:  
  name: postgres-service  
spec:  
  selector:  
    app: postgres  
  ports:  
    - protocol: TCP  
      port: 5432  
      targetPort: 5432  
---  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: postgres  
  labels:  
    app: postgres  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: postgres  
  template:  
    metadata:  
      labels:  
        app: postgres  
    spec:  
      containers:  
        - name: postgres  
          image: postgres:14  
          ports:  
            - containerPort: 5432  
          env:  
            - name: POSTGRES_USER  
              value: "admin"  
            - name: POSTGRES_PASSWORD  
              value: "admin123"  
            - name: POSTGRES_DB  
              value: "sampdb"
```

Redis

```
apiVersion: v1  
kind: Service  
metadata:  
  name: redis-service  
spec:  
  selector:  
    app: redis  
  ports:  
    - protocol: TCP  
      port: 6379
```

```
    targetPort: 6379
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis
  labels:
    app: redis
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: redis
          image: redis:7
          ports:
            - containerPort: 6379
```

Nginx

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
```

```
- containerPort: 80
```

Busybox

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox
spec:
  containers:
    - name: busybox
      image: busybox
      command: ["sleep", "3600"]
      imagePullPolicy: IfNotPresent
```

8 Push yaml to Github

Sunday, 7 September 2025 2:11 pm

1. Download Git <https://git-scm.com/downloads/win>

Open git bash or use ur terminal. To use git via ur terminal, do the following.

Add Git to System PATH

1. Search for "Environment Variables" in Windows Start Menu
Open: "Edit the system environment variables"
2. In the System Properties window:
 - o Click "Environment Variables..."
3. Under System variables, find and select Path, then click Edit...
4. Click New, and paste the Git bin path.
The default is usually one of:
 - o C:\Program Files\Git\bin
 - o or C:\Program Files\Git\cmdIf you're unsure, open Git Bash and run:

where git

and copy the path.

5. Click OK on all dialogs to save.
6. Restart VS Code completely (close and reopen).
7. Open the VS Code terminal (PowerShell or CMD) and try:

git --version

2. Initialise git

git init

3. Create a .gitignore (optional but good practice)

Create a .gitignore file if needed (e.g., if you don't want to commit .env files or logs):

echo "*.env" >> .gitignore

- o This will create a .gitignore file - a file that tells Git which files and folders to skip when committing changes. These files won't be tracked, added or pushed to your repo
- o Examples of some files are r*.log, .env (secrets, PW, API keys), venv

4. Save your github repo URL so you can later push/pull from it easily

By convention, the first remote connection is called origin. (origin is the nickname for your repo)

- Set remote named origin

git remote add origin <https://github.com/mikiangelica/k8s-manifests.git>

- If you want to rename:

git remote set-url origin <https://github.com/mikiangelica/k8s-manifests.git>

- If you want to add a new remote (with new name)

```
git remote add k8s https://github.com/mikiangelica/sampapp.git
```

- Check connections:

```
git remote -v
```

Would show something like the below:

```
git remote -v
origin https://github.com/mikiangelica/k8s-manifests.git (fetch)
origin https://github.com/mikiangelica/k8s-manifests.git (push)
```

- This just shows where those actions would go if you did run them

5. Set Github credentials

```
git config --global user.name mikiangelica
git config --global user.email mikiangelica@gmail.com
```

6. Push files to Github

- Be in the correct folder (in the folder with the YAML manifests are)
- Add your YAML files to Git - git add tells Git that you would like the following files to be tracked and included in the next commit
 - Add all files in current folder

```
git add .
```

- Add specific files

```
git add webapp.yaml redis.yaml
```

- Commit your changes - git commit saves a snapshot of all the staged files into Git's local history - but only locally (in ur computer), still doesn't send anything to github yet; only in ur local repo.

```
git commit -m "Initial commit of k8s YAML manifests"
```

- Push to Git

```
git push -u origin master
```

*-u is --set-upstream, so it sets the default remote and branch for your local branch. So it will push your local master branch to the origin remote (connects remote conn to github repo URL) and sets the upstream tracking ref so in future you can just type git push/pull instead of git push/pull origin master

To check which branch you're in rn

```
git branch
```

*What are Git Branches?

- A branch is like a parallel Universe of your project
- It lets you work on features, bug fixes or experiments separately from the main (or master) branch
- Each branch is an independent line of development
- When you're done with the changes in that branch, you can merge it back into main
- Why use multiple Git branches?
 - Feature dev - feature/add-login - work on new features without disturbing main code
 - Bug Fix - fix/postgres-connection - isolates fixes so you can test safely
 - Experiment - test/redis-cluster - Try new ideas without risking the main app
 - Integration or release - release/v1.2 - prepare prod releases separately

- Hotfix - hotfix/broken-prod - Urgent fixes to prod env

To see default branch for the respective remote name:

```
git remote show origin
```

To switch branch

```
git checkout main
```

*Switches you to the main branch

To check branches

```
git branch # local branches
```

```
git branch -r # remote branches
```

```
git branch -a # local + remote
```

7. New Branch

- Create and change to the new branch

```
git checkout -b calico-istio-testing
```

How the new branch will look:

```
# git branch # adds local branch calico-istio-testing
* calico-istio-testing
  master
# git branch -a # adds local branch calico-istio-testing
* calico-istio-testing
  master
  remotes/origin/master
# git branch -r * no change as not connected to git yet
  origin/master
```

Netshoot yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: netshoot
  namespace: default
  labels:
    app: netshoot
spec:
  containers:
    - name: netshoot
      image: nicolaka/netshoot:latest
      command: [ "sleep", "3600" ] # keeps the pod running for testing
  restartPolicy: Always
```

- Add, Commit and push to github

```
git add netshoot.yaml
```

```
git commit -m "Add netshoot pod for Calico/Istio testing"
```

```
git push -u origin calico-istio-testing
```

*You will see on github UI, that the new branch is created with your files

- Open pull request on github

To check all commits:

```
git log --oneline --graph --all
```

```
* 38c4c3e (HEAD -> calico-istio-testing, origin/calico-istio-testing) Add netshoot pod for  
Calico/Istio testing  
* 1fb2bb4 (origin/master, master) Initial commit of k8s YAML manifests
```

Changes to any one file

If you make a change to any file and want to push it to Git

- Edit your file
- Change to respective branch, add, commit, push

```
git checkout master
```

```
git add nginx.yaml
```

```
git commit -m "Change to nginx.yaml"
```

```
git push
```

*The change shld reflect on github itself

Revert changes above

- Commit and push

```
git revert HEAD # creates a new commit that undoes changes from last one
```

```
git push (if in correct branch)
```

```
git push origin master (if diff branch)
```

8. Merge branch into master

As of now, we have pushed the new netshoot pod yaml manifest to github under the new branch calico-istio-testing but it is not yet merged into master branch.

- On Github UI, click on "Compare & pull request"
- Review changes and click "Create pull request"
- Click "Merge pull requests"

9. Delete branch after merge:

- On Github (remote)

- Just click on Delete branch on the branch

- OR on CLI: # git push origin --delete calico-istio-testing

*FYI after deleting, if u do git branch -r it will still show the deleted branch bc Git hasn't updated your local list of remote branches yet (it is a cached view)

To refresh:

```
# git fetch --prune
```

To make it auto-prune every fetch:

```
# git config --global fetch.prune true
```

- On Local Machine

- `# git checkout master` # get out of the branch u wna delete
- `# git branch -d branch_name` (e.g. `git branch calico-istio-testing`)

9 Access k3s cluster from local

Sunday, 7 September 2025 9:29 pm

1. Download the latest .exe for Windows:
 - Example: <https://dl.k8s.io/release/v1.30.1/bin/windows/amd64/kubectl.exe>
2. Move kubectl.exe to a folder (e.g., C:\kubectl) and add that folder to your **System PATH**:
 - Right-click on *This PC* → Properties → Advanced system settings → Environment Variables → Edit Path
3. Confirm installation:

kubectl version --client

4. Make .kube dir
 1. mkdir %USERPROFILE%\.kube
 2. cd %USERPROFILE%\.kube
 3. SSH into master
 - a. Sudo su -
 - b. Cp /etc/rancher/k3s/k3s.yaml /home/miki/k3s.yaml
 - c. Chmod +r /home/miki/k3s.yaml (actly I did 777)
 - d. Chown miki:miki /home/miki/k3s.yaml
 4. Back on local cmd prompt:

```
Scp miki@192.168.81.128:/home/miki/k3s.yaml %USERPROFILE%/.kube/config
```
5. Edit .kube/config file so that the IP is the master node's IP
6. Save and test with kubectl get no
7. **Bonus... I like to do on Git Bash. Easier.

```
mikia@miki MINGW64 ~/OneDrive/Documents/SampApp/sampapp (master)
$ k get no
NAME      STATUS    ROLES          AGE     VERSION
master1   Ready     control-plane,master  20d    v1.33.4+k3s1
worker1   Ready     <none>        20d    v1.33.4+k3s1
```

10 Deploy YAMLs on cluster

Sunday, 7 September 2025 10:13 pm

1. Git clone

On the local term (git bash), do:

```
# git clone https://github.com/mikiangelica/k8s-manifests.git
```

Now check there shld be a k8s-manifests folder with all the yaml manifests.

2. Create objects

```
kubectl apply -f .
```

3. Validate

```
kubectl get po -owide
```

11 Calico

Sunday, 7 September 2025 10:23 pm

Technology:

Calico (CNI)

Calico

- East-west traffic (Internal)
- Pod to Pod (L3)
- Without calico, pods can't even talk to each other at all
- Can control which pod talks to who by label/NS/IP (netpols)

Calico GlobalNetpol

- Can control egress (what traffic leaves cluster e.g. you can control if you don't want pods/svcs to communicate with outside world except GitHub)

Ingress

- North-south traffic (to external)
- Only controls ingress (traffic inbound)
- Cannot control egress

ServiceMesh Istio

- North-south traffic
- Replaces Ingress
- mTLS
- More granular control over traffic and how traffic moves
- Metrics

1. Check for flannel

```
Journalctl -u k3s | grep flannel # there shld be flannel output if there is embedded flannel  
ls /var/lib/rancher/k3s/agent/etc/cni/net.d/ # shld have a 10-flannel.conflist file
```

2. Test current connectivity

```
k exec -it netshoot -- sh  
ping <pod-IP>  
curl <pod-IP>:<port>
```

*redis:6379, postgres:5432, sampapp:5000, nginx:80

3. Drain both nodes

4. Stop k3s svc

```
Master node: systemctl stop k3s  
Worker node: systemctl stop k3s-agent
```

5. Disable and remove flannel

```
rm -rf /var/lib/rancher/k3s/agent/etc/cni # delete cni directories  
rm -rf /etc/cni/net.d # doesn't exist  
rm -rf /opt/cni # doesn't exist  
rm -rf /var/lib/rancher/k3s/agent/etc/flannel (can optionally delete -- AFTER disabling flannel)  
ip a | grep flannel
```

```
Ps aux | grep flannel # check that flannel is no longer running
```

```
Ip link delete flannel.1
```

Edit the /etc/systemd/system/k3s.service.env (master) OR /etc/systemd/system/k3s-agent.service.env (worker) file:

Add in : K3S_NO_FLANNEL=true

6. Start k3s svc

```
Master node: systemctl start k3s
```

```
Worker node: systemctl start k3s-agent
```

Seems like flannel kept coming back after each restart of k3s. So I decided to just uninstall k3s and reinstall with calico

1. Uninstall k3s and cleanup

Master Node:

```
sudo /usr/local/bin/k3s-uninstall.sh
```

Worker Node:

```
sudo /usr/local/bin/k3s-agent-uninstall.sh
```

Clean up:

```
sudo rm -rf /etc/rancher /var/lib/rancher /var/lib/kubelet /etc/cni /opt/cni /var/lib/cni
```

```
Rm -rf ~/.kube/config (if u wanna reset cluster config on ur local)
```

2. Install k3s without flannel

Master node:

```
curl -sfL https://get.k3s.io | INSTALL_K3S_EXEC="server --flannel-backend=none --disable-network-policy" sh -
```

Worker node:

```
curl -sfL https://get.k3s.io | K3S_URL=https://192.168.81.128:6443  
K3S_TOKEN="K1080e8d3f7b27ad61ef4ca4d60c0f53705c52f3aeb3c17a4d782cedcdd865c1430::server:60f6d6681f7e54098d82c7f01d32f64b" sh -
```

Confirm that there is no flannel:

```
Ip a | grep -i flannel
```

```
Ps aux | grep flannel
```

3. Push calico to Git

In Git bash, go into the SampApp/k8s-manifest and create new branch

```
Git checkout -b calico
```

Next, download calico from browser

1. Go to:

<https://raw.githubusercontent.com/projectcalico/calico/v3.30.3/manifests/calico.yaml>

2. Press Ctrl + S to save the file as calico.yaml
3. Move it to your k8s-manifests Git repo directory

Push to git

| |
|---|
| Git add calico.yaml |
| Git commit -m "Add calico yaml" |
| Git push -u origin calico # create new branch in git and push |

Delete calico branch

*Delete on UI

| |
|--|
| Git fetch --prune (will match UI and remove calico from remote branch) |
| Git pull origin master (update new merged master branch) |
| Git branch -d calico # make sure ur not in calico branch |

4. Download calico

I still cant kgn or kgp because k3s cannot fully initialise because there is no CNI.
So I could either use the calico.yaml from git or download directly to master:

| |
|--|
| curl -LO https://raw.githubusercontent.com/projectcalico/calico/v3.30.3/manifests/calico.yaml |
|--|

Apply calico

| |
|--|
| KUBECONFIG=/etc/rancher/k3s/k3s.yaml kubectl apply -f calico.yaml --validate=false |
|--|

Check on calico pod progress (1 controller and 1 node pod per node)

| |
|--|
| KUBECONFIG=/etc/rancher/k3s/k3s.yaml kubectl get pods -A |
|--|

5. Config for master node to access cluster

Btw add this line to end of .bashrc then source the file.

| |
|--------------------------------------|
| KUBECONFIG=/etc/rancher/k3s/k3s.yaml |
|--------------------------------------|

Now test with kubectl get no

6. Config for local to access cluster

Copy contents of /etc/rancher/k3s/k3s.yaml from master node to your local's ~/.kube/config

| |
|--|
| Sudo su - and copy contents of above file into /home/miki/k3s.yaml |
|--|

| |
|--|
| Change ownership and permissions and scp into local and change IP to master node IP. |
|--|

Calico shld be present now.

```
k get po -A -owide | grep calico
kube-system   calico-kube-controllers-7959b6fcdb-Fvcjz   1/1   Running   9 (123m ago)   12d   172.18.235.188   worker1   <none>      <none>
kube-system   calico-node-9dg8n   1/1   Running   9 (123m ago)   12d   192.168.81.128   master1   <none>      <none>
kube-system   calico-node-g83cv   1/1   Running   9 (123m ago)   12d   192.168.81.129   worker1   <none>
```

*There was a network disruption and nginx, postgres and netshoot pods had ImagePullBackoff error. I reconnected and ran "sudo k3s ctr image pull docker.io/library/nginx:1.25" and now all 3 pods are running. Might be because I reconnected to network

12 Network Policies

Monday, 8 September 2025 8:56 pm

Default Deny-all Ingress

```
# netpol-default-deny.yaml

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
  namespace: default
spec:
  podSelector: {}
  matchLabels:
    role: db
  policyTypes:
    - Ingress
```

Netshoot to Sampapp

```
# netpol-allow-netshoot-to-sampapp.yaml

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-netshoot-to-sampapp
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: sampapp
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: netshoot
```

Sampapp to db

```
# netpol-allow-sampapp-to-db.yaml

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-sampapp-to-db
  namespace: default
spec:
  podSelector:
```

```

    matchLabels:
      app: redis
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
      matchLabels:
        app: sampapp
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-sampapp-to-postgres
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: postgres
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
      matchLabels:
        app: sampapp

```

Busybox to nginx

```

# netpol-allow-busybox-to-nginx.yaml

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-busybox-to-nginx
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: nginx
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
      matchLabels:
        app: busybox

```

Policies made:

1. Default Deny-all Ingress

Go into the netshoot pod

```
Kubectl exec -it netshoot -n default -- sh
```

Try to connect to redis and postgres (shldnt work)

```
Nc -zv redis-service 6379
```

Nc -zv postgres 5432

2. Netshoot to Sampapp

Go into the netshoot pod

```
Kubectl exec -it netshoot -n default -- sh
```

Try to connect to sampapp (*shld work)

Curl sampapp-service:5000

3. Sampapp to db (redis and postgres)

****Btw I edited the sampapp yaml here so that there is a netshoot container in it.**

***Make sure to specify container below or it will default to first container sampapp

```
Kubectl exec -it sampapp-pod -n default -c netshoot -- sh
```

Nc -zv 172.16.235.136 6379 (redis)

Nc -zv 172.16.137.73 5432 (postgres)

Ran into an issue where it could not reach redis and postgres. The labels and pods were all correct. Turns out dns issue because it could reach the pod ips:

```
~ # nc -zv 172.16.235.136 6379
Connection to 172.16.235.136 6379 port [tcp/redis] succeeded!
~ # nc -zv 172.16.137.73 5432
Connection to 172.16.137.73 5432 port [tcp/postgresql] succeeded!
```

**After some troubleshooting I realise it is not a dns issue as Im able to use nslookup to the services. It is not a netpol issue perse because I can connect via the pod IPs - but because I try to connect via the service (which means going through the service cluster IP that forwards traffic to the relevant pod) -- so it is not directly trying to connect to the pod with the specified label in my netpol.

4. Busybox to nginx

Make sure there is a nginx-service there and exec into busybox pod

```
Kubectl exec -it busybox -- sh
```

Connect to nginx

Wget --spider nginx-service

OR

Wget nginx-service

Expected output:

```
/ # wget --spider nginx-service
```

Connecting to nginx-service (10.43.104.188:80)

remote file exists

```
/ # wget nginx-service
```

Connecting to nginx-service (10.43.104.188:80)

saving to 'index.html'

index.html 100%

ETA

'index.html' saved

13 PVC for postgres

Saturday, 13 September 2025 11:50 pm

Technology:

Postgres in K3s

We use PVCs for services that are stateful, like Postgres.

Stateful = Keeps important data that must persist across restarts

E.g.

Postgres > Stores databases/tables

MySQL > Stores user records

Kafka/Zookeeper > Store offsets/metadata

Stateless = Can be killed/restarted anytime without losing critical data (e.g. Nginx, frontend API)

They require persistent storage so that they have a stable, durable volume that survives pod restarts and migrations... OTHERWISE the data is stored on the pod's ephemeral (short term) container filesystem and if the pod restarts, all data is LOST!!

A.K.A Without PVC, your Postgres DB would vanish every time the pod dies.

To bind a PVC to postgres, we need a storageclass (SC), Persistent Volume (PV), and Persistent Volume Claim (PVC).

How they work together:

1. You define storage class (e.g. gp2)
2. Postgres PVC says "I need 10Gi from gp2"
3. K8s (via CSI driver) provisions a PV (e.g. EBS Volume)
4. The PVC is bound to that PV

Creating the PVC for Postgres

1. First, we check if there is an SC in the cluster with:

```
# kubectl get sc
```

| NAME | PROVISIONER | RECLAIMPOLICY | VOLUMEBINDINGMODE |
|----------------------|-----------------------|---------------|---------------------------------|
| ALLOWVOLUMEEXPANSION | AGE | | |
| local-path (default) | rancher.io/local-path | Delete | WaitForFirstConsumer false 6d6h |

This SC will dynamically create PVs for you whenever you make a PVC so don't need create PV.

Most times don't need to create PV unless in static provisioning cases... like u want to handcraft the PV and manually bind to a PVC etc.

-- Rule of thumb > If an SC exists in cluster, only need to create PVC. If no SC, need to create SC, PV and PVC

3. Replace postgres yaml from a deployment to a statefulset

Statefulsets:

- Guarantee stable pod names e.g. postgres-1 postgres-2
- Creates PVCs per replica using VolumeClaimTemplates
- Ensures the pod reattaches to the same volume when rescheduled.

```

apiVersion: v1
kind: Service
metadata:
  name: postgres-service
  namespace: default
spec:
  selector:
    app: postgres
  ports:
    - port: 5432
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: postgres
  namespace: default
spec:
  selector:
    matchLabels:
      app: postgres
  serviceName: "postgres"
  replicas: 1
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:15
          env:
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-secret
                  key: password
          ports:
            - containerPort: 5432
          volumeMounts:
            - name: postgres-storage
              mountPath: /var/lib/postgresql/data
  volumeClaimTemplates:
    - metadata:
        name: postgres-storage
      spec:
        accessModes: [ "ReadWriteOnce" ]
        storageClassName: "local-path"
        resources:
          requests:
            storage: 1Gi

```

Push above yaml manifest to github... create new branch, add, commit, push. Then create pull req on github and merge branch. Then back on CLI, git fetch --prune and git branch -d postgres.

Apply the new manifest...

Faced multiple issues from here:

ISSUE 1: Pod, PVC and PV not coming up

CAUSE: Provisioner pod wasn't even running - causing PVCs to be stuck in Pending state and failing to provision

FIX: Redeployed the local-path-provisioner into cluster by reapplying manifest

Troubleshooting Steps Overview: Noticed Pod was not coming up. Checked PVC and PV and realised they weren't coming up either. Then checked the provisioner pod (responsible for creating the local PVs based on PVC requests) and saw it wasn't running.

Reference: Troubleshooting Log 2

ISSUE 2: Pod still not coming up with 'calico' failure error (ContainerCreating state) - Calico Pods Issue

CAUSE: Calico was running on BGP mode.

FIX: Switched to VXLAN mode and reapplied VXLAN calico manifest with same initial version and these network interface calico errors cleared

Troubleshooting Steps Overview: Checked logs of pod, and saw that it had network interface (calico) errors, Checked calico pods & their logs, cluster status. Tested connectivity with netshoot and got same calico error. Reapplied calico manifest but same error observed. Checked if CNI binaries exist on worker node (calico, calico-ipam) for kubelet to call calico. Checked connectivity to API server (connected). Checked for CNI config file and calico crd existence (both fine). Checked kubeconfig file in worker node and did a curl to the ClusterIP API (reachable). Reapplied calico manifest again and now calico pods are not running. Checked logs and noticed it says BGP is not ready. Switched to VXLAN mode by editing default IPPool, reapplied calico manifest with vxlan mode with correct calico version and restarted pods.

Reference: Troubleshooting Log 2

ISSUE 3: CoreDNS failure found

CAUSE: Config map misconfiguration

FIX: Amended Config map

Troubleshooting Steps Overview: During pod to pod connectivity test in nettest ns, ping was failing by pod name - but via pod IP was fine. Checked coreDNS pod and svc. Checked DNS from test pods and logs of coredns pod, which showed that coreDNS is forwarding queries upstream to possibly host VM's DNS which is not reachable from inside CoreDNS pod. Amended coredns configmap so that it forwards to 8.8.8.8 1.1.1.1 instead. Restarted deployment and DNS was working.

Reference: Troubleshooting Log 2

ISSUE 4: (RELATED TO ISSUE 1) Local-path-storage pods not coming up

CAUSE: Broken config.

FIX: Reapply manifests and also create secret.

Troubleshooting Steps Overview: All local-path-storage pods were CrashLoopBackOff. Checked logs and saw error: cm is missing in ns. Removed finalizers block from ns config and force deleted broken namespace. Had to apply ns.json via API finalize endpoint to rly delete ns as it was stuck. Reapplied local-path-storage manifest and all local-path-storage pods are up - as well as the sc, pvc, pv of postgres. But postgres pod still CreateContainerConfigError. Checked logs and saw that it could not find postgres-secret secret. Had to create secret and verify. Restarted pod and all is well.

Reference: Troubleshooting Log 2

14 Cronjob for backups for Postgres

Tuesday, 16 September 2025 12:09 am

Over here we create a cronjob that connects to Postgres DB using the Secret created for it (matches postgres credentials), and collects dumps to store in the pvc mount /backups. This /backups is located in 1 of the worker nodes in my lab cluster - but in Prod clusters, this /backups is located in a dedicated DB node for the cluster. Hence why we have DB nodes.

1. Create pvc

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pg-backup-pvc
  namespace: default
spec:
  storageClassName: local-path
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
~
```

```
# k apply -f pg-backup-pvc.yaml
```

2. Create secret

```
kubectl create secret generic pg-backup-secret \
> --from-literal=POSTGRES_USER=postgres \
> --from-literal=POSTGRES_PASSWORD=admin123 \
> --from-literal=POSTGRES_DB=postgres
secret/pg-backup-secret created
```

*The POSTGRES_USER, PASSWORD, DB here must match the ones Postgres is actually running with (from STS secret) or else pg_dump will fail to connect.

This secret (pg-backup-secret) is for clients (like your backup job) to connect to the running Postgres DB.

3. Create cronjob

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: pg-backup
  namespace: default
spec:
  schedule: "0 2 * * *" # Every day at 2 AM
  jobTemplate:
    spec:
      template:
        spec:
          containers:
```

```

- name: pg-backup
  image: postgres:15
  imagePullPolicy: IfNotPresent
  envFrom:
    - secretRef:
        name: pg-backup-secret
  volumeMounts:
    - name: backup-storage
      mountPath: /backups
  command:
    - /bin/sh
    - -c
    - |
      FILE=/backups/backup-$(date +\%F-\%H%M).sql
      PGASSWORD=$POSTGRES_PASSWORD pg_dump -h postgres-service -U $POSTGRES_USER
$POSTGRES_DB > $FILE
      echo "$Backup saved to $FILE"
  restartPolicy: OnFailure
  volumes:
    - name: backup-storage
      persistentVolumeClaim:
        claimName: pg-backup-pvc
~
```

**Over here, we mount the PVC pg-backup-pvc at /backups.

- The pg_dump command writes SQL dumps into /backups/backup-<date>.sql
- That means the dump file is persisted on the PV provisioned by your local-path SC
- So even if the pod dies, the file stays on disk (wherever your local-path-provisioner stores PVs on the node.. Usually /var/lib/rancher/k3s/storage/...)

After creating the above, PVC was in Pending state as it the SC is WaitForFirstConsumer. So we created a manual trigger to test the cronjob immediately instead of waiting til 2AM.

- It clones my cronjob into a one time job so I can validate:
 - That the pvc binds successfully
 - the pod runs pg_dump
 - that a backup file gets written into your /backups mount

```
# kubectl create job --from=cronjob/pg-backup pg-backup-test
job.batch/pg-backup-test created
```

After this, I see that pvc is bound and pv is there

```
# k get pvc
NAME           STATUS VOLUME          CAPACITY ACCESS MODES
STORAGECLASS   VOLUMEATTRIBUTESCLASS AGE
pg-backup-pvc  Bound  pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb  1Gi     RWO
local-path     <unset>       5m30s
postgres-storage-postgres-0 Bound  pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d  1Gi     RWO
local-path     <unset>       22h

# k get pv
NAME           CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM
STORAGECLASS   VOLUMEATTRIBUTESCLASS REASON AGE
pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb  1Gi     RWO     Delete     Bound  default/pg-
backup-pvc      local-path  <unset>            2m9s
```

| | | | | |
|--|------------|---------|--------|-------|
| pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d | 1Gi | RWO | Delete | Bound |
| default/postgres-storage-postgres-0 | local-path | <unset> | | 15h |

15 GitHub Resync

Thursday, 18 September 2025 5:36 am

I made some changes directly to the cluster a few times during troubleshooting, so for this section I will be rearranging my github repo and taking my latest cluster yaml and configs and pushing them to GitHub.

Restructuring my GitHub repo

Intended Structure

```
SampApp/
  └── sampapp/          # All infra + manifests + automation
    ├── README.md
    ├── app/              # Your actual app code (Python etc)
    │   ├── requirements.txt
    │   ├── app.py
    │   └── ...
    ├── manifests/
    ├── terraform/
    ├── ansible/
    ├── Dockerfiles/
    └── images/
.gitignore
```

Edited Dockerfile

```
FROM python:3.10-slim

ENV PYTHONDONTWRITEBYTECODE=1
ENV PYTHONUNBUFFERED=1

WORKDIR /app

COPY app/requirements.txt .

RUN pip install --no-cache-dir -r requirements.txt

COPY app/ .

EXPOSE 5000

CMD ["python", "app.py"]
```

- Paths are relative to the build context!!
- I run docker build in SampApp/ dir, and when I run docker build I ran with sampapp as my build context (see the end), so in my Dockerfile, when I specify "app/" it is relative to sampapp/

Ran into the following issue:

ISSUE 1: ERROR BUILDING IMAGE

CAUSE: Build context was too big so it cancelled

FIX: To only include sampapp dir (context build)

Troubleshooting Steps Overview: Analysed error message from build command, reduced build context to just application directory

Reference: Troubleshooting Log 1

Resyncing my GitHub repo

```
# kubectl get all -n default -oyaml > manifests/base/default-all.yaml  
# kubectl get sc,pv,pvc -oyaml > manifests/storage/storage.yaml  
# kubectl get netpol -Aoyaml > manifests/netpols/netpols.yaml  
# kubectl get cm -Aoyaml > manifests/base/configmaps.yaml  
# kubectl get secrets -Aoyaml > manifests/base/secrets.yaml
```

README

Notes on Extracted Manifests

These manifests are direct exports from my cluster and may include system fields. In a production GitOps workflow, I would clean them using kubectl-net or yq.

Notes on Secrets

Secrets are also exported in raw base64-encoded form for completeness.

In a real production environment, **storing raw secrets in Git is a no-no**

Instead, a team would use Sealed Secrets/External Secrets Operator/HashiCorp Vault/SOPS+GitOps

This repo demonstrates my learning and workflow process. In a prod scenario I would use a proper secret management solution.

Push to Git

```
# git add manifests/ README.md  
# git commit -m "Add synced cluster YAMLs and updated README notes on secrets"  
# git push (didn't work because I completely restructured my local dirs)  
# git push origin master --force (added the new files forcefully)
```

On Github it still showed the old files...

```
# git ls-files (still showed the old files)  
(Checked if files are still actually there - they weren't)  
# git add .  
# git rm -r --cached . (untracks everything Git still thinks exists (incl old files))  
# git add . -k(stages all actual files/folders that still exists locally)  
# git commit -m "Remove old tracked files and sync repo structure"  
# git push origin master (syncs with github)
```

Now updated.

16 Ingress

Saturday, 13 September 2025 11:54 pm

Technology:

Ingress

Components in Ingress setup

1. Ingress Controller (ingress-nginx-controller Pod)
 - Traffic Cop / Reverse proxy (NGINX running inside K8s)
 - Watches the Ingress resources in the cluster
 - When the pod starts, it runs with an argument like:

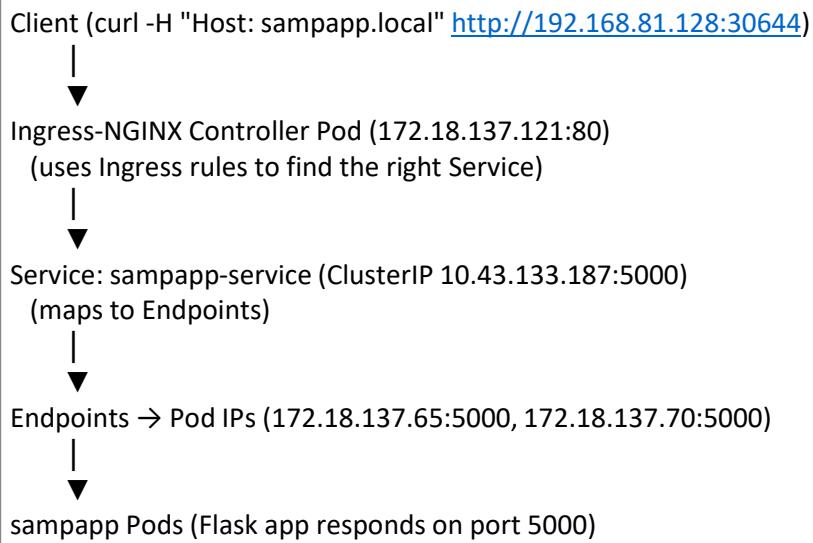
```
--ingress-class=nginx  
--controller-class=k8s.io/ingress-nginx
```
 - So the above means that it watches for ingress objects with ingressClassName= nginx and it syncs those rules into its own internal NGINX config
 - Actually accepts requests (on NodePort/LoadBalancer ports like :30644) and decides where to send them
 - It does not talk directly to Pods, it only talks to services
 - Role: 'entry point into your cluster'
2. Ingress Resource
 - Just configuration (rules)
 - Ingress resource will have a field (ingressClassName: nginx), which tells K8s that this ingress belongs to the controller that manages the class nginx.
 - Says things like:
 - If Host: sampapp.local and path / -> Send them to sampapp-service:5000
 - This is how the controller know which Service to target
 - Role: 'traffic rules' for the controller to follow
3. IngressClass (nginx)
 - This is a cluster-wide object that maps a name (nginx) to a specific controller

```
apiVersion: networking.k8s.io/v1  
kind: IngressClass  
metadata:  
  name: nginx  
spec:  
  controller: k8s.io/ingress-nginx
```
 - It says "Any Ingress with ingressClassName: nginx shld be handled by the Ingress-NGINX controller"
4. Service (sampapp-service)
 - Stable front door for your app's Pods
 - Has a ClusterIP (10.x.x.x) and points to Pod endpoints (172.x.x.x)
 - Service has selectors (labels like app=sampapp) which map it to Pods
 - Auto maintains an endpoints object listing all matching pods
 - Role 'load balancer inside the cluster'
5. App Pods (sampapp-xxx_)
 - These are the actual Flask containers running your code
 - They bind to port 5000 and serve requests

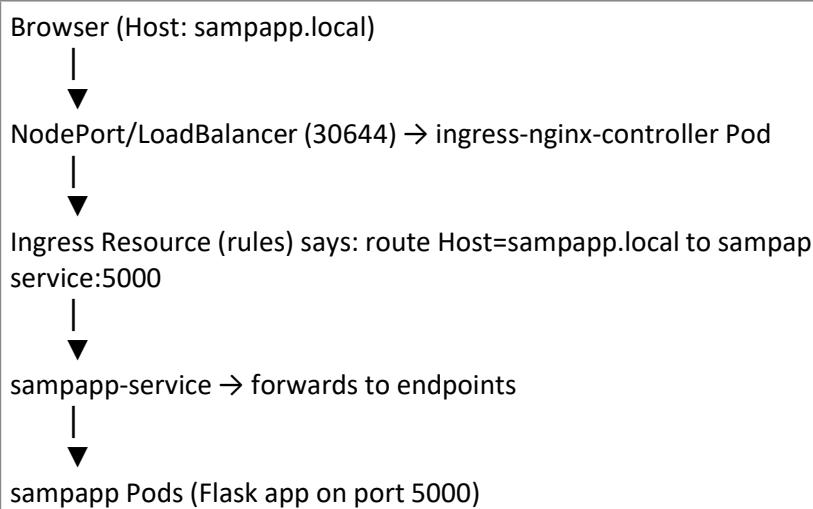
- Role 'where the real app lives'

Ingress Controller takes instruction from Ingress resource and carries it out. For IngressController and Ingress Resource to connect, they need IngressClass as the medium.

End to end Connection



Flow of a request



Implementation

1. Download version 3.19.0 of Helm
<https://github.com/helm/helm/releases>
2. Moved Helm to an existing path in env variables PATH
3. Install ingress-nginx (controller)


```
# helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
# helm repo update
# helm install ingress-nginx ingress-nginx/ingress-nginx
```

Ran into the various issues:

ISSUE 1: Installation failed

CAUSE: Existing resources preventing installation due to config mismatch

FIX: Deleted all current ingress resources

Troubleshooting Steps Overview: Analysed error message from install command, deleted existing ingress resource and verified helm release is deployed, and ingress pods, svc and resources are up. Mapped host locally (sampapp.local) since no external IP. Tested connectivity and it reached ingress-nginx controller! But next issue...

Reference: Troubleshooting Log 3

ISSUE 2 : Error 503 (Either ingress or svc issue)

CAUSE: Request is reaching ingress-nginx controller but can't find a backend svc to route to (svc/Ingress issue)

FIX: Edited ingress to match port of ingress svc

Troubleshooting Steps Overview: Check ingress and svc port and IP, found a port mismatch [5000 vs 80], Tested connectivity and now no more 503 error... now it is 504 error :(

Reference: Troubleshooting Log 3

ISSUE 3 : 504 error gateway timeout

CAUSE: No response was received within timeout duration and connection closed.

FIX: Deleted all evicted pods

Troubleshooting Steps Overview: Checked svc endpoints, pod IPs and labels. Svc and Endpoints were fine but there were many failed pods. Deleted all evicted pods til left with 2 running pods - matched EP IPs to these 2 running pods. Tested curl again and still 504 error.

Reference: Troubleshooting Log 3

ISSUE 4 Ingress svc external IP pending, svc loadbalancer helpers failing, local-path-provisioner is broken

CAUSE: No external IP for Ingress service, causing loadbalancer helpers to fail and local-path-provisioner to break

FIX: Quick workaround - use NodePort

Troubleshooting Steps Overview: Edited ingress-nginx-controller svc to change type to NodePort. Tested connectivity from inside ingress-controller pod and still got 504 error.

Reference: Troubleshooting Log 3

ISSUE 5 Nodes unhealthy (CPU and disk)

CAUSE: Pressure from disk also causing CPU pressure

FIX: Cleared disk space, then deleted all Failed pods (ephemeral storage) which alleviated CPU pressure too

Troubleshooting Steps Overview: Checked node health and saw that they had disk and CPU pressure. Cleared disk space, deleted failed pods and verified healthy Node pressure again

Reference: Troubleshooting Log 3

ISSUE 5 FIXED

Now back to checking ingress

1. Check node health

```
kubectl get nodes -o wide  (all shld be Ready)  
kubectl describe node master1 | grep -A5 Conditions  (all shld be False)  
kubectl describe node worker1 | grep -A5 Conditions  (all shld be False)
```

2. Check that all ingress controller pods are running

```
# kubectl get pods -n ingress-nginx -o wide  
NAME           READY   STATUS    RESTARTS   AGE   IP  
NODE   NOMINATED-NODE  READINESS GATES  
ingress-nginx-controller-6b96cf684f-qz8s9  1/1   Running   0   20m  
172.18.137.121  master1 <none>      <none>
```

3. Describe pod

```
# kubectl describe pod -n ingress-nginx <ingress-nginx-pod-name>
```

4. Check Ingress object

```
# kubectl get ingress -n default  
# kubectl describe ingress sampapp-ingress -n default
```

5. Check that service ports for ingress are 80:xxxxxx/TCP and 443:yyyyyy/TCP

```
# k get svc -n ingress-nginx  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)  
AGE  
ingress-nginx-controller      NodePort      10.43.230.85  <none>  
80:30644/TCP,443:30589/TCP  10h  
ingress-nginx-controller-admission  ClusterIP  10.43.251.154 <none>  
443/TCP                      10h
```

Take note of the NodePort

```
# kubectl get svc sampapp-service -n default -oyaml | egrep "port:  
|targetPort:"
```

```
# kubectl get ep sampapp-service-n defualt -owide
```

Shld be service port 5000/TCP, ep shld be IPs of sampapp pods

6. Check backend pods

```
# kubectl get pods -n default -l app=sampapp -o wide  
# kubectl logs -n default <one-sampapp-pod> -c sampapp
```

7. Test Ingress from inside cluster

```
# kubectl run curltest --rm -it --image=curlimages/curl --restart=Never -- \  
curl http://sampapp-service.default.svc.cluster.local:5000
```

And from outside

```
# curl -H "Host: sampapp.local" http://<node-ip>:<nodeport>
```

Curl was still not working!

ISSUE 6 Netpol

CAUSE: My test netpols were blocking traffic!

FIX: Deleted all netpols (will be implementing properly in future)

Troubleshooting Steps Overview: Was checking connectivity from ingress-nginx pod when I remembered the netpols, then deleted all netpols. Connectivity is now working.

Reference: Troubleshooting Log 3

Push ingress yaml to Github

```
# kubectl get ingress -A -o yaml > ingress/all-ingresses.yaml  
# kubectl get ingressclass -o yaml > ingress/all-ingressclasses.yaml  
# kubectl get ingress sampapp-ingress -n default -o yaml >  
ingress/sampapp-ingress.yaml
```

Went back to sampapp dir and:

```
# git add .  
# git commit -m "Update with Ingress"  
# git push
```

17 CI/CD w Jenkins & ArgoCD

Saturday, 13 September 2025 11:54 pm

Technology:

GitOps

GitHub - Source of Truth for manifests

CI - Jenkins (build image)

CD - ArgoCD (applies manifests)

Pros: Auditable, GitOps best practice, clean separation of roles

Cons: More moving parts (Jenkins + ArgoCD)

Jenkins is a CI/CD Tool. It can build and deploy, which makes things simple if we just have Jenkins. It will:

- Checkout repo > Build code
- Build & Push docker image
- Update your deployment YAML with new image tag
 - e.g. patch sampapp-deployment.yaml image to myrepo/sampapp:<newtag>
- Commit and Push that YAML change back to GitHub
- Apply k8s manifests directly to cluster (Jenkins needs a kubeconfig with access to cluster)

But generally this is not ideal for GitOps. Jenkins has cluster access. It is harder to audit and rollback as manifests may drift from GitHub.

With a CD Tool like ArgoCD, Jenkins will only have the CI role:

Jenkins will:

- Checkout repo > Build code
- Build & Push docker image
- Update your deployment YAML with new image tag
 - e.g. patch sampapp-deployment.yaml image to myrepo/sampapp:<newtag>
- Commit and Push that YAML change back to GitHub

ArgoCD will:

- Detect the change in GitHub
- Sync the new manifest into the cluster

So with this setup:

- Jenkins (CI+CD) -> Jenkins directly

Jenkins (CI side: build & push image)

1. Using Helm, deploy Jenkins on cluster

```
# helm repo add jenkins https://charts.jenkins.io
# helm repo update
# helm install jenkins jenkins/jenkins -n jenkins --create-namespace
```

2. Got Jenkins admin pw (on windows via Git Bash)

```
# kubectl get secret jenkins -n jenkins -o jsonpath="{.data.jenkins-admin-password}" | base64 --decode
```

&& echo

3. Expose Jenkins via Port-forward (easiest and fastest for dev/test)

Later once ingress-nginx is stable, can expose Jenkins behind an Ingress rule (jenkins.local).

```
# kubectl port-forward svc/jenkins -n jenkins 8080:8080
```

4. Login to Jenkins :

Open <http://localhost:8080> in browser.

Username: admin

PW: the PW above

5. Install recommended Plugins in Jenkins

Must-Have Plugins for Your Setup

1. Core CI/CD

- **Pipeline** → lets you use Jenkinsfile for declarative pipelines.
- **Git** → integrates with Git repositories.
- **GitHub Integration** → enables GitHub webhooks, builds on push.
- **GitHub Branch Source** → auto-detects branches and PRs.

2. Container & Kubernetes

- **Docker Pipeline** → allows docker.build() and docker.push() inside Jenkinsfile.
- **Kubernetes** (a.k.a. Kubernetes plugin) → run Jenkins agents as pods in your cluster (optional, but very useful).
- **Kubernetes Credentials Provider** → stores secrets in Kubernetes and syncs them into Jenkins.

3. Credentials & Secrets

- **Credentials Binding** → lets you use Docker registry passwords, GitHub tokens, etc. safely in pipelines.
- **Config File Provider** → store kubeconfigs, YAMLs, or .npmrc securely for pipelines.

4. Notifications / Quality of Life

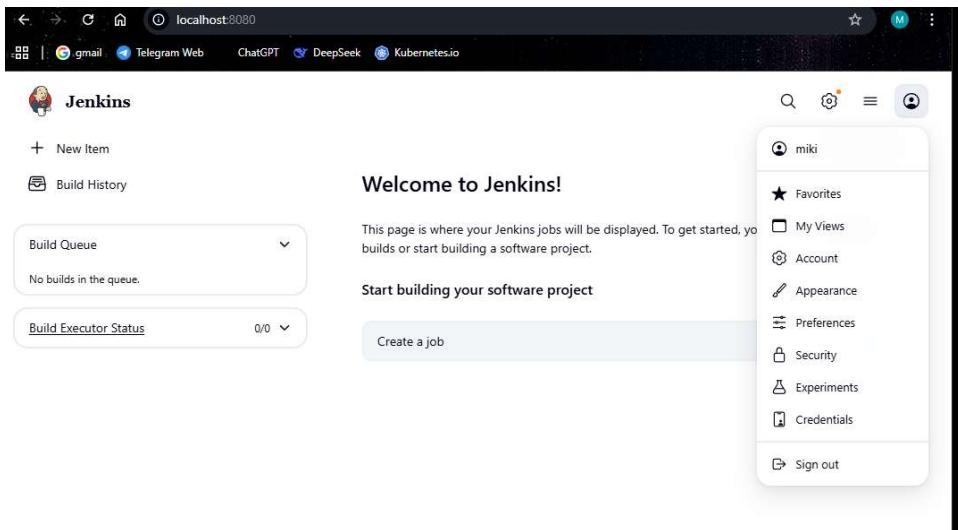
- **Blue Ocean** → a modern UI for pipelines (optional, but makes pipelines much clearer).
- **Email Extension** → send pipeline result emails.
- **Slack Notification** (if you use Slack).

How the plugins Fit Into Your Flow

- a. **GitHub Plugins** → Jenkins watches your sampapp repo.
- b. **Pipeline + Docker Plugins** → Jenkins builds/pushes Docker image to your registry.
- c. **Kubernetes Plugin** → optionally, Jenkins jobs run in ephemeral pods inside your k8s cluster.
- d. **Credentials Binding** → safely pass registry login (e.g., DockerHub password, GitHub PAT).
- e. **Config File Provider** → handle kubeconfigs or ArgoCD CLI tokens.
- f. **ArgoCD** → not a Jenkins plugin; ArgoCD just watches your repo for updated manifests.
Jenkins doesn't deploy — it only commits manifest changes.

6. Created new admin user

Manage Jenkins > Users > Create User



ArgoCD (CD - Deploy)

1. Install ArgoCD

```
# kubectl create namespace argocd
# kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
# kubectl -n argocd get pods
# all should become Running/Ready (api-server, repo-server, server, dex, redis, notifications)
```

2. Access ArgoCD UI

```
# kubectl -n argocd port-forward svc/argocd-server 8080:80
```

Go to <http://localhost:8080>

3. Get the initial admin password

```
# kubectl -n argocd get secret argocd-initial-admin-secret \
-o jsonpath="{.data.password}" | base64 --decode && echo
```

ISSUE 1: EPHEMERAL STORAGE OF NODES CAUSING ALL PODS TO FAIL

CAUSE: Ephemeral storage was very high, causing CPU pressure too.

FIX: Capacity Add (Disk)

Troubleshooting Steps Overview: All pods were failing with error messages of ephemeral storage issues. Checked node health and this confirmed it. Evaluated and confirmed that disk space needed to expand to accommodate more workloads. Expanded disk and restored cluster to working state.

Reference: Troubleshooting Log 4

ISSUE 2: LOCAL-PATH-PROVISIONER-PODS IN CRASHLOOPBACKOFF (RBAC)

CAUSE: Local-path-provisioner does not have RBAC permission to read the local-path-config CM in its namespace

FIX: Gave service account the correct RBAC role

Troubleshooting Steps Overview: While checking status of all pods, saw this pod failing and checked logs. Verified that all PVs, PVCs are fine. Gave SA the correct RBAC role and restarted pods.

Reference: Troubleshooting Log 4

Continue with installing ArgoCD

1. Delete argocd namespace with all its resources (since there are many broken argoCD pods due to the ephemeral storage issue last time)

```
# k delete ns argocd --grace-period=0 --force
```

2. Create ns again

```
# k create ns argocd
```

3. Apply official argocd manifest

```
# kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

This installs all core components:

- argocd-server (UI + API)
- argocd-repo-server (git ops)
- argocd-application-controller (reconciler)
- argocd-dex-server (SSO/auth)
- argocd-redis (caching)

4. Verify pods are running

```
# kubectl get pods -n argocd
NAME                  READY   STATUS    RESTARTS   AGE
argocd-application-controller-0      1/1    Running   0          57s
argocd-applicationset-controller-54f96997f8-m2bjb 1/1    Running   0          57s
argocd-dex-server-798cbff4c7-pg447     1/1    Running   0          57s
argocd-notifications-controller-644f66f7df-qw6tc 1/1    Running   0          57s
argocd-redis-6684c6947f-2lkwq        1/1    Running   0          57s
argocd-repo-server-6fccc5759b-lz7hl 1/1    Running   0          57s
argocd-server-64d5fc5bd58-8qbdt     1/1    Running   0          57s
```

5. Get ArgoCD initial admin PW

```
# k get secret argocd-initial-admin-secret -n argocd -o jsonpath=".data.password" | base64 -d
&& echo
```

6. Since there's no LoadBalancer here, we expose via NodePort:

```
# kubectl patch svc argocd-server -n argocd -p '{"spec": {"type": "NodePort"}}'
service/argocd-server patched
```

```
# k get svc argocd-server -n argocd
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)           AGE
argocd-server  NodePort  10.43.157.58 <none>       80:31805/TCP,443:32619/TCP  6m8s
```

7. Access ArgoCD UI

<https://<VM-IP>:<Node-Port>>

<https://192.168.81.128:31805> (master IP)

Back up entire cluster to GitHub again (before we test out the CI/CD)

1. Application Manifests (GitOps style)
➤ The following will capture all that should run but not the PVC contents)

```

# All resources in all namespaces (except cluster-wide CRDs etc.)
kubectl get all --all-namespaces -o yaml > cluster-all.yaml

# ConfigMaps and Secrets
kubectl get cm,secrets --all-namespaces -o yaml > cluster-configs.yaml

# Storage resources
kubectl get sc,pv,pvc --all-namespaces -o yaml > cluster-storage.yaml

# Network policies
kubectl get netpol --all-namespaces -o yaml > cluster-netpols.yaml

```

Set up a GitHub > Jenkins > ArgoCD pipeline

1. Created a Jenkins file in repo root

```

pipeline {
    agent any
    stages {
        stage('Checkout') {
            steps {
                git branch: 'master',
                url: 'https://github.com/mikiangelica/k8s-manifests.git'
            }
        }
        stage('Lint Manifests') {
            steps {
                sh 'kubectl kustomize k8s-manifests/ > /dev/null || true'
            }
        }
        stage('Build & Push Image') {
            steps {
                sh ""
                docker build -t mikiangelica/sampapp:${BUILD_NUMBER} .
                docker push mikiangelica/sampapp:${BUILD_NUMBER}
                ""
            }
        }
        stage('Update Manifests') {
            steps {
                sh ""
                sed -i "s|image:.*$|image: mikiangelica/sampapp:${BUILD_NUMBER}|"
                sampapp/manifests/app/deployment.yaml
                git config --global user.email "mikiangelica@gmail.com"
                git config --global user.name "jenkins"
                git commit -am "Update image to build ${BUILD_NUMBER}"
                git push origin master
                ""
            }
        }
    }
}

```

The above will create a new pipeline job in Jenkins and connect it to my GitHub repo (manifests + app code)

2. ArgoCD : set up an application

Point ArgoCD to the GitHub repo and path with manifests

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: sampapp
  namespace: argocd
spec:
  project: default
  source:
    repoURL: https://github.com/mikiangelica/k8s-manifests.git
    targetRevision: master
    path: manifests
  destination:
    server: https://kubernetes.default.svc
    namespace: default
  syncPolicy:
    automated:
      prune: true
      selfHeal: true
```

Apply

```
# k apply -f argocd-app.yaml
application.argoproj.io/sampapp created

# k get app -A
NAMESPACE  NAME    SYNC STATUS  HEALTH STATUS
argocd     sampapp  Unknown   Healthy
```

What the columns mean

- SYNC STATUS
 - Synced → manifests in cluster match what's in GitHub.
 - OutOfSync → something in GitHub differs from cluster.
 - Unknown → ArgoCD hasn't synced yet or doesn't know the state.
- HEALTH STATUS
 - Healthy → the deployed pods/services are running fine.
 - Progressing → deployment is rolling out.
 - Degraded → something failed (crashloop, not ready).
 - Unknown → ArgoCD hasn't checked health yet.

ISSUE 3 : UNKNOWN SYNC STATUS

CAUSE: argocd Netpol that came with manifest was blocking communication to argocd redis pod

FIX: Deleted all argocd netpols too.

Troubleshooting Steps Overview: Tried various DNS and connectivity tests within argocd NS between argocd repo server pod, argocd application controller and redis. Checked logs and configs of services and pods. Tried to create new netpol to explicitly allow connection but it didn't work. Deleted all netpols

Reference: Troubleshooting Log 4

NOW my cluster sync status is Synced.

```
# k get app sampapp -n argocd
NAME      SYNC STATUS  HEALTH STATUS
sampapp   Synced       Healthy
```

NOW LETS MAKE A PIPELINE.

Push code to GitHub > Jenkins pulls repo, builds Docker image, pushes to DockerHub, commits updated manifests back to GitHub > ArgoCD detects repo changes, syncs manifests and deploys new image to cluster

1. First add credentials to Jenkins for GitHub and DockerHub. We will create the Access Tokens in both before adding to Jenkins.

In Github UI:

- Go to **GitHub → Settings** (top-right corner, click your profile picture → *Settings*).
- On the left sidebar, scroll down to **Developer settings**.
- Click **Personal access tokens → Tokens (classic)**.
- Click **Generate new token → Generate new token (classic)**.
 - (You can also use “Fine-grained tokens”, but for Jenkins the classic ones are simpler.)

Configure the token

- **Note:** Give it a name like *Jenkins Pipeline Token*.
- **Expiration:** Choose a duration (90 days, 1 year, or “No expiration” if it’s a private lab).
- **Scopes to check** (minimum for Jenkins CI/CD):
 - repo (Full control of private repositories — needed if Jenkins will clone/push code)

- workflow (if you ever want Jenkins to trigger GitHub Actions)
- [read:org](#) (optional, if your repo is under an org and Jenkins needs to read it)

Generate and Copy token

In DockerHub UI:

- Go to **DockerHub** → Click on profile > Account Settings > Personal access tokens
- Click on Generate New access token
 - Access token description: jenkins-ci
 - Expiration date: 90 days
 - Access Permissions: Read, Write, Delete

Generate and Copy token

dckr_pat_AnjFz8Qj6fQtssHvYcwAVoaQ_5c

In Jenkins,

I add credentials

Manage Jenkins > Credentials > (system) Global:

- **GitHub:**
 - Kind: Username with password
 - Scope: Global
 - Username: mikiangelica
 - Password: PAT
 - ID: github-pat
 - Description: GitHub PAT for Jenkins pipelines
- **Docker Hub:**
 - Kind: Username with password
 - Scope: Global
 - Username: mikiangelica
 - Password: PAT
 - ID: dockerhub-creds
 - Description: DockerHub PAT for Jenkins pipelines

The screenshot shows the Jenkins Global credentials (unrestricted) page. At the top, there's a navigation bar with Jenkins logo, Manage Jenkins, Credentials, System, and Global credentials (unrestricted). Below the navigation is a search bar and a 'Add Credentials' button. The main area is titled 'Global credentials (unrestricted)' and contains a table with two rows of credentials.

| ID | Name | Kind | Description |
|-----------------|--|------------------------|-------------------------------------|
| dockerhub-creds | mikiangelica/***** (DockerHub PAT for Jenkins pipelines) | Username with password | DockerHub PAT for Jenkins pipelines |
| github-pat | mikiangelica/***** (GitHub PAT for Jenkins pipelines) | Username with password | GitHub PAT for Jenkins pipelines |

At the bottom left, there are icons for S, M, and L.

2. Create a new pipeline job in Jenkins

- Go to Jenkins Dashboard > New Item
- Enter name : sampapp-pipeline
- Choose Pipeline and click OK

Configure SCM (GitHub Repo)

Inside the job:

- a. Scroll to **Pipeline** section.
- b. Select **Pipeline script from SCM**.
- c. **SCM** → Git.
- d. Enter your **Repository URL** (e.g., <https://github.com/mikiangelica/sampapp.git>).
- e. **Credentials** → choose the GitHub PAT credentials you added.
- f. **Branch Specifier** → */master (since your branch is master).

This tells Jenkins to pull your code + manifests from GitHub.

The screenshot shows the Jenkins dashboard. At the top, there's a search bar, a user icon, and several global filters. Below that is a 'Build History' section with a 'New Item' button. In the center, there's a table for the 'Build Queue'. The table has columns: S (Status), W (Workload), Name (sorted), Last Success, Last Failure, Last Duration, and F (File). There's one entry: 'sampapp-pipeline' with status 'S' (green), workload 'W' (yellow), last success 'N/A', last failure 'N/A', last duration 'N/A', and file icon 'F' (blue star). To the left of the table is a 'Build Queue' dropdown showing 'No builds in the queue.' Below the table is a 'Build Executor Status' section with '0/0' builds. At the bottom, there are icons for 'Icon: S M L' and a 'More' button.

Back up Cluster

```
# Base resources
$ kubectl get all --all-namespaces -o yaml > base/all.yaml

# Storage
$ k get sc,pv,pvc -Ao yaml > ./storage/storage.yaml

# Configs and secrets
$ k get cm,secrets -Aoyaml > configs.yaml

# Netpols
$ k get netpol -Aoyaml > netpols/netpols.yaml

# Apps
$ kubectl get applications.argoproj.io -A -o yaml > argocd/apps.yaml

# Appprojects
$ kubectl get appprojects.argoproj.io -A -o yaml > argocd/appprojects.yaml

# CRDs
$ kubectl get crd -o yaml > crds/all-crds.yaml

# NS
$ k get ns -oyaml > base/namespace.yaml

# SA, Roles, Rolebindings, ClusterRoles, ClusterRoleBindings
kubectl get sa,role,rolebinding,clusterrole,clusterrolebinding -A -o yaml > rbac/rbac.yaml

# Ingress
$ kubectl get ingress -A -o yaml > networking/ingresses.yaml
```

Backup to github.

Add Jenkinsfile to repo root (same level as /manifests, Dockerfiles/ etc)

```
$ cat Jenkinsfile
pipeline {
    agent any
    environment {
        IMAGE_NAME = "mikiangelica/sampapp"
        IMAGE_TAG = "v1"
    }

    stages {
        stage('Checkout Code') {
            steps {
                git branch: 'master',
                    credentialsId: 'github-pat',
                    url: 'https://github.com/mikiangelica/k8s-manifests.git'
            }
        }

        stage('Build & Push Docker Image') {
            steps {
                withCredentials([usernamePassword(credentialsId: 'dockerhub-creds',
                    usernameVariable: 'DOCKER_USER',
                    passwordVariable: 'DOCKER_PASS')]) {
                    sh """
                        docker build -t $IMAGE_NAME:$IMAGE_TAG -f Dockerfiles/Dockerfile .
                        echo $DOCKER_PASS | docker login -u $DOCKER_USER --password-stdin
                        docker push $IMAGE_NAME:$IMAGE_TAG
"""
                }
            }
        }
        stage('Update Manifests') {
            steps {
                sh """
                    sed -i 's|image: .*|image: $IMAGE_NAME:$IMAGE_TAG|'
manifests/base/deployment.yaml
                git config user.email "mikiangelica@gmail.com"
                git config user.name "admin"
                git add manifests/base/deployment.yaml
                git commit -m "Update image to $IMAGE_NAME:$IMAGE_TAG" || echo "No changes
to commit"
                git push origin master
"""
            }
        }

        stage('Trigger ArgoCD Sync') {
            steps {
                withCredentials([string(credentialsId: 'argocd-token', variable: 'ARGOCD_TOKEN')])
            }
        }
    }
}
```

```
sh """
curl -k -H "Authorization: Bearer $ARGOCD_TOKEN" \
-X POST https://192.168.81.128:31805/api/v1/applications/sampapp/sync
"""
"""

    }
}
}
}
}
```

Take note of the following:

- dockerhub-creds → with the ID of your DockerHub credential in Jenkins.
 - github-pat → with the ID of your GitHub credential in Jenkins.
 - your-dockerhub-username/sampapp → your DockerHub repo.
 - <ARGOCD_SERVER> → your ArgoCD API server (e.g. argocd-server.argocd.svc.cluster.local:443 or NodePort IP). - how I access UI basically
 - <ARGOCD_TOKEN> → you'll need to generate a token (argocd account generate-token).

Get argocd server IP

```
$ k get svc -n argocd argocd-server
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
argocd-server  NodePort  10.43.157.58  <none>        80:31805/TCP,443:32619/TCP  4d
```

- **ClusterIP:** 10.43.157.58 → only reachable *inside* the cluster.
 - **NodePort mapping:**
 - Port 80 is exposed at **NodePort 31805**
 - Port 443 is exposed at **NodePort 32619**

Get argocd token

By default the **admin account in ArgoCD does not allow API key/token generation**, unless you explicitly enable it in the config. For the purpose of this lab we will use admin account. Normally for actual setup we wld have a dedicated Jenkins Service account.

- Edit argocd-cm config map so that admin account can generate token

```
# kubectl edit configmap argocd-cm -n argocd
```

Under data:, add:

data:

accounts.admin: apiKey

- Restart deploy

```
# kubectl rollout restart deployment argocd-server -n argocd
```

- Go to ArgoCD UI > Settings > Accounts > Admin > Generate New (token)

Created new Secret Text credential on Jenkins with above token called argocd-token

Create deployment.yaml specifically for my app since I removed all initial files and my current app manifest is in all.yaml.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sampapp
  namespace: default
spec:
```

```

replicas: 1
selector:
  matchLabels:
    app: sampapp
template:
  metadata:
    labels:
      app: sampapp
spec:
  containers:
    - name: sampapp
      image: mikangelica/sampapp:v1 # <-- Jenkins will update this line
      ports:
        - containerPort: 8080

```

Push above to Git.

Test the Pipeline in Jenkins

Go to Jenkins Job > "Build Now"

- Watch the stages: Checkout → Build → Push → Update Manifests → Trigger ArgoCD.
- Confirm on GitHub that deployment.yaml gets updated with the new tag.
- Confirm in ArgoCD UI that the app syncs and your pods pull the new image from DockerHub.

```

#1 (Sep 25, 2025, 9:56:49 AM) Add description Keep this build forever Started 2 min 27 sec ago
Took 1 min 50 sec Started by user Jenkins Admin This run spent: 1 min 31 sec waiting; 1 min 50
sec build duration; 1 min 50 sec total from scheduled to completion. Revision:
f8bf15ec15040f22b06a4f90a63c59253e509269 Repository:
https://github.com/mikangelica/sampapp.git refs/remotes/origin/master Revision:
f8bf15ec15040f22b06a4f90a63c59253e509269 Repository: https://github.com/mikangelica/k8s-manifests.git refs/remotes/origin/master The following steps that have been detected may have
insecure interpolation of sensitive variables (click here for an explanation): sh: [DOCKER_PASS]
No changes.

```

ISSUE 4: BUILD FAILED 1 (same tag version)

CAUSE: Current deployment.yaml and image tag in DockerHub are same - v1

FIX: Changed deployment.yaml's image tag to oldtag instead of v1

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Changed deployment.yaml's image tag to oldtag instead of v1.

Reference: Troubleshooting Log 4

ISSUE 5: BUILD FAILED 2 (Indentation Error)

CAUSE: Indentation error in Jenkinsfile

FIX: Added 4 tabs in Jenkinsfile before image to handle indentation

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Added 4 tabs in Jenkinsfile before image to handle indentation

Reference: Troubleshooting Log 4

ISSUE 6: BUILD FAILED 3 (Docker not found)

CAUSE: Docker not found bc jenkins agent does not by default include Docker

FIX: Made jenkins-agent file for jenkins to use

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Made jenkins-agent file for jenkins to use with docker. Registered this custom agent in Jenkins UI and Jenkinsfile.

Reference: Troubleshooting Log 4

ISSUE 7: BUILD FAILED 4 (Didn't push above Jenkinsfile changes to GitHub - silly)

CAUSE: Jenkins detects from GitHub and I didn't push it there!

FIX: Pushed image

Troubleshooting Steps Overview: I realised it after I pressed Build Now button for the 4th time...

Reference: Troubleshooting Log 4

ISSUE 8: BUILD FAILED 5 (Could not find user)

CAUSE: Amended user in Dockerfile.jenkins-agent

FIX: Amended user in Dockerfile.jenkins-agent from admin to jenkins

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Then amended Dockerfile.jenkins-agent USER at end of file. Repushed Dockerfile to DockerHub

Reference: Troubleshooting Log 4

ISSUE 9: BUILD FAILED 6 (No docker daemon)

CAUSE: No Docker daemon running inside the Jenkins agent container

FIX: Added hostpath volume to jenkins pod template for jenkins agent (hosts' docker daemon socket file)

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Added hostpath volume to jenkins pod template for jenkins agent (hosts' docker daemon socket file)

Reference: Troubleshooting Log 4

ISSUE 10: BUILD FAILED 7 (Permission denied error)

CAUSE: Usually the above socket file is owned by root and group docker on the host.

But in our container, Jenkins run as user jenkins (not root, and not in docker group)

So when docker CLI tries to talk to the socket, the kernel blocks it > permission denied.

FIX: Removed last USER jenkins line from Dockerfile.jenkins-agent-build

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Removed last USER jenkins line from Dockerfile.jenkins-agent-build and repushed to DockerHub

Reference: Troubleshooting Log 4

ISSUE 11: BUILD FAILED 8

CAUSE: Wrong path in application Dockerfile (because I had moved my Dockerfile to a subdir after creating above custom jenkins agent Dockerfile)

FIX: Amended path to my sampapp Dockerfile

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Amended path to my sampapp Dockerfile

Reference: Troubleshooting Log 4

ISSUE 11: BUILD FAILED 8 (Dockerfile: no such file or dir)

CAUSE: Wrong path in application Dockerfile (because I had moved my Dockerfile to a subdir after creating above custom jenkins agent Dockerfile)

FIX: Amended path to my sampapp Dockerfile

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Amended path to my sampapp Dockerfile

Reference: Troubleshooting Log 4

ISSUE 12: BUILD FAILED 8 (Dockerfile: no such file or dir)

CAUSE: No github credentials in Jenkinsfile

FIX: Added github credentials in Jenkinsfile (github PAT)

Troubleshooting Steps Overview: Checked console output of build in Jenkins. Added github credentials in Jenkinsfile (github PAT)

Reference: Troubleshooting Log 4

Repushed Jenkinsfile to Github. Had to push twice more (because 10th try I missed out 1 } in the Jenkinsfile which caused it to fail)

Finally 11th time's the charm! It succeeded....

```
$ k get app -n argocd sampapp
NAME      SYNC STATUS  HEALTH STATUS
sampapp   Synced     Healthy
```

So what did this pipeline do exactly?

1. Checkout code
 - a. Jenkins checked out 2 repos (sampapp and k8s-manifests)
2. Build & Pushed Docker Images
 - a. Inside our custom docker-build agent (with Docker CI + buildx and mounted Docker socket):
 - i. Ran docker build -t mikangelica/sampapp:v1 -f Dockerfiles/sampapp/Dockerfile . >> Built a new Docker image of app
 - ii. Logged into Docker Hub using Jenkins-stored credentials (dockerhub-creds).
 - iii. Pushed the new image >> docker.io/mikangelica/sampapp:v1
3. Update Manifests
 - a. Edited manifests/base/deployment.yaml with sed to replace the old iamge with the new one (mikangelica/sampapp:v1)
 - b. Committed the change back to the k8s-manifests repo using Github PAT.
 - c. Pushed the commit > updated repo now contains the new image ref
 - d. GitOps source of truth (manifests repo) points to the latest image
4. Trigger ArgoCD Sync
 - a. Jenkins called the ArgoCD API with the bearer token to sync the sampapp app to match the latest manifest in the k8s-manifests repo.

One thing to note here:

In ArgoCD UI we can see that the last sync was 3 days ago even though we just successfully ran the

pipeline

The screenshot shows the ArgoCD interface for the 'sampapp' application. The left sidebar has options like 'Applications', 'Settings', 'User Info', and 'Documentation'. The main area shows 'APP HEALTH' as 'Healthy'. 'SYNC STATUS' shows 'Synced' to 'master (cbb108a)' with a green checkmark. 'LAST SYNC' shows 'Sync OK' to '2edba24' with a green checkmark, occurring 3 days ago. Below this, there's a 'PIPELINE' section with two stages: 'sampapp' and 'application', each with a green icon and '4 days' duration.

We also see that the pods in the cluster did not get redeployed/restarted:

| | | | | |
|---------|-------------------------|-----|---------|--------------|
| default | sampapp-f6c65db94-l72nz | 2/2 | Running | 6 (113m ago) |
| 7d4h | | | | |
| default | sampapp-f6c65db94-wwfg4 | 2/2 | Running | 6 (113m ago) |
| 7d4h | | | | |

Why?

Because even though we changed the deployment.yaml to image:oldtag and made it so that Jenkins wld update the image to v1, there weren't actually any real changes to the app source code.

And so, ArgoCD sees that there is nothing new to apply hence there was no redeployment.

And the actual pods were already labelled v1 actually (they were deployed with v1 before we changed to oldtag), so no manifest diff > no new commit > no Pod restart

K8s Deployments only roll Pods if the image tag changed fr (or if you forced a rollout).

So the pipeline did work but just that ArgoCD saw that there wasn't actually anything to update.

**Can consider trying to resync pipeline with new manifest changes and image tag

18 Ansible to provision a new node

Saturday, 13 September 2025 11:54 pm

Technology:

Ansible (Infrastructure Automation)

1. Set up Ansible VM

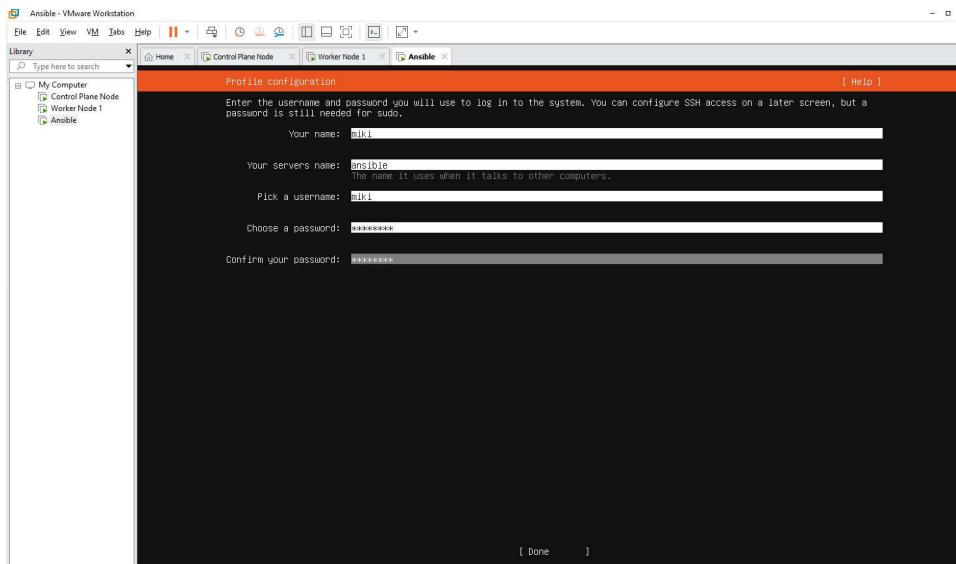
My computer is Windows and Ansible doesn't run natively on Windows. Either we use WSL or spin up a Linux VM for Ansible. Let's do the latter!

I gave 2CPU, 1GB RAM and 10GB Disk.

1. Click on "Create a New Virtual Machine"
2. Select Installer disc image file (iso)
 - Click on Browse and navigate to the directory I saved the ISO Image:
➤ C:\Users\mikia\Downloads\ubuntu-22.04.5-live-server-amd64.iso
3. Name VM (Ansible)
4. Select "Store virtual disk as a single file"
5. Customize Hardware > New CD/DVD (SATA) change to Auto detect
6. Click Finish and let OS boot up

2. System Bootup

- a) Click on first option in GRUB menu (Try or install Ubuntu Server)
- b) Keep clicking Done for all the options offered
- c) Ensure IP is on same L2 network as k3s nodes
- d) Name the server as your name, nodename (ansible), your username and pw
- e) Select Install OpenSSH server
- f) Wait for configurations to complete then click Reboot Now
- g) When screen shows "failed to mount /cdrom", press enter
- h) Login to server and shutdown now
- i) Click on VM settings and edit CD/DVD (SATA) to Use physical drive (autodetect)
- j) Power on VM again



It kept hanging at the username configuration page so I gave it 2GB RAM instead.

Now boots up.

3. Set up ssh for localhost to this ansible node

- Copy pasted ssh pubkey from localhost to ansible node
- Made sure file is one liner
- Changed ownership & permissions
- Remove # from AuthorizedKeysFile line in sshd_config
- Restarted ssh daemon
- Added ansible to ~/.ssh/config file so I can ssh ansible instead of IP

4. Set up ssh for ansible node to master1 & worker1 too

```
# ssh-keygen -t rsa -b 4096 -C "ansible-control" -N "" -f ~/ssh/id_rsa
```

Generating public/private rsa key pair.

Your identification has been saved in /home/miki/.ssh/id_rsa

Your public key has been saved in /home/miki/.ssh/id_rsa.pub

The key fingerprint is:

SHA256:uHw1xjRtAheXoUROOBd3SuqHannyfo0gAM7gX6k2pc0 ansible-control

The key's randomart image is:

```
+--[RSA 4096]---
```

```
|   .oBo=o. |
| .. o*.B.o |
| .+ . oB + |
| . o =. + = |
| . B..S B . |
| *E..=.o |
| ..o *... o |
|   o + o . |
|     .o. |
+---[SHA256]---
```

Ansible-control pubkey

```
ssh-rsa
```

```
AAAAB3NzaC1yc2EAAAQABAAQDkmEJEqziRIRpsbdbIOkC6AeU/mEKV8Lw9ewfM2kNKaf3o1jQR
```

```
wF4v2bjq94WY3XSOOrPgqv22DF1PKSxFWIUv44Gx1oOLYSoNEWNt+HmS1sVmHCiLm5HVO641/ntopNL  
WRPnqla10msrZgLWbjnW6  
+j9gAR1b23Bvl9JyRKz897rpN+reksQCeTZnsBbcjdJX35Ihlq71YL8Eo7Uifglb/3aeRSjk/Nh8/N/UpY1sJGJTn/  
76JJMRoGE6oAiA1dNtsVilZFao/5K4Hj7quYURqTrrkVsklt/a0rZLGdMQuYJtgLhZqBP7oBRiu63yoHQOa7qL  
S3ARqrXjWqn7FPgTFo3B0s2qSFPRiBzAzLibOH+XVj3SPsKLI+JYOR9auVsNOlxEOtT5KVfO+u2Ui4pb/lnCot  
7sXo9xCtK3rOuUev9cZbw4m1QgBqhkY+AD1cOOlt3TsyDxfB7qslkEVT+j+oqDcrTV1WsITySJdyggEgJQQc0  
TatfYdYHbveczq1NA0XMwRoPYK3Rom69Qn4VoKhqjgbXisZ3LeJmEe+HuwfR6Fh5dsLZ9LhrEgRd4qCD4py  
Tk8pF4vBHDplqMIgdg4p/LX5ZLc14PzSrYEIzCWkZ6AhY2hUo45wHCDWxYgm2vWWU3FD/rb89bRVBbA4  
ch5hKzf7UDLxajlpnFNZYgQ== ansible-control
```

Make one liner

```
sed -i 's/\r$//' /home/miki/.ssh/authorized_keys
```

Created the `~/.ssh/config` file to ssh via hostname

```
# Master node  
Host master1  
    HostName 192.168.81.128  
    User miki  
    IdentityFile ~/.ssh/id_rsa  
  
# Worker node 1  
Host worker1  
    HostName 192.168.81.129  
    User miki  
    IdentityFile ~/.ssh/id_rsa  
  
# Worker node 2  
Host worker2  
    HostName 192.168.81.134  
    User miki  
    IdentityFile ~/.ssh/id_rsa
```

5. Installing Ansible

Ansible only needs to run on **one machine**. It uses SSH to talk to other nodes.

For Ubuntu/Debian:

```
sudo apt update  
sudo apt install -y software-properties-common  
sudo add-apt-repository --yes --update ppa:ansible/ansible  
sudo apt install -y ansible
```

Check version:

```
root@ansible:~# ansible --version  
ansible [core 2.17.14]  
  config file = /etc/ansible/ansible.cfg  
  configured module search path = ['/root/.ansible/plugins/modules',  
 '/usr/share/ansible/plugins/modules']  
  ansible python module location = /usr/lib/python3/dist-packages/ansible  
  ansible collection location = /root/.ansible/collections:/usr/share/ansible/collections  
  executable location = /usr/bin/ansible  
  python version = 3.10.12 (main, Aug 15 2025, 14:32:43) [GCC 11.4.0] (/usr/bin/python3)  
  jinja version = 3.0.3
```

```
libyaml = True
```

6. Set up worker2 VM

- Set up SSH between worker2 and localhost
- Set up SSH between worker2 and Ansible

7. Set up Ansible Inventory in the following dir (common ansible dir):

```
#mkdir ~/ansible-k8s/
```

Place the following files in the above dir:

- inventory.ini → inventory file with your nodes
- playbook.yml → your playbooks
- group_vars/, roles/, etc. (later if you structure more complex projects)

In **inventory.ini** :

```
[k8s_masters]
master1 ansible_host=192.168.81.128 ansible_user=miki

[k8s_workers]
worker1 ansible_host=192.168.81.129 ansible_user=miki
worker2 ansible_host=192.168.81.134 ansible_user=miki
```

Explanation:

- [k8s_masters] = group name.
- master1 = alias for your node (can be anything meaningful).
- ansible_host=192.168.81.128 = the actual IP address of the node.
- ansible_user=ubuntu = the SSH user you log in with.

8. Check connectivity:

```
ansible all -i inventory.ini -m ping
```

You should see pong back from each host.

```
miki@ansible:~/ansible-k8s$ ansible all -i inventory.ini -m ping
[WARNING]: Platform linux on host worker2 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
```

```
worker2 | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python3.10"
    },
    "changed": false,
    "ping": "pong"
}
```

```
[WARNING]: Platform linux on host master1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
```

```
master1 | SUCCESS => {
    "ansible_facts": {
```

```

    "discovered_interpreter_python": "/usr/bin/python3.10"
},
"changed": false,
"ping": "pong"
}
[WARNING]: Platform linux on host worker1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
worker1 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python3.10"
},
"changed": false,
"ping": "pong"
}

```

In the above, Ansible read the inventory.ini file and saw master1 & worker1. Then it used SSH to connect to each host, discovered the python interpreter on each node, Ran the ping module which tests connection & python availability. Both returned pong means Ansible can successfully connect, authenticate, and execute commands on both nodes! :D

Test something slightly more advanced:

```

miki@ansible:~/ansible-k8s$ ansible all -i inventory.ini -m shell -a "hostname && uptime"
[WARNING]: Platform linux on host worker1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
worker1 | CHANGED | rc=0 >>
worker1
17:53:34 up 1:28, 1 user, load average: 0.60, 0.28, 0.20
[WARNING]: Platform linux on host worker2 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
worker2 | CHANGED | rc=0 >>
worker2
17:53:34 up 4 min, 1 user, load average: 0.02, 0.07, 0.03
[WARNING]: Platform linux on host master1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
master1 | CHANGED | rc=0 >>
master1
17:53:34 up 1:29, 1 user, load average: 0.06, 0.04, 0.08

```

9. Create folder structure

```

ansible-k8s/
|   |-- inventory/
|   |   |-- inventory.ini

```

```
└── playbooks/
    └── add-worker.yml
└── roles/
    ├── common/
    │   └── tasks/main.yml
    └── worker/
        └── tasks/main.yml
```

10. Create the main playbook

```
---
- name: Configure and join new Kubernetes worker node
  hosts: k8s_workers
  become: true
  vars:
    k3s_token:
      "K1080e8d3f7b27ad61ef4ca4d60c0f53705c52f3ae3c17a4d782cedcdd865c1430::server:60f6d6681f7e
      54098d82c7f01d32f64b"
    master_ip: "192.168.81.128"
  roles:
    - common
    - worker
```

Find node token on master1:

```
sudo cat /var/lib/rancher/k3s/server/node-token
```

Create common role task main playbook (~/ansible-k8s/roles/common/tasks/main.yml)

```
---
- name: Update apt packages
  apt:
    update_cache: yes
    upgrade: yes

- name: Install dependencies
  apt:
    name:
      - curl
      - apt-transport-https
      - ca-certificates
      - software-properties-common
    state: present
```

Create worker role task main playbook (~/ansible-k8s/roles/worker/tasks/main.yml)

```
---
- name: Install K3s agent (worker)
  shell: |
    curl -sfL https://get.k3s.io | K3S_URL=https://{{ master_ip }}:6443 K3S_TOKEN={{ k3s_token }} sh -
  args:
    creates: /etc/systemd/system/k3s-agent.service
  register: k3s_agent_install
```

```
- name: Show K3s installation output
  debug:
    var: k3s_agent_install.stdout
```

11. Run the playbook

To run playbook

```
cd ~/ansible-k8s
ansible-playbook -i inventory/inventory.ini playbooks/add-worker.yaml
```

Ran into issues

ISSUE 2: Playbook not running because role not found

CAUSE: roles dir in wrong path

FIX: Moved roles dir to under playbooks dir

Troubleshooting Steps Overview: Saw error upon trying to run playbook to add worker2 into cluster is related to role 'common' not being found, moved roles dir to path and looks like playbook can detect it now

Reference: Troubleshooting Log 1

ISSUE 3: Playbook not running because passwordless sudo not configured for user

CAUSE: ansible user does not have passwordless sudo

FIX: Added passwordless sudo for ansible user

Troubleshooting Steps Overview: Error during playbook run states Missing sudo password, added user to visudo file for passwordless sudo

Reference: Troubleshooting Log 1

Run playbook:

```
miki@ansible:~/ansible-k8s$ ansible-playbook -i inventory/inventory.ini playbooks/add-worker.yaml
```

PLAY [Configure and join new Kubernetes worker node]

```
*****
*****
```

TASK [Gathering Facts]

```
*****
*****
```

[WARNING]: Platform linux on host worker2 is using the discovered Python interpreter at /usr/bin/python3.10, but future installation of another Python interpreter could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference_appendices/interpreter_discovery.html for more information.

ok: [worker2]

[WARNING]: Platform linux on host worker1 is using the discovered Python interpreter at /usr/bin/python3.10, but future installation of another Python interpreter could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference_appendices/interpreter_discovery.html for more information.

ok: [worker1]

TASK [common : Update apt packages]

```
*****
*****
```

changed: [worker2]

```
changed: [worker1]
```

```
TASK [common : Install dependencies]
```

```
changed: [worker2]  
changed: [worker1]
```

```
TASK [worker : Install K3s agent (worker)]
```

```
ok: [worker1]  
changed: [worker2]
```

```
TASK [worker : Show K3s installation output]
```

```
ok: [worker1] => {  
    "k3s_agent_install.stdout": "skipped, since /etc/systemd/system/k3s-agent.service exists"  
}  
ok: [worker2] => {  
    "k3s_agent_install.stdout": "[INFO] Finding release for channel stable\n[INFO] Using v1.33.5+k3s1  
as release\n[INFO] Downloading hash https://github.com/k3s-io/k3s/releases/download/v1.33.5+k3s1/sha256sum-amd64.txt\n[INFO] Downloading binary https://github.com/k3s-io/k3s/releases/download/v1.33.5+k3s1/k3s\n[INFO] Verifying binary download\n[INFO] Installing k3s to /usr/local/bin/k3s\n[INFO] Skipping installation of SELinux RPM\n[INFO] Creating /usr/local/bin/kubectl symlink to k3s\n[INFO] Creating /usr/local/bin/crictl symlink to k3s\n[INFO] Creating /usr/local/bin/ctr symlink to k3s\n[INFO] Creating killall script /usr/local/bin/k3s-killall.sh\n[INFO] Creating uninstall script /usr/local/bin/k3s-agent-uninstall.sh\n[INFO] env: Creating environment file /etc/systemd/system/k3s-agent.service.env\n[INFO] systemd: Creating service file /etc/systemd/system/k3s-agent.service\n[INFO] systemd: Enabling k3s-agent unit\n[INFO] systemd: Starting k3s-agent"  
}
```

```
PLAY RECAP
```

```
*****  
*****  
worker1      : ok=5  changed=2  unreachable=0  failed=0  skipped=0  rescued=0  ignored=0  
worker2      : ok=5  changed=3  unreachable=0  failed=0  skipped=0  rescued=0  ignored=0
```

Verify:

```
$ kubectl get no -owide  
NAME      STATUS ROLES      AGE VERSION      INTERNAL-IP      EXTERNAL-IP      OS-IMAGE  
KERNEL-VERSION      CONTAINER-RUNTIME  
master1   Ready   control-plane,master   26d   v1.33.4+k3s1   192.168.81.128   <none>      Ubuntu  
22.04.5 LTS 5.15.0-156-generic containerd://2.0.5-k3s2  
worker1   Ready   <none>      26d   v1.33.4+k3s1   192.168.81.129   <none>      Ubuntu 22.04.5  
LTS 5.15.0-156-generic containerd://2.0.5-k3s2  
worker2   Ready   <none>      84s   v1.33.5+k3s1   192.168.81.134   <none>      Ubuntu 22.04.5  
LTS 5.15.0-157-generic containerd://2.1.4-k3s1
```

SUCCESS!

19 Istio ServiceMesh

Saturday, 13 September 2025 11:55 pm

Technology:

Istio ServiceMesh

1. Confirm cluster access

```
root@master1:~# kubectl get no
NAME    STATUS  ROLES      AGE     VERSION
master1 Ready   control-plane,master  26d    v1.33.4+k3s1
worker1 Ready   <none>       26d    v1.33.4+k3s1
worker2 Ready   <none>       8m16s  v1.33.5+k3s1
```

2. Install istioctl locally

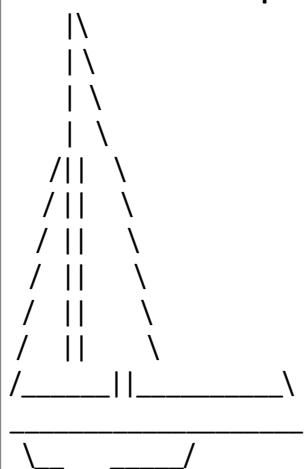
- Localhost already has kubeconfig with full admin access to cluster
 - Can version-control and rerun istio installs easily from there
 - Avoid messing with control plane environment

```
# mkdir -p ~/istio && cd ~/istio
# curl -kLO https://github.com/istio/istio/releases/download/1.27.1/istio-1.27.1-win.zip
% Total  % Received % Xferd Average Speed Time Time Time Current
          Dload Upload Total Spent Left Speed
0 0 0 0 0 0 0 0 --:--:--:--:--:--:--:-- 0
100 29.1M 100 29.1M 0 0 4861k 0 0:00:06 0:00:06 --:-- 5860k
# unzip istio-1.27.1-win.zip -d istio-1.27.1
# export PATH=$HOME/istio/istio-1.27.1/istio-1.27.1/bin:$PATH
# istioctl version
Istio is not present in the cluster: no running Istio pods in namespace "istio-system"
client version: 1.27.1
```

- Add istioctl path to environment variables (system properties) so it is persistent

3. Install the Istio Control Plane (Default Profile)

```
# istioctl install --set profile=demo -v
```



- ✓ Istio core installed
- ✓ Istiod installed
- ✓ Ingress gateways installed
- ✓ Egress gateways installed
- ✓ Installation complete

demo = a full-featured profile (Ingress, Egress, Prometheus, etc.)
Perfect for learning and testing.

It'll create:

- Namespace istio-system
- Control plane pods (istiod, istio-ingressgateway, etc.)
- Webhooks, CRDs, etc.

Verify pods are there now:

```
# k get po -n istio-system
NAME           READY STATUS RESTARTS AGE
istio-egressgateway-74bbcd4975d-5xftt 1/1   Running 0      80s
istio-ingressgateway-7c5fd9bc88-zq95g 1/1   Running 0      80s
istiod-597958d8ff-x64wr    1/1   Running 0      116s
```

4. Label application namespace

To enable automatic sidecar injection (Envoy proxy)

Check which ns the app is deployed in

```
# kubectl get po -A | grep sampapp
default      sampapp-f6c65db94-l72nz          2/2   Running   14
(3h5m ago)  13d
default      sampapp-f6c65db94-wwfg4         2/2   Running   14
(3h5m ago)  13d
```

```
# kubectl label namespace default istio-injection=enabled
namespace/default labeled
```

5. Restart app deployment

```
# k rollout restart deployment sampapp -n default
deployment.apps/sampapp restarted
```

Wait a couple mins and you'll see that sampapp has 3 containers now instead of 2

| | | | | |
|-------------------------|-----|---------|---|------|
| sampapp-cfdc49c4b-bddhl | 3/3 | Running | 0 | 113s |
| sampapp-cfdc49c4b-f9knr | 3/3 | Running | 0 | 3m |

```
$ k describe po sampapp-cfdc49c4b-bddhl -n default | grep -i "image:"
Image: docker.io/istio/proxyv2:1.27.1
Image: docker.io/istio/proxyv2:1.27.1
Image: mikiangelica/sampapp:v1
Image: nicolaka/netshoot
```

6. Verify that app is part of the mesh

```
# istioctl proxy-status
NAME                CLUSTER   ISTIOD      VERSION
SUBSCRIBED TYPES
istio-egressgateway-74bbd4975d-5xftt.istio-system  Kubernetes
istiod-597958d8ff-x64wr  1.27.1  3 (CDS,LDS,EDS)
istio-ingressgateway-7c5fd9bc88-zq95g.istio-system  Kubernetes
istiod-597958d8ff-x64wr  1.27.1  3 (CDS,LDS,EDS)
sampapp-cfdc49c4b-bddhl.default          Kubernetes  istiod-597958d8ff-
x64wr  1.27.1  4 (CDS,LDS,EDS,RDS)
sampapp-cfdc49c4b-f9knr.default          Kubernetes  istiod-597958d8ff-
x64wr  1.27.1  4 (CDS,LDS,EDS,RDS)
```

Awesome it worked.

7. Check app service

```
# k get svc -n default sampapp-service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
sampapp-service  ClusterIP  10.43.133.187  <none>      5000/TCP  26d
```

8. Create Istio Gateway manifest

In manifests/istio/sampapp-gateway.yaml

```
apiVersion: gateway.networking.k8s.io/v1
kind: Gateway
metadata:
  name: sampapp-gateway
  namespace: default
spec:
  gatewayClassName: istio
  listeners:
    - name: http
      protocol: HTTP
      port: 80
      allowedRoutes:
        namespaces:
          from: Same
```

```
# k apply -f sampapp-gateway.yaml
gateway.networking.istio.io/sampapp-gateway created
```

9. Create VirtualService

In manifests/istio/sampapp-virtualservice.yaml

```
apiVersion: networking.istio.io/v1beta1
kind: VirtualService
metadata:
  name: sampapp
  namespace: default
spec:
  hosts:
    - "*"
```

```

gateways:
- sampapp-gateway
http:
- match:
- uri:
  prefix: /
route:
- destination:
  host: sampapp-service # match svc
  port:
    number: 5000 # match svc

```

```
# k apply -f sampapp-virtualservice.yaml
virtualservice.networking.istio.io/sampapp created
```

10. Check that both are up

```
# k get gateway,vs
NAME          CLASS ADDRESS PROGRAMMED AGE
gateway.gateway.networking.k8s.io/sampapp-gateway istio      False   66s

NAME          GATEWAYS HOSTS AGE
virtualservice.networking.istio.io/sampapp ["sampapp-gateway"] ["*"] 24m
```

Note that for gateway, PROGRAMMED is False and there is no Address as it is not yet bound to an istio ingress gateway service.

11. Link to istio ingress gateway svc

Verify istio ingress gateway class

```
# k get gatewayclass
NAME      CONTROLLER      ACCEPTED AGE
istio     istio.io/gateway-controller True   51m
istio-remote istio.io/unmanaged-gateway True   51m
```

Check that istio ingress gateway svc exists

```
# k get svc -n istio-system
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP PORT(S)
AGE
istio-egressgateway ClusterIP  10.43.143.163 <none>    80/TCP,443/TCP
52m
istio-ingressgateway LoadBalancer 10.43.71.123 <pending>
15021:32477/TCP,80:31171/TCP,443:31707/TCP,31400:30106/TCP,15443:30688/TC
P 52m
istiod     ClusterIP  10.43.73.122 <none>
15010/TCP,15012/TCP,443/TCP,15014/TCP           52m
```

Patch istio ingress gateway svc to use NodePort instead of LoadBalancer

```
# kubectl patch svc istio-ingressgateway -n istio-system -p '{"spec": {"type": "NodePort"}}'
```

Check ingressgateway again and take note of the "80:**31171**/TCP"

```
# kubectl get svc -n istio-system istio-ingressgateway
NAME          TYPE      CLUSTER-IP   EXTERNAL-IP PORT(S)
AGE
istio-ingressgateway  NodePort  10.43.71.123 <none>
15021:32477/TCP,80:31171/TCP,443:31707/TCP,31400:30106/TCP,15443:30688/TC
P 53m
```

12. Wait for gateway to become programmed

ISSUE 4: GATEWAY NOT BECOMING PROGRAMMED "TRUE"

CAUSE: Wrong service type and VS are no longer attached to in 1.27+

FIX: Changed from loadbalancer to Nodeport, and VS to HTTPRoute

Troubleshooting Steps Overview: Checked yaml of gateway and saw error, changed from VS to HTTPRoute. Checked gateway svc type, tried to patch (failed). Deleted gateway (which owns svc) and applied new gateway yaml that states NodePort for svc

Reference: Troubleshooting Log 1

Programmed = True

```
$ k get svc sampapp-gateway-istio
NAME          TYPE      CLUSTER-IP   EXTERNAL-IP PORT(S)           AGE
sampapp-gateway-istio  NodePort  10.43.252.215 <none>
15021:30857/TCP,80:31747/TCP  22s

$ k get gateway
NAME      CLASS ADDRESS           PROGRAMMED AGE
sampapp-gateway  istio  sampapp-gateway-istio.default.svc.cluster.local  True    41s
```

13. Test access to app

On browser, go to:

<http://192.168.81.128:31747>

Should be able to access the app!



Feedback Form

| |
|---------------|
| Your Name |
| Your Email |
| Your Feedback |
| Submit |

Can use any of the nodes in clusters' IPs to access.

20 Prometheus + Grafana

Saturday, 13 September 2025 11:55 pm

Technology:

Prometheus

Grafana

Kiali

1. Create namespace for monitoring

```
# kubectl create namespace monitoring
```

2. Install kube-prometheus-stack via Helm

(includes Prometheus, Grafana, Alertmanager) - most popular setup

```
# helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
```

```
# helm repo update
```

```
# helm install prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring \
--set grafana.enabled=true \
--set prometheus.prometheusSpec.serviceMonitorSelectorNilUsesHelmValues=false
```

NAME: prometheus

LAST DEPLOYED: Sun Oct 5 10:02:58 2025

NAMESPACE: monitoring

STATUS: deployed

REVISION: 1

NOTES:

kube-prometheus-stack has been installed. Check its status by running:

```
kubectl --namespace monitoring get pods -l "release=prometheus"
```

Get Grafana 'admin' user password by running:

```
kubectl --namespace monitoring get secrets prometheus-grafana -o
jsonpath="{.data.admin-password}" | base64 -d ; echo
```

Access Grafana local instance:

```
export POD_NAME=$(kubectl --namespace monitoring get pod -l
"app.kubernetes.io/name=grafana,app.kubernetes.io/instance=prometheus" -oname)
kubectl --namespace monitoring port-forward $POD_NAME 3000
```

Visit <https://github.com/prometheus-operator/kube-prometheus> for instructions on how to create & configure Alertmanager and Prometheus instances using the Operator.

Verify pods are up:

| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE |
|--|-----------|---------|----------|----------------|-----------------------|-----------------------|
| NOMINATED NODE | READINESS | GATES | | | | |
| alertmanager-prometheus-kube-prometheus-alertmanager-0 | 2/2 | Running | 0 | 109s | 172.18.189.84 | worker2 <none> <none> |
| prometheus-grafana-674cf8cb44-hgjmf | 3/3 | Running | 0 | 172.18.189.83 | worker2 <none> <none> | 2m12s |
| prometheus-kube-prometheus-operator-77c7674c65-8776m | 1/1 | Running | 0 | 2m12s | 172.18.189.82 | worker2 <none> <none> |
| prometheus-kube-state-metrics-7c5fb9d798-g9cn4 | 1/1 | Running | 0 | 172.18.189.81 | worker2 <none> <none> | 2m12s |
| prometheus-prometheus-kube-prometheus-prometheus-0 | 2/2 | Running | 0 | 108s | 172.18.189.85 | worker2 <none> <none> |
| prometheus-prometheus-node-exporter-64vtm | 1/1 | Running | 0 | 192.168.81.134 | worker2 <none> <none> | 2m12s |
| prometheus-prometheus-node-exporter-7vbl5 | 1/1 | Running | 0 | 192.168.81.129 | worker1 <none> <none> | 2m12s |
| prometheus-prometheus-node-exporter-qbnpb | 1/1 | Running | 0 | 192.168.81.128 | master1 <none> <none> | 2m12s |

3. Enable istio metrics scraping

Istio exposes metrics automatically via istiod, istio-proxy etc.. Just need to enable Prometheus scraping annotations

```
# k label ns istio-system istio-injection=enabled --overwrite
namespace/istio-system labeled

# kubectl apply -f
https://raw.githubusercontent.com/istio/istio/release-1.27/samples/addons/prometheus.yaml

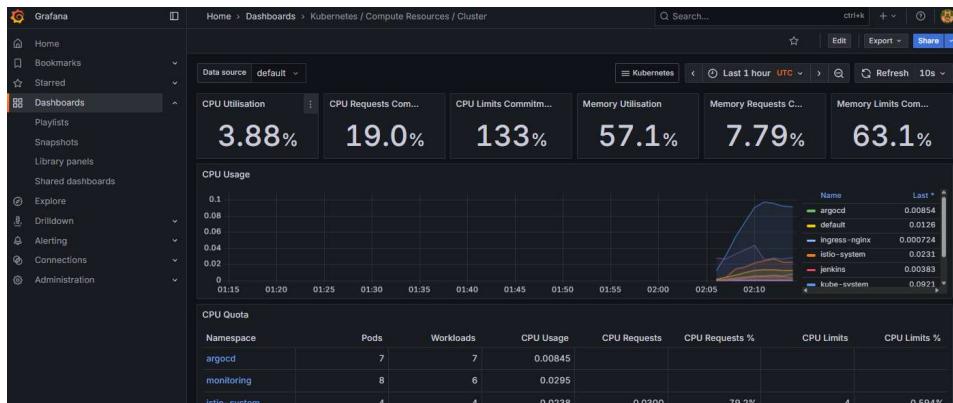
serviceaccount/prometheus created
configmap/prometheus created
clusterrole.rbac.authorization.k8s.io/prometheus created
clusterrolebinding.rbac.authorization.k8s.io/prometheus created
service/prometheus created
deployment.apps/prometheus created
```

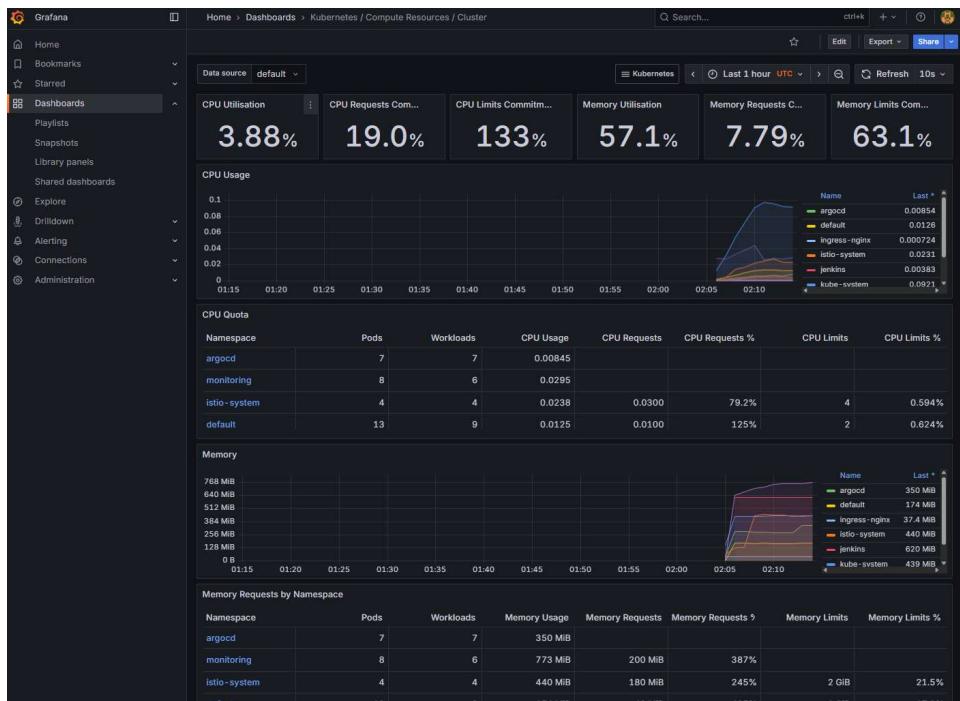
6. Access Grafana Dashboard

Port-forward grafana

```
kubectl -n monitoring port-forward svc/prometheus-grafana 3000:80 &
```

Go to : <http://localhost:3000>





7. Install Kiali

Kiali is an observability console for Istio servicemesh used to visualise the service mesh topology, monitor traffic, analyse metrics, view distributed traces and validate Istio configs.

Helps us understand app health, identify bottlenecks and troubleshoot issues within the complex architecture of a microservices environment.

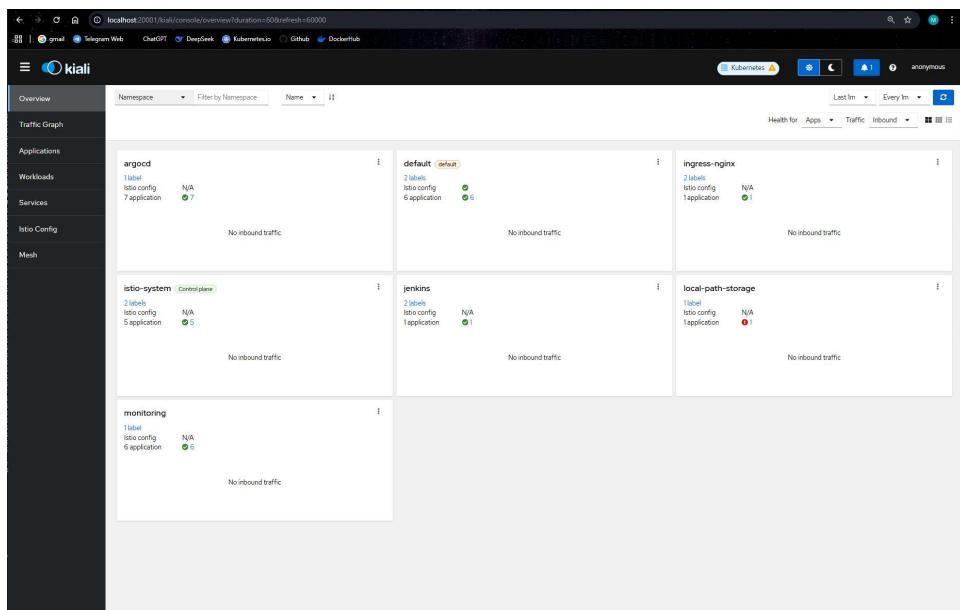
```
# kubectl apply -f
https://raw.githubusercontent.com/istio/istio/release-1.27/samples/addons/kiali.yaml
serviceaccount/kiali created
configmap/kiali created
clusterrole.rbac.authorization.k8s.io/kiali created
clusterrolebinding.rbac.authorization.k8s.io/kiali created
service/kiali created
deployment.apps/kiali created

# k get all -A | grep kiali
istio-system      pod/kiali-56f54f58f9-5xsgj           1/1   Running    0
103s
istio-system      service/kiali                         ClusterIP  10.43.72.137 <none>
20001/TCP,9090/TCP                                     103s
istio-system      deployment.apps/kiali                  1/1   1       1       103s
istio-system      replicaset.apps/kiali-56f54f58f9        1       1       1       103s
```

Portforward

```
kubectl port-forward svc/kiali 20001:20001 -n istio-system &
```

On browser, go to <http://localhost:20001>



Troubleshooting Log 1

Sunday, 28 September 2025 11:04 pm

SECTION: GITHUB RESYNC

ISSUE 1: ERROR BUILDING IMAGE

```
# docker buildx build -f sampapp/Dockerfiles/sampapp/Dockerfile --load -t mikiangelica/sampapp:v2 .
[+] Building 6.8s (6/9)
docker:desktop-linux
=> [internal] load build definition from Dockerfile
0.0s
=> => transferring dockerfile: 294B
0.0s
=> [internal] load metadata for docker.io/library/python:3.10-slim
0.9s
=> [internal] load .dockerignore
0.0s
=> => transferring context: 2B
0.0s
=> [1/5] FROM docker.io/library/python:3.10-
slim@sha256:f8081b61393cc16b2339d377bb523a5646cdd28fc6105612b3893c405d27f730
5.8s
=> => resolve docker.io/library/python:3.10-
slim@sha256:f8081b61393cc16b2339d377bb523a5646cdd28fc6105612b3893c405d27f730
0.0s
=> => sha256:fceec8b8a90ef53ff8f6211082b139436e3b3f0f16adf7752fe0d9e6681a3231 250B / 250B
0.6s
=> => sha256:6dc114d5e12c107652ba61ef068cfacc611596390cbafa6cd21206a670e4e264 14.09MB /
14.09MB
=> => sha256:d02c566428bb557aab93bf3d9d768b9abb145010f1f8a6756db930f0902f60f7 1.29MB /
1.29MB
=> => sha256:ce1261c6d567efa8e3b457673eeeb474a0a8066df6bb95ca9a6a94a31e219dd3 15.73MB /
29.77MB
=> [internal] load build context
1.1s
=> => transferring context: 45.81MB
1.0s
=> CANCELED [2/5] WORKDIR /app
0.0s
ERROR: failed to build: failed to solve: Canceled: context canceled
```

^^ Above means build context was too big so it was cancelled.

So the fix was to only include the sampapp dir (context build)

```
# docker buildx build \
-f sampapp/Dockerfiles/sampapp/Dockerfile \
-t mikiangelica/sampapp:v2 \
--load \
sampapp/
```

SECTION: ANSIBLE TO PROVISION A NEW NODE

ISSUE 2: Playbook not running because role not found

Got this error

```
miki@ansible:~/ansible-k8s$ ansible-playbook -i inventory/inventory.ini playbooks/add-worker.yaml
ERROR! the role 'common' was not found in /home/miki/ansible-
k8s/playbooks/roles:/home/miki/.ansible/roles:/usr/share/ansible/roles:/etc/ansible/roles:/home/miki
/ansible-k8s/playbooks
```

The error appears to be in '/home/miki/ansible-k8s/playbooks/add-worker.yaml': line 9, column 7, but may be elsewhere in the file depending on the exact syntax problem.

The offending line appears to be:

roles:

- common
- ^ here

Moved the roles/ dir to under playbooks/ dir and the error msg changed, meaning now it can detect

SECTION: ANSIBLE TO PROVISION A NEW NODE

ISSUE 3: Playbook not running because passwordless sudo not configured for user

Got this error

```
miki@ansible:~/ansible-k8s$ ansible-playbook -i inventory/inventory.ini playbooks/add-worker.yaml

PLAY [Configure and join new Kubernetes worker node]
*****
*****
*****
```

TASK [Gathering Facts]

```
*****
*****
```

fatal: [worker2]: FAILED! => {"msg": "Missing sudo password"}

fatal: [worker1]: FAILED! => {"msg": "Missing sudo password"}

PLAY RECAP

```
*****
*****
```

worker1 : ok=0 changed=0 unreachable=0 failed=1 skipped=0 rescued=0 ignored=0

worker2 : ok=0 changed=0 unreachable=0 failed=1 skipped=0 rescued=0 ignored=0

Went into each host (master and both workers) to add passwordless sudo for user miki

```
# sudo visudo
# add "miki ALL=(ALL) NOPASSWD:ALL" to end of file
Then save and exit (Ctrl+O, Enter, Ctrl+X in nano).
```

Test sudo command

```
ansible k8s_workers -i inventory/inventory.ini -m shell -a "sudo whoami"
```

The above command hung

```
# ssh worker1 (from ansible)

# sudo -n whoami
sudo: a password is required

# sudo -l -U miki
Matching Defaults entries for miki on worker1:
  env_reset, mail_badpass, secure_path=/usr/local/sbin\:/usr/local/bin\:/usr/sbin\:/usr/bin\:/sbin
\:/bin\:/snap/bin, use_pty

User miki may run the following commands on worker1:
  (ALL) NOPASSWD: ALL
  (ALL : ALL) ALL
```

The sudo -l output proves miki *is* allowed NOPASSWD, but sudo -n whoami still requiring a password means something else on the node is overriding or taking precedence.

Added an explicit sudoers file

This will create a small file that explicitly grants NOPASSWD with a predictable filename so it sorts last

```
# sudo visudo -f /etc/sudoers.d/99-ansible-miki
# Add "miki ALL=(ALL) NOPASSWD:ALL"
```

Test locally on node

```
sudo -n whoami && echo OK | echo FAIL:$?
root
OK
```

Check SSH and sudo from ansible node

```
# ansible all -i inventory/inventory.ini -m command -a "sudo -n whoami"
[WARNING]: Platform linux on host master1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
master1 | CHANGED | rc=0 >>
root
[WARNING]: Platform linux on host worker2 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
worker2 | CHANGED | rc=0 >>
root
[WARNING]: Platform linux on host worker1 is using the discovered Python interpreter at
/usr/bin/python3.10, but future installation of another Python interpreter
could change the meaning of that path. See https://docs.ansible.com/ansible-core/2.17/reference\_appendices/interpreter\_discovery.html for more information.
worker1 | CHANGED | rc=0 >>
root
```

Now run again

```
# ansible-playbook -i inventory/inventory.ini playbooks/add-worker.yaml
```

SECTION: ISTIO SERVICEMESH

ISSUE 4: GATEWAY NOT BECOMING PROGRAMMED "TRUE"

Below error

```
$ kubectl get gateway sampapp-gateway -n default -o yaml | grep -A10 status:  
status:  
conditions:  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: Resource accepted  
  observedGeneration: 1  
  reason: Accepted  
  status: "True"  
  type: Accepted  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: 'Assigned to service(s) sampapp-gateway-istio.default.svc.cluster.local:80,  
  but failed to assign to all requested addresses: address pending for hostname  
  "sampapp-gateway-istio.default.svc.cluster.local"'  
  observedGeneration: 1  
  reason: AddressNotAssigned  
  status: "False"  
  type: Programmed  
listeners:  
- attachedRoutes: 0  
  conditions:  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: No errors found  
  observedGeneration: 1  
  reason: Accepted  
  status: "True"  
  type: Accepted  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: No errors found  
  observedGeneration: 1  
  reason: NoConflicts  
  status: "False"  
  type: Conflicted  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: No errors found  
  observedGeneration: 1  
  reason: Programmed  
  status: "True"  
  type: Programmed  
- lastTransitionTime: "2025-10-04T20:13:07Z"  
  message: No errors found  
  observedGeneration: 1  
  reason: ResolvedRefs  
  status: "True"  
  type: ResolvedRefs  
name: http  
supportedKinds:  
- group: gateway.networking.k8s.io  
  kind: HTTPRoute  
- group: gateway.networking.k8s.io
```

```
kind: GRPCRoute
```

In Istio **1.27+**, the new Gateway API **does not attach VirtualServices anymore**.

It only attaches **HTTPRoute** or **GRPCRoute** resources.

Created HTTPRoute instead and deleted the VS

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: sampapp
  namespace: default
spec:
  parentRefs:
  - name: sampapp-gateway
  rules:
  - matches:
    - path:
      type: PathPrefix
      value: /
  backendRefs:
  - name: sampapp-service # must match your Service name
  port: 5000 # must match your Service port
```

⌚ What's Happening

GatewayClass istio → Accepted = True

Istiod logs show it's actively reconciling your gateway

Gateway status = Programmed=False (AddressNotAssigned)

 This means Istio **did** pick up your Gateway, but it cannot bind it to any Service (address).

💡 Why That Happens

Since Istio 1.27+ uses the **new Kubernetes Gateway API**, your Gateway spec:

```
gatewayClassName: istio
```

tells Istiod to automatically **spawn a service** named:

sampapp-gateway-istio

in the **same namespace** (default).

That service is responsible for binding a NodePort or LoadBalancer IP —
but in your logs:

failed to assign to all requested addresses: address pending for hostname "sampapp-gateway-istio.default.svc.cluster.local"

...it means that service never got created properly.

Checking gateway-istio service is LoadBalancer

```
# k get deploy,svc -n default
```

| NAME | READY | UP-TO-DATE | AVAILABLE | AGE |
|--------------------------|-------|------------|-----------|-----|
| deployment.apps/nginx | 1/1 | 1 | 1 | 26d |
| deployment.apps/postgres | 1/1 | 1 | 1 | 26d |
| deployment.apps/redis | 1/1 | 1 | 1 | 26d |

```
deployment.apps/sampapp 2/2 2 2 26d
deployment.apps/sampapp-gateway-istio 1/1 1 1 4h58m
```

| NAME | TYPE | CLUSTER-IP | EXTERNAL-IP | PORT(S) | AGE |
|-------------------------------|--------------|------------------------------|-------------|----------|-----|
| service/kubernetes | ClusterIP | 10.43.0.1 | <none> | 443/TCP | 26d |
| service/nginx-service | ClusterIP | 10.43.104.188 | <none> | 80/TCP | 26d |
| service/postgres-service | ClusterIP | 10.43.178.155 | <none> | 5432/TCP | 26d |
| service/redis-service | ClusterIP | 10.43.146.71 | <none> | 6379/TCP | 26d |
| service/sampapp-gateway-istio | LoadBalancer | 10.43.34.51 | <pending> | | |
| | | 15021:32516/TCP,80:32683/TCP | 4h58m | | |
| service/sampapp-service | ClusterIP | 10.43.133.187 | <none> | 5000/TCP | 26d |

Patch to NodePort

```
# kubectl patch svc sampapp-gateway-istio -n default \
-p '{"spec": {"type": "NodePort"}}'
```

Still showed up as LoadBalancer

```
# k get svc sampapp-gateway-istio
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
sampapp-gateway-istio  LoadBalancer  10.43.34.51  <pending>  15021:32516/TCP,80:32683/TCP
5h1m
```

Istio's **Gateway controller** is the *owner* of that Service (ownerReferences → gateway.networking.k8s.io/v1beta1), so any manual edit (like changing LoadBalancer → NodePort) is automatically reverted.

kubectl delete gateway sampapp-gateway -n default

New GW yaml

```
apiVersion: gateway.networking.k8s.io/v1
kind: Gateway
metadata:
  name: sampapp-gateway
  namespace: default
  annotations:
    networking.istio.io/service-type: "NodePort" # ⚡ tells Istio to use NodePort
spec:
  gatewayClassName: istio
  listeners:
  - name: http
    port: 80
    protocol: HTTP
    allowedRoutes:
      namespaces:
        from: Same
```

Check

```
$ k get svc sampapp-gateway-istio
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
sampapp-gateway-istio  NodePort  10.43.252.215  <none>  15021:30857/TCP,80:31747/TCP  22s

$ k get gateway
```

| NAME | CLASS | ADDRESS | PROGRAMMED AGE |
|-----------------|-------|---|----------------|
| sampapp-gateway | istio | sampapp-gateway-istio.default.svc.cluster.local | True 41s |

Troubleshooting Log 2

Sunday, 21 September 2025 3:43 pm

*****ALL BELOW ISSUES ARE FROM PVC FOR POSTGRES SECTION.**

4 ISSUES

ISSUE 1: Pod, PVC and PV not coming up.

Ran into an ISSUE where after applying, no pv was made. Pod was not coming up and the pvc was pending state with error saying ProvisioningFailed.

1. Checked the provisioner pod status:

```
# kubectl get po -n local-path-storage  
No resources found in local-path-storage namespace.
```

!!! Provisioner pod isn't even running -- that's why PVCs are stuck in Pending and failing to provision.
This pod is important as it is the one actually creating the local PVs.
Without that pod, no one is acting on the PVC requests.

1. Redeploy the local-path-provisioner into cluster: (below steps is for rancher/k3s)
Normally it is pre-installed but if missing, can reapply manifest

```
# kubectl apply -f https://raw.githubusercontent.com/rancher/local-path-provisioner/master/deploy/local-path-storage.yaml
```

```
namespace/local-path-storage created  
serviceaccount/local-path-provisioner-service-account created  
role.rbac.authorization.k8s.io/local-path-provisioner-role created  
Warning: resource clusterroles/local-path-provisioner-role is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.  
clusterrole.rbac.authorization.k8s.io/local-path-provisioner-role configured  
rolebinding.rbac.authorization.k8s.io/local-path-provisioner-bind created  
Warning: resource clusterrolebindings/local-path-provisioner-bind is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.  
clusterrolebinding.rbac.authorization.k8s.io/local-path-provisioner-bind configured  
deployment.apps/local-path-provisioner created  
Warning: resource storageclasses/local-path is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.  
storageclass.storage.k8s.io/local-path configured  
configmap/local-path-config created
```

^^Above manifest will:

- Create local-path-storage namespace
- Deploy the local-path-provisioner pod
- Bind it to your SC (local-path)

ISSUE 2: Pod still not coming up with 'calico' failure error (ContainerCreating state) - Calico Pods Issue
 Next ISSUE, after applying above manifest, pod is just in ContainerCreating state.

Events:

| Type | Reason | Age | From | Message |
|---------|------------------------|-----------------------|-------------------|---|
| Normal | Scheduled | 9m3s | default-scheduler | Successfully assigned local-path-storage/local-path-provisioner-7d6dddf9dd-7hhgx to master1 |
| Warning | FailedCreatePodSandBox | 9m3s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "e2bd24f5fc422df3c85ce9880d1eda6265de9ea7e95ab2896799bd0d5cbf4d28": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 8m52s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "11b27fe7c8fde4061d54627b9531dc9ad93518a9ead1fb947f71973ac7b4ebd": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 8m41s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "108a10b7d50addee59c81156bafc45a4fccd9c23441f6fb3ee4e91d12e7778114": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 8m28s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "9755e3a358659242db0a6b1da3577e4d9a2822f49c6ab9404d6ed48e91ad1e97": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 8m13s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "9a3481abb8a92206fc2a21f6eab3ccf9f2f091f3738293b3ea89ab184b0b441c": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 8m2s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "ac04750ab380279e62b83f8d5a6eefad39ecf1ffef9b119112f89767e4a43f3": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 7m46s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "18b0b1b51e7d2454d74655518e41cac1e119664583731cd466e8470c08fd7d3d": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 7m31s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "8fba02df518f3d586b24610d3f0ddc9b425e53b4106552fcadcfed4a110cec1a": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 7m18s | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "34d789f5e90931248da4026fb8563d9c685cc409e976031bd32eb4fadf6b9d8b": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |
| Warning | FailedCreatePodSandBox | 3m29s (x17 over 7m5s) | kubelet | (combined from similar events): Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "aef83303396dff043d10fe805d439f29cefc3fb0f1a20236f7fa81e99d902e1": plugin type="calico" failed (add): error getting ClusterInformation: connection is unauthorized: Unauthorized |

From above, the pod isn't starting because it can't even get a network interface.

1. Checked calico pods -- all were Running

| | | | | |
|-------------|--|-----|---------|------|
| kube-system | calico-kube-controllers-7959b6fcf8-wbcvc | 1/1 | Running | 0 |
| 6d7h | | | | |
| kube-system | calico-node-lbzg7 | 1/1 | Running | 0 |
| kube-system | calico-node-wb6lq | 1/1 | Running | 0 |
| | | | | 6d7h |

1. Checked one calico pod logs and no errors were seen.

1. Checked cluster status

```
# kubectl get clusterinformation.crd.projectcalico.org default -oyaml
apiVersion: crd.projectcalico.org/v1
kind: ClusterInformation
metadata:
  annotations:
    projectcalico.org/metadata: '{"creationTimestamp":"2025-09-08T10:54:36Z"}'
  creationTimestamp: "2025-09-08T10:54:36Z"
  generation: 1
  name: default
  resourceVersion: "1915"
  uid: 25f4597b-b245-4a9f-a542-998ca4d8ac61
spec:
  calicoVersion: v3.30.3
  clusterGUID: 6af15618022245369ad9a7f5dc5e4c7c
  clusterType: k8s,bgp,kdd
  datastoreReady: true
```

1. Checked if I can go into netshoot AFTER **force** deleting it

```
# k delete po netshoot --grace-period=0 --force
```

(if can go in, means that new pods can be created with networking right now)

```
# kubectl exec -it netshoot -- sh
~ #
```

BUT for my case, I recreated and it got the same calico error as above log events for netshoot pod.

1. Tried to reapply calico yaml but still same error as above when trying to restart netshoot pod.

1. Checked for CNI binaries on worker node (shld show calico, calico-ipam. If missing > kubelet can't call calico)

```
# ls -l /opt/cni/bin/ | grep calico
-rwxr-xr-x 1 root root 79640052 Sep  8 10:51 calico
-rwxr-xr-x 1 root root 79640052 Sep  8 10:51 calico-ipam
```

1. Check connectivity to API server... (AKA master node IP but can check with following command):

```
# kubectl config view --minify -o jsonpath='{.clusters[0].cluster.server}'
https://192.168.81.128:6443
```

```
# curl -vk https://192.168.81.128:6443/healthz
* Trying 192.168.81.128:6443...
* Connected to 192.168.81.128 (192.168.81.128) port 6443 (#0)
* ALPN, offering h2
* ALPN, offering http/1.1
* TLSv1.0 (OUT), TLS header, Certificate Status (22):
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.2 (IN), TLS header, Certificate Status (22):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.2 (IN), TLS header, Finished (20):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, Encrypted Extensions (8):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, Request CERT (13):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, Certificate (11):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, CERT verify (15):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, Finished (20):
* TLSv1.2 (OUT), TLS header, Finished (20):
* TLSv1.3 (OUT), TLS change cipher, Change cipher spec (1):
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.3 (OUT), TLS handshake, Certificate (11):
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.3 (OUT), TLS handshake, Finished (20):
* SSL connection using TLSv1.3 / TLS_AES_128_GCM_SHA256
* ALPN, server accepted to use h2
* Server certificate:
* subject: O=k3s; CN=k3s
* start date: Sep 8 10:09:29 2025 GMT
* expire date: Sep 8 10:09:29 2026 GMT
* issuer: CN=k3s-server-ca@1757326169
* SSL certificate verify result: self-signed certificate in certificate chain (19), continuing anyway.
* Using HTTP2, server supports multiplexing
* Connection state changed (HTTP/2 confirmed)
* Copying HTTP/2 data in stream buffer to connection buffer after upgrade: len=0
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* Using Stream ID: 1 (easy handle 0x555a88c279f0)
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
> GET /healthz HTTP/2
> Host: 192.168.81.128:6443
> user-agent: curl/7.81.0
> accept: */*
>
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.3 (IN), TLS handshake, Newsession Ticket (4):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* Connection state changed (MAX_CONCURRENT_STREAMS == 250)!
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
```

```

* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
* TLSv1.2 (IN), TLS header, Supplemental data (23):
< HTTP/2 401
< audit-id: 1e53d3b7-f894-4988-8cf3-4a7656ae33af
< cache-control: no-cache, private
< content-type: application/json
< content-length: 157
< date: Mon, 15 Sep 2025 01:22:59 GMT
<
* TLSv1.2 (IN), TLS header, Supplemental data (23):
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {},
  "status": "Failure",
  "message": "Unauthorized",
  "reason": "Unauthorized",
  "code": 401
* Closing connection 0
* TLSv1.2 (OUT), TLS header, Supplemental data (23):
* TLSv1.3 (OUT), TLS alert, close notify (256):
}root@worker1:~# kubectl config view --minify -ojsonpath='{.clusters[0].cluster.server}'
error: current-context must exist in order to minify
root@worker1:~#
logout
miki@worker1:~$
logout
Connection to 192.168.81.129 closed.

```

^^The output above shows that we **successfully** connected to the API server. TLS handshake succeeded and server responded so there's no firewall or routing issue.

The response says :

```

"
  "status": "Failure",
  "message": "Unauthorized",
  "reason": "Unauthorized",
  "code": 401

```

Which is normal, because curl didn't send a client cert/token. (curl isn't authenticated)

1. Check that the CNI config file exists

```

root@worker1:~# cat /etc/cni/net.d/10-calico.conflist

{
  "name": "k8s-pod-network",
  "cniVersion": "0.3.1",
  "plugins": [
    {
      "type": "calico",
      "log_level": "info",
      "log_file_path": "/var/log/calico/cni/cni.log",
      "datastore_type": "kubernetes",
      "nodename": "worker1",
      "mtu": 0,
    }
  ]
}

```

```

"ipam": {
    "type": "calico-ipam"
},
"policy": {
    "type": "k8s"
},
"kubernetes": {
    "kubeconfig": "/etc/cni/net.d/calico-kubeconfig"
}
},
{
    "type": "portmap",
    "snat": true,
    "capabilities": {"portMappings": true}
}
]

```

^^Should be a valid JSON file referencing calico as type.
If missing or corrupt > kubelet can't set up networking

1. Checked if calico crds are present. If these aren't applied correctly, calico won't run. Below shows that the crds are healthy and loaded.

```
# kubectl get crd | grep projectcalico
bgpconfigurations.crd.projectcalico.org      2025-09-15T02:29:01Z
bgpfilters.crd.projectcalico.org             2025-09-08T10:48:34Z
bgppeers.crd.projectcalico.org              2025-09-15T02:29:01Z
blockaffinities.crd.projectcalico.org        2025-09-15T02:29:01Z
caliconodestatuses.crd.projectcalico.org     2025-09-15T02:29:01Z
clusterinformations.crd.projectcalico.org    2025-09-15T02:29:01Z
felixconfigurations.crd.projectcalico.org    2025-09-15T02:29:01Z
globalnetworkpolicies.crd.projectcalico.org   2025-09-15T02:29:01Z
globalnetworksets.crd.projectcalico.org       2025-09-15T02:29:01Z
hostendpoints.crd.projectcalico.org          2025-09-15T02:29:01Z
ipamblocks.crd.projectcalico.org            2025-09-15T02:29:01Z
ipamconfigs.crd.projectcalico.org           2025-09-15T02:29:01Z
ipamhandles.crd.projectcalico.org          2025-09-15T02:29:01Z
ippools.crd.projectcalico.org                2025-09-15T02:29:01Z
ipreservations.crd.projectcalico.org         2025-09-15T02:29:01Z
kubecollectorsconfigurations.crd.projectcalico.org 2025-09-15T02:29:01Z
networkpolicies.crd.projectcalico.org        2025-09-15T02:29:01Z
networksets.crd.projectcalico.org            2025-09-15T02:29:01Z
stagedglobalnetworkpolicies.crd.projectcalico.org 2025-09-08T10:48:34Z
stagedkubernetesnetworkpolicies.crd.projectcalico.org 2025-09-08T10:48:34Z
stagednetworkpolicies.crd.projectcalico.org    2025-09-08T10:48:34Z
tiers.crd.projectcalico.org                  2025-09-08T10:48:34Z
```

1. Checked the kubeconfig file in worker node. In the output from #8, we cat the file that is stated under kubernetes, kubeconfig

```
# cat /etc/cni/net.d/calico-kubeconfig
# Kubeconfig file for Calico CNI plugin. Installed by calico/node.
apiVersion: v1
kind: Config
```

```

clusters:
- name: local
  cluster:
    server: https://10.43.0.1:443
    certificate-authority-data:
"LS0tLS1CRUdjTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUJkekNDQVlyZ0F3SUJBZ0lCQURBS0JnZ3Foa2pPUFFRR
EFqQWpNU0V3SHdZRFZRUUREQmhyTTNNdGMyVnkKZG1WeUxXTmhRREUzTIRjek1qWXhOamt3SGhjTk
1qVXdPVEE0TVRBd09USTVXaGNOTXpVd09UQTJNVEF3T1RJNQpXakFqTVNFd0h3WURWUVFEREJock0zT
XRjMlZ5ZG1WeUxXTmhRREUzTIRjek1qWXhOamt3V1RBVEJnY3Foa2pPCIBRSUJCZ2dxaGtqT1BRTUJCd05
DQUFRSHNMZmF4aIN5T3U0a1VacVINK211UGUwUzdGRTd1ZEERm05jNmZLaDIKQytibTRqZ2RKUVYrV2I
yWkFZc3MwZlZdEFxN05PVHNqUzBDYXJMEdcxMXlvMEI3UURBT0JnTIZIUTHCQWY4RQpCQU1DQXFrd0
R3WURWUjBUQVFIL0JBVXdBd0VCL3pBZEJnTIZIUTRFRmdRTI0b1R1cVVtckRybWNFcXQrdm5GCk5NNm
EzWVF3Q2dZSUtvWkl6ajBFQXdJRFNBQXdsUUlnZXpIVGFNeFFPUDQyQ0dCOUFLUGtQVDluWXVaZy8wRI
QKQ2MzeXIQdJFVUVDSVFEQmdZWjVjRTVucEZvd3h3QkNhOEhZaENxK0NKc3A4TGFuN3IUZzVrVWFQ
09Ci0tLS0tRU5EIENFUJRkIDQVRFLS0tLSOK"
users:
- name: calico
  user:
    token:
eyJhbGciOiJSUzI1NlslmtpZCI6IjjDTV9rZzNneFc5RKVldWpuUTh2cktHVC1acmFIRDlfalRscFBqUUJ5Qmcif
Q.eyJhdWQiOlsiaHR0cHM6Ly9rdWJlcmlGVzLmRlZmF1bHQuc3ZjLmNsdXNOZXlubG9jYWwiLCJrM3Mi
SwiZXhwIjoxNzU3OTg1MzQ0LCJpYXQiOjE3NTc4OTg5NDQsImlzcyI6lmh0dBzOi8va3ViZXJuZXRIcy5kZW
ZhdWx0LnN2Yy5jbHVzdGVyLmxvY2FsiwianRpIjoiYTQ1OGU0NTMtZTk1Yi00ZTkyLWEyY2ItMG12Y2Y0Nzk
5MzQ2liwia3ViZXJuZXRIcy5pbyl6eyJuYW1lc3BhY2UiOjJrdWJlXN5c3RlbSlsInNlcnPzY2VhY2NvdW50lp7I
m5hbWUiOijYwxpY28tY25pLBsdWdpbilsInVpZCI6ljA2NDMyNTIzLWQyN2MtNDg2ZS05ZmNkLTkxY2Q
wMWQ1MTc2OCJ9fSwibmJmljoxNzU3ODk4OTQ0LCJzdWliOjzeXN0ZW06c2VydmljZWfjY291bnQ6a3Vi
ZS1zeXN0ZW06Y2FsaWNvLWNuaS1wbHVnaW4ifQ.MKmSIQ1dvVF7Rt93RA-
PjGQxeprjrsjUm7ceC4rXO2OCYh1av3TNdbsFcxMSGok41H7z0n3j4jnnowi8l3YXRz0C26XkUk2MPsTgzVP
glj3xs2Fhl5oZ_TgSk3g3_mRUggS_SrKh_UgUd7JImHnSPG6_BYpAx4wQNvrEH-
BX3flcl6vYymrsURabETLmhnustrqCYhS7
_M2DmwNztswaL7D45W_ho2EwdhS2NDDZ7Ez8mXqZCQ2TDmV3SGqf3Qn23bL8v_B1gLjSTRJmcCS4Ub
K5BcJHOYrb7MXWY6asHx3GQzIZLw5gqarskMUvCWCz1T1jOqcu62MqZGlbrN5Q
contexts:
- name: calico-context
  context:
    cluster: local
    user: calico

```

^^We see in the cluster server section, the IP stated does not match our master node IP.

The IP stated is the Kubernetes service ClusterIP for kubernetes.default.svc. Normally this is fine bc calico CNI shld be able to reach the API server via the cluster DNS and kube-proxy. But in this case, either kube-proxy isn't working properly or the service IP isn't reachable on the worker node which breaks calico CNI API calls.

11. I also did a curl to the ClusterIP API:

```

curl -k https://10.43.0.1:443/healthz
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {},
  "status": "Failure",
  "message": "Unauthorized",
  "reason": "Unauthorized",

```

"code": 401

^^From above we can see that there is a 401 unauthorized error **not a timeout** -

- so the service CIDR is **reachable from worker node**.
- Kube-proxy is routing traffic correctly
- The 401 Unauthorized is just bc curl isn't passing credentials
- So networking to the API svc is fine. That means calico **shld be able to talk to the API via 10.43.0.1**

1. I reapplied calico manifest again and checked pods, they are not coming up now

```
# kubectl get po -A | grep -i calico
kube-system      calico-kube-controllers-7498b9bb4c-9cfq9           0/1   Terminating   0
10m
kube-system      calico-node-gprjf                         0/1   Running     0   6m28s
kube-system      calico-node-s4txj                         0/1   Init:2/3    0   4m7s
```

1. Checked logs of one of the calico pod

```
Warning Unhealthy 3m55s (x2 over 3m57s) kubelet      Readiness probe failed: calico/node is not ready: BIRD is not ready: Error querying BIRD: unable to connect to BIRDv4 socket: dial unix /var/run/calico/bird.ctl: connect: connection refused
Warning Unhealthy 3m46s      kubelet      Readiness probe failed: 2025-09-15 01:52:15.865 [INFO][209] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 3m36s kubelet Readiness probe failed: 2025-09-15 01:52:25.876 [INFO][245] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 3m26s kubelet Readiness probe failed: 2025-09-15 01:52:35.866 [INFO][280] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 3m16s kubelet Readiness probe failed: 2025-09-15 01:52:45.872 [INFO][307] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 3m6s kubelet Readiness probe failed: 2025-09-15 01:52:55.873 [INFO][341] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 2m56s kubelet Readiness probe failed: 2025-09-15 01:53:05.863 [INFO][367] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 2m46s kubelet Readiness probe failed: 2025-09-15 01:53:15.869 [INFO][393] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 2m36s kubelet Readiness probe failed: 2025-09-15 01:53:25.869 [INFO][420] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
Warning Unhealthy 17s (x15 over 2m34s) kubelet (combined from similar events): Readiness probe failed: 2025-09-15 01:55:44.668 [INFO][855] node/health.go 202: Number of node(s) with BGP peering established = 0
calico/node is not ready: BIRD is not ready: BGP not established with 192.168.81.129
```

^^ From above we can see that it says BGP (BIRD) is not ready...

Calico pods are stuck Not Ready bc BGP (BIRD) cannot establish any peers.

In a 2 node cluster, Calico defaults to using BGP peering between nodes but BGP failing > calico readiness probe fails > CNI fails > New pods (workloads and even calico itself) cant get networking

For small clusters, BGP is not really needed so we can switch calico to run in VXLAN mode (overlay

network). BGP is mainly for bigger, prod clusters where you want explicit routing.

1. Change to VXLAN mode

Edit the default IPPool

```
# kubectl edit ippool default-ipv4-ippool -n kube-system
```

Change

```
vxlanMode: Never
```

To:

```
vxlanMode: Always
```

And natOutgoing shld be true.

Then restart calico pods

```
# kubectl delete po -n kube-system -l k8s-app=calico-node
```

1. I had restarted my VMs (cuz I continued troubleshooting next dday) and now all calico pods were up but when I described 1 calico-node pod it shot out this error

Events:

| Type | Reason | Age | From | Message |
|--|---|-----|-------------------|--|
| Normal | Scheduled | 11h | default-scheduler | Successfully assigned kube-system/calico-node-4xrnn to worker1 |
| Warning | FailedCreatePodSandBox | 11h | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to start sandbox |
| "f955d7cce9bfcea0d9139017e2c89956c841454e841416e1172ee7217b845f07" | : failed to get sandbox image "rancher/mirrored-pause:3.6": failed to pull image "rancher/mirrored-pause:3.6": failed to pull and unpack image "docker.io/rancher/mirrored-pause:3.6": failed to resolve reference "docker.io/rancher/mirrored-pause:3.6": failed to authorize: failed to fetch anonymous token: Get " https://auth.docker.io/token?scope=repository%3Arancher%2Fmirrored-pause%3Apull&service=registry.docker.io ": net/http: TLS handshake timeout | | | |
| Warning | FailedCreatePodSandBox | 11h | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to start sandbox |
| "667043e5c17804cdb892fa3a90f164c25db34a7097fe7d077bcbaa25876843d2" | : failed to get sandbox image "rancher/mirrored-pause:3.6": failed to pull image "rancher/mirrored-pause:3.6": failed to pull and unpack image "docker.io/rancher/mirrored-pause:3.6": failed to copy: httpReadSeeker: failed open: failed to authorize: failed to fetch anonymous token: Get " https://auth.docker.io/token?scope=repository%3Arancher%2Fmirrored-pause%3Apull&service=registry.docker.io ": net/http: TLS handshake timeout | | | |
| Warning | FailedCreatePodSandBox | 11h | kubelet | Failed to create pod sandbox: rpc error: code = Unknown desc = failed to start sandbox |
| "b86eef5503797237463e24bfb5d6ab1453cf0ee5973be957516fb63649ee4015" | : failed to get sandbox image "rancher/mirrored-pause:3.6": failed to pull image "rancher/mirrored-pause:3.6": failed to pull and unpack image "docker.io/rancher/mirrored-pause:3.6": failed to copy: httpReadSeeker: failed open: failed to authorize: failed to fetch anonymous token: Get " https://auth.docker.io/token?scope=repository%3Arancher%2Fmirrored-pause%3Apull&service=registry.docker.io ": dial tcp 3.88.79.16:443: i/o timeout | | | |

```

Warning FailedCreatePodSandBox 11h kubelet      Failed to create pod sandbox: rpc error: code = Unknown desc = failed to start sandbox
"18fa76bf5588026acdb2648a3ded632beaa516c39bbc455d35df5fbe83ee541a": failed to get sandbox image "rancher/mirrored-pause:3.6": failed to pull image "rancher/mirrored-pause:3.6": failed to pull and unpack image "docker.io/rancher/mirrored-pause:3.6": failed to resolve reference "docker.io/rancher/mirrored-pause:3.6": failed to do request: Head
"https://registry-1.docker.io/v2/rancher/mirrored-pause/manifests/3.6https://registry-1.docker.io/v2/rancher/mirrored-pause/manifests/3.6Warning Unhealthy        4m43s kubelet      Readiness probe failed: calico/node is not ready: felix is not ready: Get "http://localhost:9099/readiness
Warning Unhealthy        4m42s kubelet      Readiness probe failed: calico/node is not ready: felix is not ready: readiness probe reporting 503

```

Felix is the main dataplane agent inside each calico-node pod. Readiness probe is failing because felix can't start its internal API on port 9099.

Common causes - IPPool misconfiguration(wrong mode like CrossSubnet > BGP dependency)

```
# kubectl get ippool default-ipv4-ippool -oyaml
apiVersion: crd.projectcalico.org/v1
kind: IPPool
metadata:
annotations:
  projectcalico.org/metadata: '{"uid":"0dbc429-e106-4774-bc49-3f615231c793","creationTimestamp":"2025-09-15T13:35:22Z"}'
  creationTimestamp: "2025-09-15T13:35:22Z"
```

```

generation: 1
name: default-ipv4-ippool
resourceVersion: "15462"
uid: aacb61b2-15e3-4675-93a5-6de217b6520f
spec:
  allowedUses:
    - Workload
    - Tunnel
  blockSize: 26
  cidr: 172.16.0.0/16
  ipipMode: Never
  natOutgoing: true
  nodeSelector: all()
  vxlanMode: CrossSubnet

```

1. Checked that kernel vxlan support is running in both nodes with lsmod | grep vxlan (shld show up) -- if doesn't show up need to use modprobe to load module.
2. Patched ippool vxlanMode to Always

```

kubectl patch ippool default-ipv4-ippool \
> --type merge -p '{"spec": {"vxlanMode": "Always", "ipipMode": "Never", "natOutgoing": true}}'

```

Restarted pods and now this error:

| Events: | | | | |
|----------------|------------------|-----|-------------------|--|
| Type | Reason | Age | From | Message |
| Normal | Scheduled | 37s | default-scheduler | Successfully assigned kube-system/calico-node-cd71d to worker1 |
| Normal | Pulled | 31s | kubelet | Container image "docker.io/calico/cni:v3.25.0" already present on machine |
| Normal | Created | 31s | kubelet | Created container: upgrade-ipam |
| Normal | Started | 31s | kubelet | Started container upgrade-ipam |
| Normal | Pulled | 30s | kubelet | Container image "docker.io/calico/cni:v3.25.0" already present on machine |
| Normal | Created | 30s | kubelet | Created container: install-cni |
| Normal | Started | 30s | kubelet | Started container install-cni |
| Normal | Pulled | 27s | kubelet | Container image "docker.io/calico/node:v3.25.0" already present on machine |
| Normal | Created | 27s | kubelet | Created container: mount-bpffs |
| Normal | Started | 27s | kubelet | Started container mount-bpffs |
| Normal | Pulled | 26s | kubelet | Container image "docker.io/calico/node:v3.25.0" already present on machine |
| Normal | Created | 26s | kubelet | Created container: calico-node |
| Normal | Started | 26s | kubelet | Started container calico-node |
| Warning | Unhealthy | 25s | kubelet | Readiness probe failed: calico/node is not ready: felix is not ready: Get "http://localhost:9099/readiness": dial tcp 127.0.0.1:9099: connect: connection refused |

Calico pod is starting but readiness is still failing.. The problem is now inside felix itself.

But pods themselves are Running

I created 2 test pods and did a ping test to each other to check pod connectivity

```
kubectl run test1 --image=busybox:1.28 --restart=Never -- sleep 3600  
kubectl run test2 --image=busybox:1.28 --restart=Never -- sleep 3600  
Kubectl exec -it test1 -- ping <test2 IP>
```

OK this didn't work - I realise because of my netpol default deny all. Ping works in other ns.

1. Suspect maybe due to version mismatch (first time I downloaded calico was 3.30.3 version. After that when I installed it was 3.25x version. So the first crds may have been 3.30.3 and the older version could not work with these crds.)

```
# kubectl delete -f  
https://raw.githubusercontent.com/projectcalico/calico/v3.30.3/manifests/calico.yaml  
# kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.30.3/manifests/calico-vxlan.yaml
```

Now pods are running and checking the felix container itself with

```
# kubectl logs <pod-name> -n kube-system -c calico-node | tail -n 50
```

Checked Readiness of pods directly

```
kubectl -n kube-system get po calico-node-9dgn8 \  
> -o jsonpath='{.status.conditions[?(@.type=="Ready")].status}{"\n"}'  
True
```

ISSUE 3: CoreDNS failure found

**Due to Config Map misconfig

I saw that I couldn't ping by name when doing the pod to pod connectivity test in nettest ns. But I could ping by pod IP.

1. Checked coreDNS pod is running

```
# kubectl get po -n kube-system -l k8s-app=kube-dns -owide  
  
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE  
READINESS GATES  
coredns-64fd4b4794-csng7 1/1 Running 2 (66m ago) 7d4h 172.16.235.130 worker1 <none>  
<none>
```

^^Pod is running

1. Checked svc is running

```
# kubectl get svc kube-dns -n kube-system  
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE  
kube-dns ClusterIP 10.43.0.10 <none> 53/UDP,53/TCP,9153/TCP 7d4h
```

^^Svc exists at 10.43.0.10:53/UDP

1. Checked DNS from pod

```
# k exec -it netshoot -- nslookup kubernetes.default.svc.cluster.local  
;; communications error to 10.43.0.10#53: timed out
```

```
;; communications error to 10.43.0.10#53: timed out  
;; communications error to 10.43.0.10#53: timed out  
;; no servers could be reached
```

```
command terminated with exit code 1
```

^^Pods cannot reach the kube-dns Service IP > Service routing or kube-proxy/iptables issue.

1. Tested dns resolution from a dnsutils pod

```
# kubectl run -it dnsutils --image=infoblox/dnstable --restart=Never --rm --command -- sh  
If you don't see a command prompt, try pressing enter.  
dnstable# nslookup kubernetes.default.svc.cluster.local  
Server: 10.43.0.10  
Address: 10.43.0.10#53  
  
** server can't find kubernetes.default.svc.cluster.local.localdomain: SERVFAIL  
  
dnstable# kubectl -n kube-^C  
dnstable# ^C  
dnstable# E0915 22:53:42.978134 4392 v2.go:104] "Unhandled Error" err="write on closed stream 0"  
pod "dnsutils" deleted  
pod default/dnsutils terminated (Error)
```

^^ CoreDNS svc is reachable so service/clusterIP routing is fine. But coreDNS fails to resolve queries (SERVFAIL)

1. Checked logs of coredns pod

```
# kubectl -n kube-system logs -l k8s-app=kube-dns  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.override  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.server  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.override  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.server  
[ERROR] plugin/errors: 2 kubernetes.default.svc.cluster.local.localdomain. A: read udp  
172.16.235.130:60498->192.168.81.2:53: i/o timeout  
[ERROR] plugin/errors: 2 kubernetes.default.svc.cluster.local.localdomain. A: read udp  
172.16.235.130:58302->192.168.81.2:53: i/o timeout  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.override  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.server  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.override  
[WARNING] No files matching import glob pattern: /etc/coredns/custom/*.server
```

^^ Above shows that coreDNS is trying to resolve kubernetes.default.svc.cluster.local

- Instead of answering from the in-cluster K8s plugin, it's forwarding queries upstream (to 192.168.81.2:53)
- That upstream (probably host VM's DNS) is not reachable from inside the CoreDNS pod
- So DNS lookup fails with SERVFAIL

1. Amend configmap config

```
# kubectl edit configmap -n kube-system coredns
```

```
apiVersion: v1  
data:  
  Corefile: |  
    .:53 {  
      errors
```

```

health
ready
kubernetes cluster.local in-addr.arpa ip6.arpa {
    pods insecure
    fallthrough in-addr.arpa ip6.arpa
}
hosts /etc/coredns/NodeHosts {
    ttl 60
    reload 15s
    fallthrough
}
prometheus :9153
# OLD : forward . /etc/resolv.conf
# NEW : forward . 8.8.8.8 1.1.1.1
cache 30
loop
reload
loadbalance
import /etc/coredns/custom/*.override
}
import /etc/coredns/custom/*.server
NodeHosts: |
    192.168.81.128 master1
    192.168.81.129 worker1
kind: ConfigMap
metadata:
annotations:
    objectset.rio.cattle.io/applied:
H4sIAAAAAAA/4yQwWrzMBCEX0Xs2fEf20nsX9BDybH02IMva2kdq1Z2g6SkBJN3L8IUCiVtbyNGOzvfzoAn90lhOmHQcKmgAIsJQc+wl0CD8wQaSr1t1PzKSilFIUilix4JfRoXHQjtdZHTuafAlCgq488xUSi9wK2AybEFDXvhwR2e8QQFHCnh50ZkloTJCcf8IP6NTlqUyuCkNJiSp9LJP5czoLjryztTWB0uE2iYmvjFuVSFenJsHx6tFf41gvGY6Y0Eshz/9D2e0OSZfIJVvMZExwzusSf/I9SlcQQNvaG6a+r/XVdV7abBddPtsN9W66EediON7aberM22zaHf6t0tcPsIAAD//8Ix+PfoAQAA
    objectset.rio.cattle.io/id: ""
    objectset.rio.cattle.io/owner-gvk: k3s.cattle.io/v1, Kind=Addon
    objectset.rio.cattle.io/owner-name: coredns
    objectset.rio.cattle.io/owner-namespace: kube-system
creationTimestamp: "2025-09-08T10:09:39Z"
labels:
    objectset.rio.cattle.io/hash: bce283298811743a0386ab510f2f67ef74240c57
name: coredns
namespace: kube-system
resourceVersion: "559"
uid: 8f9e5d04-3d23-49a6-8307-6e9e05e7efd6

```

1. Restart deployment

```
# k rollout restart deploy -n kube-system coredns
```

1. Check again

```
# kubectl run -it dnsutils --image=infoblox/dnstablets --restart=Never --rm --command -- sh
If you don't see a command prompt, try pressing enter.
dnstablets# nslookup kubernetes.default.svc.cluster.local
```

```
Server: 10.43.0.10
Address: 10.43.0.10#53
```

```
Name: kubernetes.default.svc.cluster.local
Address: 10.43.0.1
```

ISSUE 4: (RELATED TO ISSUE 1) Local-path-storage pods not coming up

Broken config. Need to reapply manifests and also create secret.

- Upon checking local-path-storage pods, all were not up

```
# k get all -n local-path-storage
NAME READY STATUS RESTARTS AGE
pod/local-path-provisioner-5b5f8bd697-44ss2 0/1 CrashLoopBackOff 3 (20s ago) 65s
pod/local-path-provisioner-7d6dddf9dd-gdbkq 0/1 CrashLoopBackOff 10 (2m7s ago) 28m

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/local-path-provisioner 0/1 1 0 21h

NAME DESIRED CURRENT READY AGE
replicaset.apps/local-path-provisioner-5b5f8bd697 1 1 0 65s
replicaset.apps/local-path-provisioner-7d6dddf9dd 1 1 0 21h
```

- Logs:

```
# k logs -n local-path-storage pod/local-path-provisioner-5b5f8bd697-44ss2
time="2025-09-15T15:07:38Z" level=fatal msg="Error starting daemon: invalid empty flag helper-pod-file and it also does not exist at ConfigMap local-path-storage/local-path-config with err: configmaps \"local-path-config\" is forbidden: User \"system:serviceaccount:local-path-storage:local-path-provisioner-service-account\" cannot get resource \"configmaps\" in API group \"\" in the namespace \"local-path-storage\""
```

Seems cm local-path-config is missing in the ns. The manifest applied may have been incomplete/mismatched.

- Force delete broken ns

```
# kubectl delete ns local-path-storage --force --grace-period=0
```

Terminating state for so long because of the finalizer (kubernetes)

When doing k get all -n local-path-storage, it showed no resources but ns was stuck in terminating state. This means that the ns is already empty but only the ns object itself is stuck in Term state bc of its own finalizer.

- Save ns object and remove finalizers block

```
# k get ns local-path-storage -o json > ns.json
```

- Apply ns.json via API finalize endpoint

```
# kubectl proxy --port=8001 & curl -s -H "Content-Type: application/json" -X PUT --data-binary @ns.json \
> http://127.0.0.1:8001/api/v1/namespaces/local-path-storage/finalize
[2] 2697
error: listen tcp 127.0.0.1:8001: bind: Only one usage of each socket address (protocol/network
```

address/port) is normally permitted.

```
{  
  "kind": "Namespace",  
  "apiVersion": "v1",  
  "metadata": {  
    "name": "local-path-storage",  
    "uid": "4bd4cf19-8b37-49f5-a78c-9692a1ff6062",  
    "resourceVersion": "20211",  
    "creationTimestamp": "2025-09-14T17:36:15Z",  
    "deletionTimestamp": "2025-09-15T15:10:14Z",  
    "labels": {  
      "kubernetes.io/metadata.name": "local-path-storage"  
    },  
    "annotations": {  
      "kubectl.kubernetes.io/last-applied-configuration": "{\"apiVersion\":\"v1\",\"kind\":\"Namespace\",\"metadata\":{\"annotations\":{},\"name\":\"local-path-storage\"}}\n    },  
    "managedFields": [  
      {  
        "manager": "kubectl-client-side-apply",  
        "operation": "Update",  
        "apiVersion": "v1",  
        "time": "2025-09-14T17:36:15Z",  
        "fieldsType": "FieldsV1",  
        "fieldsV1": {  
          "f:metadata": {  
            "f:annotations": {  
              ".": {}  
            },  
            "f:kubernetes.io/last-applied-configuration": {}  
          },  
          "f:labels": {  
            ".": {}  
          },  
          "f:kubernetes.io/metadata.name": {}  
        }  
      }  
    ],  
    {  
      "manager": "k3s",  
      "operation": "Update",  
      "apiVersion": "v1",  
      "time": "2025-09-15T15:10:21Z",  
      "fieldsType": "FieldsV1",  
      "fieldsV1": {  
        "f:status": {  
          "f:conditions": {  
            ".": {}  
          },  
          "k:{\"type\":\"NamespaceContentRemaining\":[  
            \".\" : {},  
            "f:lastTransitionTime": {},  
            "f:message": {},  
            "f:reason": {},  
            "f:status": {}  
          ]}  
        }  
      }  
    }  
  ]  
}
```

```

    "f:type": {},
  },
  "k:{\"type\":\"NamespaceDeletionContentFailure\":[
    ".": {},
    "f:lastTransitionTime": {},
    "f:message": {},
    "f:reason": {},
    "f:status": {},
    "f:type": {}
  ],
  "k:{\"type\":\"NamespaceDeletionDiscoveryFailure\":[
    ".": {},
    "f:lastTransitionTime": {},
    "f:message": {},
    "f:reason": {},
    "f:status": {},
    "f:type": {}
  ],
  "k:{\"type\":\"NamespaceDeletionGroupVersionParsingFailure\":[
    ".": {},
    "f:lastTransitionTime": {},
    "f:message": {},
    "f:reason": {},
    "f:status": {},
    "f:type": {}
  ],
  "k:{\"type\":\"NamespaceFinalizersRemaining\":[
    ".": {},
    "f:lastTransitionTime": {},
    "f:message": {},
    "f:reason": {},
    "f:status": {},
    "f:type": {}
  ]
}
},
"subresource": "status"
}
]
},
"spec": {},
"status": {
  "phase": "Terminating",
  "conditions": [
    {
      "type": "NamespaceDeletionDiscoveryFailure",
      "status": "True",
      "lastTransitionTime": "2025-09-15T15:10:19Z",
      "reason": "DiscoveryFailed",
      "message": "Discovery failed for some groups, 1 failing: unable to retrieve the complete list of server APIs: metrics.k8s.io/v1beta1: stale GroupVersion discovery: metrics.k8s.io/v1beta1"
    }
  ]
}

```

```
{
  "type": "NamespaceDeletionGroupVersionParsingFailure",
  "status": "False",
  "lastTransitionTime": "2025-09-15T15:10:19Z",
  "reason": "ParsedGroupVersions",
  "message": "All legacy kube types successfully parsed"
},
{
  "type": "NamespaceDeletionContentFailure",
  "status": "False",
  "lastTransitionTime": "2025-09-15T15:10:21Z",
  "reason": "ContentDeleted",
  "message": "All content successfully deleted, may be waiting on finalization"
},
{
  "type": "NamespaceContentRemaining",
  "status": "False",
  "lastTransitionTime": "2025-09-15T15:10:19Z",
  "reason": "ContentRemoved",
  "message": "All content successfully removed"
},
{
  "type": "NamespaceFinalizersRemaining",
  "status": "False",
  "lastTransitionTime": "2025-09-15T15:10:19Z",
  "reason": "ContentHasNoFinalizers",
  "message": "All content-preserving finalizers finished"
}
]
}
}[2]+ Exit 1          kubectl proxy --port=8001
```

NS was gone after.

1. Reapply manifest

```
# kubectl apply -f https://raw.githubusercontent.com/rancher/local-path-provisioner/master/deploy/local-path-storage.yaml

namespace/local-path-storage created
serviceaccount/local-path-provisioner-service-account created
role.rbac.authorization.k8s.io/local-path-provisioner-role created
clusterrole.rbac.authorization.k8s.io/local-path-provisioner-role unchanged
rolebinding.rbac.authorization.k8s.io/local-path-provisioner-bind created
clusterrolebinding.rbac.authorization.k8s.io/local-path-provisioner-bind unchanged
deployment.apps/local-path-provisioner created
storageclass.storage.k8s.io/local-path unchanged
configmap/local-path-config created
```

1. Check local-path-storage pods are up (previously they were not coming up)

```
# k get all -n local-path-storage
NAME                  READY STATUS RESTARTS AGE
pod/local-path-provisioner-7d6dddf9dd-8ngmq 1/1   Running 0      47s
```

```
NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/local-path-provisioner 1/1 1 1 47s
```

```
NAME DESIRED CURRENT READY AGE
replicaset.apps/local-path-provisioner-7d6dddf9dd 1 1 1 47s
```

1. Check sc, pvc, pv all are up (now they are all there! Previously they weren't)

```
# k get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE
ALLOWVOLUMEEXPANSION AGE
local-path (default) rancher.io/local-path Delete WaitForFirstConsumer false 7d5h

# k get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES
STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
postgres-storage-postgres-0 Bound pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d 1Gi RWO
local-path <unset> 22h

# k get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM
STORAGECLASS VOLUMEATTRIBUTESCLASS REASON AGE
pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d 1Gi RWO Delete Bound
default/postgres-storage-postgres-0 local-path <unset> 14h
```

1. Check postgres Pod

It has ContainerConfigError

```
# k get po | grep postgres -i
postgres-0 0/1 CreateContainerConfigError 0 30s
postgres-76d88db5cd-wxm6h 1/1 Running 2 (139m ago) 7d2h
```

Events:

| Type | Reason | Age | From | Message |
|----------------|---------------|--------------------------|-------------------|--|
| --- | --- | --- | --- | --- |
| Normal | Scheduled | 42s | default-scheduler | Successfully assigned default/postgres-0 to worker1 |
| Normal | Pulled | 13s (x5 over 42s) | kubelet | Container image "postgres:15" already present on machine |
| Warning | Failed | 13s (x5 over 42s) | kubelet | Error: secret "postgres-secret" not found |

1. Create secret and verify

```
# kubectl create secret generic postgres-secret \
> --from-literal=password=admin123
secret/postgres-secret created
```

```
# k get secret postgres-secret -oyaml
apiVersion: v1
data:
  password: YWRtaW4xMjM=
kind: Secret
```

```
metadata:  
creationTimestamp: "2025-09-15T15:57:28Z"  
name: postgres-secret  
namespace: default  
resourceVersion: "21865"  
uid: 0372ff80-5862-444d-837a-b5718e7a7a9d  
type: Opaque
```

*The Postgres STS container gets its initial credentials from `postgres-secret` for Postgres itself (at container startup)

When Postgres first runs, it bootstraps the db with these values (creates the user and db)

This is not about apps connecting through a service - it's literally how Postgres knows its admin username, pw and which db to initialize.

K8s injects the secret into the container env, so postgres pod takes the value from the Secret's `POSTGRES_PASSWORD`.

- For `POSTGRES_USER` and `POSTGRES_DB`, you can:
 - Define them explicitly in the STS and point them to keys in the Secret OR
 - Let Postgres default them (`postgres` user, db `postgres` if nothing is set)

1. Restart pod and now pod is running FINALLY!

```
# k get po | grep -i postgres  
postgres-0          1/1   Running   0           18s  
postgres-76d88db5cd-wxm6h 1/1   Running   2 (144m ago)  7d2h
```

Troubleshooting Log 3

Sunday, 28 September 2025 11:15 pm

*****ALL BELOW ISSUES ARE FROM INGRESS SECTION.**

6 ISSUES

ISSUE 1 : Command hung and already made resources in first try prevented second try from continuing, with error shown:

Error: INSTALLATION FAILED: Unable to continue with install: ClusterRole "ingress-nginx" in namespace "" exists and cannot be imported into the current release: invalid ownership metadata; annotation validation error: key "meta.helm.sh/release-namespace" must equal "ingress-nginx": current value is "default"

FIX: Delete already made resources

```
# kubectl delete clusterrole,clusterrolebinding,serviceaccount,configmap,service,deployment,daemonset --selector=app.kubernetes.io/name=ingress-nginx -A

# kubectl get all -A | grep ingress-nginx  (there was 1 pod stuck in terminating state)

# k delete po ingress-nginx-controller-6b96cf684f-n2sd7 -n default --force
# k get ingressclass
# k delete ingressclass nginx

# k get ingressclass
# k get all | grep -i ingress-nginx-admission
# kubectl get validatingwebhookconfigurations.admissionregistration.k8s.io | grep ingress
# kubectl delete validatingwebhookconfiguration ingress-nginx-admission
```

*Run install command again

1. Verify

```
# helm list -n ingress-nginx
NAME      NAMESPACE      REVISION      UPDATED      STATUS      CHART
APP VERSION
ingress-nginx  ingress-nginx  1          2025-09-18 20:00:45.3931897 +0800 +08  deployed
ingress-nginx-4.13.2  1.13.2
```

*ingress-nginx Helm release shld be **deployed**

```
# k get po -n ingress-nginx -owide
NAME      READY  STATUS  RESTARTS  AGE      IP      NODE  NOMINATED
NODE  READINESS GATES
ingress-nginx-controller-6b96cf684f-d69f8  1/1  Running  0      3h17m  172.18.137.126  master1
<none>    <none>

# k get svc -n ingress-nginx
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
ingress-nginx-controller      LoadBalancer  10.43.230.85  <pending>
```

```

80:30644/TCP,443:30589/TCP 3h27m
ingress-nginx-controller-admission ClusterIP 10.43.251.154 <none> 443/TCP
3h27m

# k get ingressclass
NAME CONTROLLER PARAMETERS AGE
nginx k8s.io/ingress-nginx <none> 3h28m
traefik traefik.io/ingress-controller <none> 10d

# k get validatingwebhookconfigurations | grep ingress
ingress-nginx-admission 1 3h28m

```

2. Test with app service

Create sampapp-ingress.yaml

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: sampapp-ingress
  namespace: default
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - host: sampapp.local
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: sampapp-service
            port:
              number: 80

```

kubectl apply -f sampapp-ingress.yaml

3. Map host locally since external IP is pending

Edit /etc/hosts (I added the bold line)

```

# Added by Docker Desktop
192.168.1.9 host.docker.internal
192.168.1.9 gateway.docker.internal
# To allow the same kube context to work on the host and the container:
127.0.0.1 kubernetes.docker.internal
127.0.0.1 sampapp.local
# End of section

```

4. Test connectivity

```

# curl -H "Host: sampapp.local" http://192.168.81.128:30644
<html>
<head><title>503 Service Temporarily Unavailable</title></head>

```

```

<body>
<center><h1>503 Service Temporarily Unavailable</h1></center>
<hr><center>nginx</center>
</body>
</html>

```

*This 503 is a good sign - means request is reaching the ingress-nginx controller but just can't find a backend svc to route to.

So ingress-controller is working fine. The problem is the Ingress or Service wiring.

ISSUE 2 : Error 503 (Either ingress or svc issue)

```
# kubectl get ingress -n default
NAME CLASS HOSTS ADDRESS PORTS AGE
sampapp-ingress nginx sampapp.local 80 6m24s
```

```
# kubectl describe ingress sampapp-ingress -n default
Name:      sampapp-ingress
Labels:    <none>
Namespace: default
Address:
Ingress Class: nginx
Default backend: <default>
Rules:
Host      Path  Backends
----      ---   -----
sampapp.local
   /  sampapp-service:80 (172.18.137.65:5000,172.18.137.70:5000)
Annotations: nginx.ingress.kubernetes.io/rewrite-target: /
```

Events:

| Type | Reason | Age | From | Message |
|--------|--------|-------|--------------------------|--------------------|
| Normal | Sync | 6m42s | nginx-ingress-controller | Scheduled for sync |

**Check that under rules, it shows sampapp.local, and that it says which service and port it's routing to

```
# kubectl get svc -n default
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.43.0.1    <none>        443/TCP    10d
nginx-service  ClusterIP  10.43.104.188 <none>        80/TCP     10d
postgres-service  ClusterIP  10.43.178.155 <none>        5432/TCP   10d
redis-service  ClusterIP  10.43.146.71  <none>        6379/TCP   10d
sampapp-service  ClusterIP  10.43.133.187 <none>        5000/TCP   10d
*From here we can see that there is a port mismatch for sampapp-service (service exposing 5000 vs ingress routing to 80)
```

*Edit sampapp-ingress.yaml so that the port is 5000 - matching svc

Test again

```
# kubectl get ingress -n default
NAME      CLASS      HOSTS      ADDRESS      PORTS      AGE
sampapp-ingress  nginx  sampapp.local  80  12m
```

```
# kubectl describe ingress sampapp-ingress -n default
Name:      sampapp-ingress
Labels:    <none>
Namespace: default
Address:
Ingress Class: nginx
Default backend: <default>
Rules:
Host      Path  Backends
----      ---   -----
sampapp.local
      /  sampapp-service:5000 (172.18.137.65:5000,172.18.137.70:5000)
Annotations: nginx.ingress.kubernetes.io/rewrite-target: /
Events:
Type  Reason  Age         From           Message
----  ----   --   -----   -----
Normal  Sync  54s (x2 over 12m)  nginx-ingress-controller  Scheduled for sync

# kubectl get svc -n default
NAME        TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes  ClusterIP  10.43.0.1   <none>       443/TCP   10d
nginx-service  ClusterIP  10.43.104.188 <none>       80/TCP   10d
postgres-service  ClusterIP  10.43.178.155 <none>       5432/TCP  10d
redis-service  ClusterIP  10.43.146.71  <none>       6379/TCP  10d
sampapp-service  ClusterIP  10.43.133.187 <none>       5000/TCP  10d

# curl -H "Host: sampapp.local" http://192.168.81.128:30644
<html>
<head><title>504 Gateway Time-out</title></head>
<body>
<center><h1>504 Gateway Time-out</h1></center>
<hr><center>nginx</center>
</body>
</html>
```

503 error solved.

ISSUE 3 : 504 error gateway timeout

| |
|--|
| <pre># k get ep sampapp-service -n default -owide Warning: v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice NAME ENDPOINTS AGE sampapp-service 172.18.137.65:5000,172.18.137.70:5000 10d *It shld list pod ips and ports (5000). If it is <none> then svc selector doesn't match pod labels</pre> |
|--|

| |
|--|
| <pre># kubectl describe svc sampapp-service -n default Name: sampapp-service Namespace: default Labels: <none> Annotations: <none> Selector: app=sampapp Type: ClusterIP IP Family Policy: SingleStack</pre> |
|--|

| | |
|--------------------------|---------------------------------------|
| IP Families: | IPv4 |
| IP: | 10.43.133.187 |
| IPs: | 10.43.133.187 |
| Port: | <unset> 5000/TCP |
| TargetPort: | 5000/TCP |
| Endpoints: | 172.18.137.65:5000,172.18.137.70:5000 |
| Session Affinity: | None |
| Internal Traffic Policy: | Cluster |
| Events: | <none> |

```
# kubectl get pods -n default --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
busybox        1/1    Running   16 (30m ago) 10d   app=busybox
netshoot       1/1    Running   11 (29m ago) 3d21h  app=netshoot
nginx-96b9d695-4rwwl  0/1    Completed   0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-64pkw  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-bmgqq  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-bwnqv  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-c6r4z  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-c7mmz  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-d8gtc  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-f4n9z  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-f5sjg  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-ggf96  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-hdch7  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-jkfrf  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-jt7v7  1/1    Running   0      3h41m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-klhfp  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-n2jp6  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-njl7q  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-nxxxg  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-q6zz6  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-s7p2m  0/1    Evicted    0      3h53m  app=nginx,pod-template-
hash=96b9d695
nginx-96b9d695-s9925  0/1    Evicted    0      3h53m  app=nginx,pod-template-
```

| | | | | | |
|---|-----|------------------------|---------------|-------|---|
| hash=96b9d695 | | | | | |
| nginx-96b9d695-srrdm | 0/1 | Evicted | 0 | 3h53m | app=nginx,pod-template-hash=96b9d695 |
| nginx-96b9d695-th6f5 | 0/1 | Evicted | 0 | 3h53m | app=nginx,pod-template-hash=96b9d695 |
| nginx-96b9d695-whw7d | 0/1 | ContainerStatusUnknown | 3 (4h48m ago) | 3d14h | app=nginx,pod-template-hash=96b9d695 |
| nginx-96b9d695-zfcwx | 0/1 | Evicted | 0 | 3h53m | app=nginx,pod-template-hash=96b9d695 |
| pg-backup-29301240-wc98k | 0/1 | Completed | 0 | 17h | |
| batch.kubernetes.io/controller-uid=56d446d1-5e88-4700-94e1-d82f26ef288e,batch.kubernetes.io/job-name=pg-backup-29301240,controller-uid=56d446d1-5e88-4700-94e1-d82f26ef288e,job-name=pg-backup-29301240 | | | | | |
| pg-backup-29302680-wzzjp | 0/1 | Completed | 0 | 4h48m | |
| batch.kubernetes.io/controller-uid=908fb3b7-85ca-41f9-b500-d9d320816056,batch.kubernetes.io/job-name=pg-backup-29302680,controller-uid=908fb3b7-85ca-41f9-b500-d9d320816056,job-name=pg-backup-29302680 | | | | | |
| pg-backup-test-74hj4 | 0/1 | Completed | 0 | 2d23h | |
| batch.kubernetes.io/controller-uid=cfc0aebe-691a-4dad-a51a-a99c5c521895,batch.kubernetes.io/job-name=pg-backup-test,controller-uid=cfc0aebe-691a-4dad-a51a-a99c5c521895,job-name=pg-backup-test | | | | | |
| postgres-0 | 1/1 | Running | 0 | 3h41m | |
| app=postgres,apps.kubernetes.io/pod-index=0,controller-revision-hash=postgres-5748f9c956,statefulset.kubernetes.io/pod-name=postgres-0 | | | | | |
| postgres-76d88db5cd-v9qdj | 1/1 | Running | 0 | 3h41m | app=postgres,pod-template-hash=76d88db5cd |
| postgres-76d88db5cd-vwt68 | 0/1 | Completed | 0 | 3h53m | app=postgres,pod-template-hash=76d88db5cd |
| postgres-76d88db5cd-wxm6h | 0/1 | Completed | 4 (4h48m ago) | 10d | app=postgres,pod-template-hash=76d88db5cd |
| redis-7b986b9f57-b4jtm | 1/1 | Running | 0 | 3h41m | app=redis,pod-template-hash=7b986b9f57 |
| redis-7b986b9f57-jsh2q | 0/1 | Completed | 4 (4h48m ago) | 10d | app=redis,pod-template-hash=7b986b9f57 |
| sampapp-6788976bbf-26f8x | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-6l2pl | 0/2 | Error | 8 | 10d | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-ftzc8 | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-g6797 | 0/2 | Error | 8 | 10d | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-hvhvq | 0/2 | Evicted | 0 | 3h53m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-jpj4w | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-lrfjs | 0/2 | Evicted | 0 | 3h53m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-n9tjk | 2/2 | Running | 0 | 3h41m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-njdgl | 0/2 | ContainerStatusUnknown | 2 | 3h53m | app=sampapp,pod-template-hash=6788976bbf |
| sampapp-6788976bbf-pcpxr | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod- |

| | | | | | |
|--------------------------|-----|-----------|---|-------|------------------|
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-qpbnh | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-r22lc | 2/2 | Running | 0 | 3h58m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-swbv7 | 0/2 | Evicted | 0 | 3h53m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-tmpm7 | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-v5pf6 | 0/2 | Evicted | 0 | 3h53m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-vfwx9 | 0/2 | Evicted | 0 | 3h53m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-vmrmf | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| sampapp-6788976bbf-zzfqq | 0/2 | Evicted | 0 | 3h41m | app=sampapp,pod- |
| template-hash=6788976bbf | | | | | |
| test1 | 0/1 | Completed | 0 | 3d1h | run=test1 |
| test2 | 0/1 | Completed | 0 | 3d1h | run=test2 |

*EP and svc are OK, but there're so many failed pods present (sampapp) so Ingress mightve briefly routed traffic to the evicted pods causing 504s.

Delete all evicted pods

```
# kubectl get pods --all-namespaces | grep Evicted | awk '{print $2 " -n " $1}' | xargs -r kubectl
delete pod  (somehow didn't delete default ns pods)
# kubectl get pods -n default | grep Evicted | awk '{print $1}' | xargs -r kubectl delete pod -n
default
```

Check that only left with the 2 running pods

```
# kubectl get pods -n default -o wide | grep sampapp
sampapp-6788976bbf-n9tjk 2/2 Running 0 3h53m 172.18.137.70 master1
<none> <none>
sampapp-6788976bbf-r22lc 2/2 Running 0 4h10m 172.18.137.65 master1
<none> <none>
```

Check that ep point only to those 2 running pods

```
# kubectl get endpoints sampapp-service -n default -o wide
Warning: v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice
NAME      ENDPOINTS          AGE
sampapp-service  172.18.137.65:5000,172.18.137.70:5000  10d
```

Check ingress wiring

```
# kubectl describe ingress sampapp-ingress -n default
Name:      sampapp-ingress
Labels:    <none>
Namespace: default
Address:
Ingress Class: nginx
Default backend: <default>
Rules:
```

| Host | Path | Backends | | |
|--|--------|--|--------------------------|--------------------|
| ---- | ----- | ----- | | |
| sampapp.local | / | sampapp-service:5000 (172.18.137.65:5000,172.18.137.70:5000) | | |
| Annotations: nginx.ingress.kubernetes.io/rewrite-target: / | | | | |
| Events: | | | | |
| Type | Reason | Age | From | Message |
| ---- | ----- | ----- | ----- | ----- |
| Normal | Sync | 21m (x2 over 32m) | nginx-ingress-controller | Scheduled for sync |

*Shld point to sampapp-service:5000 and events shld be "Scheduled for sync" without errors

Test curl again

```
# curl -H "Host: sampapp.local" http://192.168.81.128:30644
```

Still 504 error

ISSUE 4 Ingress svc external IP pending, svc loadbalancer helpers failing, local-path-provisioner is broken

2 options to fix :

- o Quick workaround (use NodePort directly)
- o Proper fix (install MetalLB for bare metal)

- For the purpose of this lab, I chose option 1

```
# k edit svc ingress-nginx-controller -n ingress-nginx
Changed:
  type: LoadBalancer -> type: NodePort
```

- Saved and checked that below has changed to NodePort:

```
# kubectl get svc -n ingress-nginx
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP  PORT(S)          AGE
ingress-nginx-controller   NodePort   10.43.230.85 <none>
80:30644/TCP,443:30589/TCP 9h
ingress-nginx-controller-admission ClusterIP 10.43.251.154 <none>    443/TCP
9h
```

Still got 504 Gateway Timeout error.

- Checked connectivity from inside ingress-controller pod

```
# kubectl -n ingress-nginx exec -it deploy/ingress-nginx-controller -- curl -s http://sampapp-service.default.svc.cluster.local:5000
```

It hung

ISSUE 5 Nodes unhealthy (CPU and disk)

- Checked nodes health

```
# df -h      -- was 95% for / dir
# top       -- CPU was 3++ (nproc is only 2)
```

- Cleared with:

```
# docker system prune -af --volumes (clear old images)
```

```
# crictl rmi --prune
# journalctl --vacuum-time=2d
```

3. After clearing, both cpu and fs usage went down to healthy levels on node (via df -h and cpu), but checking from cluster, worker1 still has DiskPressure:

```
# kubectl describe node master1 | grep -A5 Conditions
Conditions:
  Type        Status  LastHeartbeatTime          LastTransitionTime        Reason
Message
  ----      -----  -----  -----
NetworkUnavailable  False   Fri, 19 Sep 2025 05:52:39 +0800  Fri, 19 Sep 2025 05:52:39 +0800
CalicolsUp          Calico is running on this node
MemoryPressure      False   Fri, 19 Sep 2025 06:09:30 +0800  Mon, 08 Sep 2025 18:09:36 +0800
KubeletHasSufficientMemory  kubelet has sufficient memory available
DiskPressure        False   Fri, 19 Sep 2025 06:09:30 +0800  Fri, 19 Sep 2025 06:08:49 +0800
KubeletHasNoDiskPressure  kubelet has no disk pressure
```

```
# kubectl describe node worker1 | grep -A5 Conditions
Conditions:
  Type        Status  LastHeartbeatTime          LastTransitionTime        Reason
Message
  ----      -----  -----  -----
NetworkUnavailable  False   Fri, 19 Sep 2025 05:49:42 +0800  Fri, 19 Sep 2025 05:49:42 +0800
CalicolsUp          Calico is running on this node
MemoryPressure      False   Fri, 19 Sep 2025 06:09:08 +0800  Mon, 08 Sep 2025 18:14:07 +0800
KubeletHasSufficientMemory  kubelet has sufficient memory available
DiskPressure        True    Fri, 19 Sep 2025 06:09:08 +0800  Fri, 19 Sep 2025 06:07:16 +0800
KubeletHasDiskPressure  kubelet has disk pressure
```

^^This is because Kubernetes doesn't only check root usage. It tracks ephemeral storage (container writable layers, logs, emptyDirs, etc) under /var/lib/kubelet, /var/lib/containerd, or /var/lib/docker

Deleted all Failed pods

```
# kubectl get pods -A --field-selector=status.phase=Failed -o jsonpath='{range .items[*]}{.metadata.namespace}{" "}{.metadata.name}{"\n"}{end}' \
| xargs -r -n2 sh -c 'kubectl delete pod -n "$0" "$1"'
```

Checked worker1 again and it had no disk pressure

```
# kubectl describe no worker1 | grep -A5 Conditions
Conditions:
  Type        Status  LastHeartbeatTime          LastTransitionTime        Reason
Message
  ----      -----  -----  -----
NetworkUnavailable  False   Fri, 19 Sep 2025 05:49:42 +0800  Fri, 19 Sep 2025 05:49:42 +0800
CalicolsUp          Calico is running on this node
MemoryPressure      False   Fri, 19 Sep 2025 06:19:08 +0800  Mon, 08 Sep 2025 18:14:07 +0800
KubeletHasSufficientMemory  kubelet has sufficient memory available
DiskPressure        False   Fri, 19 Sep 2025 06:19:08 +0800  Fri, 19 Sep 2025 06:14:03 +0800
KubeletHasNoDiskPressure  kubelet has no disk pressure
```

ISSUE 5 FIXED

ISSUE 6 NETPOL

1. Test Ingress from inside cluster

```
# kubectl run curltest --rm -it --image=curlimages/curl --restart=Never -- \
curl http://sampapp-service.default.svc.cluster.local:5000
```

And from outside

```
# curl -H "Host: sampapp.local" http://<node-ip>:<nodeport>
```

Curl was still not working!

The **504 Gateway Timeout** means:

Ingress controller received your request → tried to connect to sampapp-service:5000 → could not get a response in time.

So the real suspects are still:

1. **App pod binding issue** → app is listening on 127.0.0.1:5000 instead of 0.0.0.0:5000.

- Run:

```
kubectl logs -n default <sampapp-pod> -c sampapp
```

Look for the bind address.

2. **Service selector mismatch** → but you already confirmed endpoints exist, so selector is fine.

3. **App crash loop or stuck** → check if the logs show errors when requests arrive.

4. **Ingress timeout settings** → default is 60s, but your curl fails fast, which means the backend isn't reachable at all.

Check for bind address - first address line 0.0.0.0 is correct.

```
# kubectl logs -n default sampapp-6788976bbf-l98p6 -c sampapp
```

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

- * Running on all addresses (0.0.0.0) (bind add)
- * Running on <http://127.0.0.1:5000> (loopback inside container)
- * Running on <http://172.18.137.94:5000> (pod's actual IP)

Press CTRL+C to quit

- * Serving Flask app 'app'
- * Debug mode: off

Tried connectivity from ingress-nginx to pod again

```
# kubectl exec -it -n ingress-nginx <ingress-nginx-pod-name> -- curl -s -H "Host: sampapp.local" \
http://sampapp-service.default.svc.cluster.local:5000
```

It hung. And suddenly it hit me!!! I rmb I set some test netpols for calico earlier... and these were blocking my traffic omg

Allow netshoot > sampapp

I deleted all netpols (will look at proper implementation later on)

And now connectivity is working...

```
# curl -H "Host: sampapp.local" http://192.168.81.128:30644
<html>
```

```
<head>
  <title>Feedback Form</title>
  <link rel="stylesheet" href="/static/style.css">
</head>
<body>
  <div class="container">
    <h2>Feedback Form</h2>
    <form method="POST">
      <input type="text" name="name" placeholder="Your Name" required>
      <input type="email" name="email" placeholder="Your Email" required>
      <textarea name="message" placeholder="Your Feedback" required></textarea>
      <button type="submit">Submit</button>
    </form>
  </div>
</body>
</html>
```

Those netpols messed with me twice now... need to be more mindful

Troubleshooting Log 4

Sunday, 28 September 2025 11:24 pm

*****ALL BELOW ISSUES ARE FROM CI/CD w JENKINS & ARGOCD SECTION.**

12 ISSUES

ISSUE 1 Ephemeral Storage of Nodes Causing All Pods to Fail

After ArgoCD, all pods in cluster started to fail. With errors hinting at ephemeral storage issues. Looking at each node, the disk space on the node itself was high and CPU was going haywire (above cpu cores)

```
root@master1:~# du -sh /var/lib/rancher/k3s/agent/containerd/
3.5G /var/lib/rancher/k3s/agent/containerd/
```

```
root@master1:~# df -h /
Filesystem Size Used Avail Use% Mounted on
/dev/mapper/ubuntu--vg-ubuntu--lv 9.8G 8.3G 970M 90% /
```

90% usage in / and /var/lib/rancher/k3s/agent/containerd alone is 3.5G - kubelet sees less than 1G free and immediately applies pressure

1. Expand virtual disk in Vmware Workstation
 - Power off both VMs
 - In Vmware Workstation, Click on Hard Disk > Expand
 - Current: 20GB > 40GB
2. Boot the VM and check new disk size
 - Check raw disk size

```
miki@master1:~$ lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
loop0      7:0    0 63.8M  1 loop /snap/core20/2599
loop1      7:1    0 63.9M  1 loop /snap/core20/2318
loop2      7:2    0  87M  1 loop /snap/lxd/29351
loop3      7:3    0 89.4M  1 loop /snap/lxd/31333
loop4      7:4    0 50.8M  1 loop /snap/snapd/25202
loop5      7:5    0 49.3M  1 loop /snap/snapd/24792
sda       8:0    0 40G  0 disk
└─sda1     8:1    0   1M  0 part
└─sda2     8:2    0 1.8G  0 part /boot
└─sda3     8:3    0 18.2G 0 part
  └─ubuntu--vg-ubuntu--lv 253:0  0 10G  0 lvm /
sr0       11:0   1 1024M 0 rom
```

- Df -h will still show old root partition size since we haven't extended it

```
miki@master1:~$ df -h
Filesystem      Size  Used Avail Use% Mounted on
tmpfs          388M  3.0M 385M  1% /run
/dev/mapper/ubuntu--vg-ubuntu--lv 9.8G  7.7G 1.6G 84% /
tmpfs          1.9G   0  1.9G  0% /dev/shm
tmpfs          5.0M   0  5.0M  0% /run/lock
```

```

/dev/sda2          1.8G 252M 1.4G 16% /boot
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/79e6e9de8c03dd811eaa879154b5954
ce7598d9cd313dc15949ba52d57292a70/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/31b93d5f3931ef02b4cb39b298b0b874
c5fc9596281e8a6f88a764ce336759ef/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/3f0577a71b8199ac0b3fbcad431ea4de
34cde190fc1593a7467be3e4521502e9/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/45b627f1adf2372fe53249430094496e
244f3025a2a64de788d2c57704875457/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/02d242e8cb36a8a4d53c6025c4afd7c1
c32858996f9ac1d7951e79085c4dd478/shm
shm              64M 28K 64M 1%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/65b8efaae9924521ccf089034d3d8839
de0d7ad7a8710a3b575f5df6410a3e2f/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/369918f2e218ff6ecb12fd6759d448582
a53d9c4ac6bad67b073f0d5912c4679/shm
shm              64M  0 64M 0%
/run/k3s/containerd/io.containerd.grpc.v1 cri/sandboxes/a72764cb74b40cf14f5823857e2722c2
ef415b025cd0f1a12e4520e11f7dd1b/shm
tmpfs            388M 4.0K 388M 1% /run/user/1000

```

3. Extend the filesystem

This is the pv, vg and lv before we do anything:

```

root@master1:~# pvs
PV   VG   Fmt Attr PSize PFree
/dev/sda3  ubuntu-vg lvm2 a-- 18.22g 8.22g

root@master1:~# vgs
VG   #PV #LV #SN Attr  VSize VFree
ubuntu-vg  1  1  0 wz--n- 18.22g 8.22g

root@master1:~# lvs
LV   VG   Attr    LSize Pool Origin Data% Meta% Move Log Cpy%Sync Convert
ubuntu-lv  ubuntu-vg -wi-ao---- 10.00g

```

We need to resize the partition first so that LVM can see the new space

```

# growpart /dev/sda 3
# pvresize /dev/sda3

```

Now pvs shld show:

```

root@master1:~# pvs
PV   VG   Fmt Attr PSize PFree
/dev/sda3  ubuntu-vg lvm2 a-- 38.22g 28.22g

```

When I run vgs, it also shows up so no vgextend is required here.

```

root@master1:~# vgs

```

```
VG #PV #LV #SN Attr VSize VFree  
ubuntu-vg 1 1 0 wz--n- 38.22g 28.22g
```

Extend LV

```
(before) root@master1:~# lvs  
LV VG Attr LSize Pool Origin Data% Meta% Move Log Cpy%Sync Convert  
ubuntu-lv ubuntu-vg -wi-ao---- 10.00g  
  
(got the lvpath from df -h /)  
root@master1:~# lvextend -l +100%FREE /dev/mapper/ubuntu--vg-ubuntu--lv  
Size of logical volume ubuntu-vg/ubuntu-lv changed from 10.00 GiB (2560 extents) to 38.22 GiB (9785  
extents).  
Logical volume ubuntu-vg/ubuntu-lv successfully resized.  
  
(after) root@master1:~# lvs  
LV VG Attr LSize Pool Origin Data% Meta% Move Log Cpy%Sync Convert  
ubuntu-lv ubuntu-vg -wi-ao---- 38.22g
```

Resize fs

```
root@master1:~# resize2fs /dev/mapper/ubuntu--vg-ubuntu--lv  
resize2fs 1.46.5 (30-Dec-2021)  
Filesystem at /dev/mapper/ubuntu--vg-ubuntu--lv is mounted on /; on-line resizing required  
old_desc_blocks = 2, new_desc_blocks = 5  
The filesystem on /dev/mapper/ubuntu--vg-ubuntu--lv is now 10019840 (4k) blocks long.
```

Check new fs size

```
root@master1:~# df -h /  
Filesystem Size Used Avail Use% Mounted on  
/dev/mapper/ubuntu--vg-ubuntu--lv 38G 7.3G 29G 21% /
```

lovely.

Now we need to restore the cluster to its working state...

4. Check health of nodes

```
# kubectl describe no master1 | grep -A5 Conditions  
Conditions:  
Type Status LastHeartbeatTime LastTransitionTime Reason  
Message  
---- ----- ----- ----- -----  
NetworkUnavailable False Sun, 21 Sep 2025 15:02:07 +0800 Sun, 21 Sep 2025 15:02:07 +0800  
CalicolsUp Calico is running on this node  
MemoryPressure False Sun, 21 Sep 2025 15:01:57 +0800 Mon, 08 Sep 2025 18:09:36 +0800  
KubeletHasSufficientMemory kubelet has sufficient memory available  
DiskPressure False Sun, 21 Sep 2025 15:01:57 +0800 Sat, 20 Sep 2025 07:57:24 +0800  
KubeletHasNoDiskPressure kubelet has no disk pressure  
  
# kubectl describe no worker1 | grep -A5 Conditions  
Conditions:
```

| Type | Status | LastHeartbeatTime | LastTransitionTime | Reason |
|----------------------------|--------|---|---------------------------------|--------|
| Message | | | | |
| NetworkUnavailable | False | Sun, 21 Sep 2025 15:02:13 +0800 | Sun, 21 Sep 2025 15:02:13 +0800 | |
| CalicolsUp | | Calico is running on this node | | |
| MemoryPressure | False | Sun, 21 Sep 2025 15:02:05 +0800 | Mon, 08 Sep 2025 18:14:07 +0800 | |
| KubeletHasSufficientMemory | | kubelet has sufficient memory available | | |
| DiskPressure | False | Sun, 21 Sep 2025 15:02:05 +0800 | Sat, 20 Sep 2025 08:03:39 +0800 | |
| KubeletHasNoDiskPressure | | kubelet has no disk pressure | | |

All are False... looking good

5. Check Calico Connectivity

Had to recreate busybox and netshoot pods then test connectivity

```
# k exec -it netshoot -- sh

ping 8.8.8.8  (outbound internet check)
nslookup kubernetes.default  (DNS check)
ping <other-pod-IP>  (pod to pod connectivity test)
curl -vk http://sampapp-service:5000  (service connectivity test)
```

Results:

```
# k exec -it netshoot -- sh
~ # ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=127 time=111 ms
nd64 bytes from 8.8.8.8: icmp_seq=2 ttl=127 time=135 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=127 time=53.2 ms
^C
--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 53.229/99.534/134.561/34.147 ms
```

```
# nslookup kubernetes.default
;; Got recursion not available from 10.43.0.10
;; Got recursion not available from 10.43.0.10
Server:    10.43.0.10
Address:   10.43.0.10#53

Name:  kubernetes.default.svc.cluster.local
Address: 10.43.0.1
;; Got recursion not available from 10.43.0.10
```

```
# k exec -it netshoot -- sh
~ # ping 172.18.235.187
PING 172.18.235.187 (172.18.235.187) 56(84) bytes of data.
64 bytes from 172.18.235.187: icmp_seq=1 ttl=63 time=0.185 ms
64 bytes from 172.18.235.187: icmp_seq=2 ttl=63 time=0.050 ms
64 bytes from 172.18.235.187: icmp_seq=3 ttl=63 time=0.063 ms
^C
```

```

--- 172.18.235.187 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2043ms
rtt min/avg/max/mdev = 0.050/0.099/0.185/0.060 ms

```

```

# k exec -it netshoot -- curl -vk sampapp-service:5000
* Host sampapp-service:5000 was resolved.
* IPv6: (none)
* IPv4: 10.43.133.187
* Trying 10.43.133.187:5000...
* Connected to sampapp-service (10.43.133.187) port 5000
* using HTTP/1.x
> GET / HTTP/1.1
> Host: sampapp-service:5000
> User-Agent: curl/8.14.1
> Accept: */*
>
* Request completely sent off
< HTTP/1.1 200 OK
< Server: Werkzeug/3.1.3 Python/3.10.18
< Date: Sun, 21 Sep 2025 07:20:53 GMT
< Content-Type: text/html; charset=utf-8
< Content-Length: 538
< Connection: close
<
<html>
<head>
  <title>Feedback Form</title>
  <link rel="stylesheet" href="/static/style.css">
</head>
<body>
  <div class="container">
    <h2>Feedback Form</h2>
    <form method="POST">
      <input type="text" name="name" placeholder="Your Name" required>
      <input type="email" name="email" placeholder="Your Email" required>
      <textarea name="message" placeholder="Your Feedback" required></textarea>
      <button type="submit">Submit</button>
    </form>
  </div>
</body>
* shutting down connection #0
</html>

```

Calico Connectivity is good.

6. Check Postgres Health

Check PV

| NAME | CAPACITY | ACCESS MODES | RECLAIM POLICY | STATUS | CLAIM |
|--|-------------------|--------------|----------------|--------|-------------|
| STORAGECLASS | VOLUME ATTRIBUTES | CLASS | REASON | AGE | |
| pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb | 1Gi | RWO | Delete | Bound | default/pg- |

| | | | | | | |
|--|--------------------|-------|--------|-------|-----------------|--|
| backup-pvc | local-path <unset> | 5d14h | | | | |
| pvc-a5cd1783-13d6-4928-94f9-5aa81e9f1aae | 8Gi | RWO | Delete | Bound | jenkins/jenkins | |
| local-path <unset> | 2d7h | | | | | |
| pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d | 1Gi | RWO | Delete | Bound | | |
| default/postgres-storage-postgres-0 | local-path <unset> | 6d6h | | | | |

Check PVC

```
$ k get pvc
NAME           STATUS VOLUME          CAPACITY ACCESS MODES
STORAGECLASS   VOLUME ATTRIBUTES CLASS AGE
pg-backup-pvc  Bound  pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb 1Gi      RWO
local-path     <unset>        5d15h
postgres-storage-postgres-0 Bound  pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d 1Gi      RWO
local-path     <unset>        6d13h
```

Check Postgres pod

```
# k get po -n default | grep postgres
postgres-0      1/1   Running       1 (32m ago) 31h
postgres-76d88db5cd-6wpwj 1/1   Running       1 (32m ago) 31h
```

Exec into postgrespod to run simple query

```
# kubectl exec -it postgres-0 -- psql -U postgres -d postgres
psql (15.14 (Debian 15.14-1.pgdg13+1))
Type "help" for help.

postgres=# 
postgres=# SELECT NOW();
now
-----
2025-09-21 07:32:23.203402+00
(1 row)

postgres=#
-----
```

Created simple pg-client.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: pg-client
  namespace: default
spec:
  containers:
    - name: pg-client
      image: postgres:15
      command: ["sleep", "3600"]
  restartPolicy: Always
```

Test

```
# kubectl exec -it pg-client -- psql -h postgres-service -U postgres -d postgres -c "SELECT version();"
```

Password for user postgres:

version

```
-----  
PostgreSQL 15.14 (Debian 15.14-1.pgdg13+1) on x86_64-pc-linux-gnu, compiled by gcc (Debian  
14.2.0-19) 14.2.0, 64-bit  
(1 row)
```

ISSUE 2: LOCAL-PATH-PROVISIONER-PODS IN CRASHLOOPBACKOFF (RBAC)

Noticed the local-path-provisioner pods were in CrashLoopBackOff state.

I described the log and got this:

```
# kubectl -n local-path-storage logs local-path-provisioner-66776cdb94-scm46  
time="2025-09-21T07:44:49Z" level=fatal msg="Error starting daemon: invalid empty flag helper-pod-  
file and it also does not exist at ConfigMap local-path-storage/local-path-config with err: configmaps  
\\"local-path-config\\" is forbidden: User \"system:serviceaccount:local-path-storage:local-path-  
provisioner-service-account\" cannot get resource \"configmaps\" in API group \"\" in the namespace  
\\"local-path-storage\\\""
```

Above shows that the local-path-provisioner pod is failing because it doesn't have RBAC permission to read the local-path-config CM in its namespace.

But in the postgres pods, PVCs, PVs all are working fine because the PVCs were already bound to the PVs before this local-path-provisioner pod failed. So all of those resources came up fine since they're just using their existing PVs. The failure of this local-path-provisioner pod only affects if you are trying to create new PVCs.

Give service account the correct RBAC role:

```
apiVersion: rbac.authorization.k8s.io/v1  
kind: ClusterRole  
metadata:  
  name: local-path-provisioner-role  
rules:  
  - apiGroups: [""]  
    resources: ["nodes", "persistentvolumeclaims", "persistentvolumes", "configmaps"]  
    verbs: ["get", "list", "watch", "create", "delete", "update", "patch"]  
  
  - apiGroups: [""]  
    resources: ["events"]  
    verbs: ["create", "patch", "update"]  
  
---  
apiVersion: rbac.authorization.k8s.io/v1  
kind: ClusterRoleBinding  
metadata:  
  name: local-path-provisioner-bind  
roleRef:  
  apiGroup: rbac.authorization.k8s.io  
  kind: ClusterRole  
  name: local-path-provisioner-role  
subjects:  
  - kind: ServiceAccount  
    name: local-path-provisioner-service-account
```

```
namespace: local-path-storage
```

Apply manifest

```
# k apply -f local-path-provisioner-role.yaml
clusterrole.rbac.authorization.k8s.io/local-path-provisioner-role configured
clusterrolebinding.rbac.authorization.k8s.io/local-path-provisioner-bind configured
```

Delete the old pods and watch the new one come up

```
# k get po -n local-path-storage
NAME           READY STATUS RESTARTS AGE
local-path-provisioner-66776cdb94-c4nv4 1/1   Running 0      3s
```

7. Check cronjob

Check pvc pv are there

```
# k get pvc
NAME           STATUS VOLUME          CAPACITY ACCESS MODES
STORAGECLASS VOLUME ATTRIBUTES CLASS AGE
pg-backup-pvc Bound  pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb 1Gi     RWO
local-path <unset>        5d15h
postgres-storage-postgres-0 Bound  pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d 1Gi     RWO
local-path <unset>        6d14h

# k get pv
NAME           CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM
STORAGECLASS VOLUME ATTRIBUTES CLASS REASON AGE
pvc-87ab800a-aab9-48a8-a2ea-3d57ab5bd9cb 1Gi     RWO     Delete    Bound default/pg-
backup-pvc      local-path <unset>        5d15h
pvc-a5cd1783-13d6-4928-94f9-5aa81e9f1aae 8Gi     RWO     Delete    Bound jenkins/jenkins
local-path <unset>        2d8h
pvc-bab4c32a-5989-4662-ac1a-5f32c2f6495d 1Gi     RWO     Delete    Bound
default/postgres-storage-postgres-0 local-path <unset>        6d6h
```

Create test job

```
# kubectl create job --from=cronjob/pg-backup pg-backup-test
```

Check that job was created and completed

```
# k get job
NAME           STATUS COMPLETIONS DURATION AGE
pg-backup-29304120 Complete 1/1      2m34s  2d5h
pg-backup-29307000 Complete 1/1      14m    52m
pg-backup-test  Complete 1/1      3s     15s
```

8. Check Ingress

➤ Verify Ingress Controller

```
# kubectl get pods -n ingress-nginx -o wide
```

- Make sure ingress-nginx-controller is **Running** and Ready = 1/1.
- Check logs for errors:

```
# kubectl logs -n ingress-nginx deploy/ingress-nginx-controller
```

➤ Verify ingress resource

```
# kubectl get ingress -A  
# kubectl describe ingress sampapp-ingress -n default
```

Confirm HOSTS = sampapp.local

Backend points to sampapp-service:5000

➤ Verify service + pods

I checked that sampapp-service endpoints matched sampapp pod ips, Service has ClusterIP and port 5000, pods are running and container listening on port 5000

```
# kubectl get svc -n default sampapp-service  
  
# k get endpoints sampapp-service  
sampapp-service 172.18.137.71:5000,172.18.235.140:5000 12d  
  
# k get po -owide | grep sampapp  
sampapp-f6c65db94-l72nz 2/2 Running 0 53m 172.18.235.140 worker1  
<none> <none>  
sampapp-f6c65db94-wwfg4 2/2 Running 0 45m 172.18.137.71 master1  
<none> <none>
```

➤ In-cluster curl

```
# kubectl exec -it -n ingress-nginx deploy/ingress-nginx-controller -- \  
curl -s -H "Host: sampapp.local" http://sampapp-service.default.svc.cluster.local:5000  
<html>  
<head>  
  <title>Feedback Form</title>  
  <link rel="stylesheet" href="/static/style.css">  
</head>  
<body>  
  <div class="container">  
    <h2>Feedback Form</h2>  
    <form method="POST">  
      <input type="text" name="name" placeholder="Your Name" required>  
      <input type="email" name="email" placeholder="Your Email" required>  
      <textarea name="message" placeholder="Your Feedback" required></textarea>  
      <button type="submit">Submit</button>  
    </form>  
  </div>  
</body>  
</html>
```

➤ From workstation curl

```
curl -H "Host: sampapp.local" http://192.168.81.128:30644  
<html>  
<head>  
  <title>Feedback Form</title>  
  <link rel="stylesheet" href="/static/style.css">  
</head>  
<body>
```

```

<div class="container">
  <h2>Feedback Form</h2>
  <form method="POST">
    <input type="text" name="name" placeholder="Your Name" required>
    <input type="email" name="email" placeholder="Your Email" required>
    <textarea name="message" placeholder="Your Feedback" required></textarea>
    <button type="submit">Submit</button>
  </form>
</div>
</body>
</html>

```

ISSUE 3 : UNKNOWN SYNC STATUS

Due to netpol

Sync status kept showing Unknown:

```
# k get app -n argocd
NAME      SYNC STATUS  HEALTH STATUS
sampapp  Unknown      Healthy
```

Argocd repo server pod running

```
# k get po -n argocd | grep repo
argocd-repo-server-6fcc5759b-lz7hl          1/1  Running  1 (13m ago)  3d4h
```

DNS test to github

```
# k run -it --rm --image=busybox:1.28 dns-test --restart=Never -- nslookup github.com
Server: 10.43.0.10
Address 1: 10.43.0.10 kube-dns.kube-system.svc.cluster.local

Name: github.com
Address 1: 20.205.243.166
pod "dns-test" deleted
```

Test dns from inside repo server pod

```
# k exec -n argocd -it argocd-repo-server-6fcc5759b-lz7hl -c argocd-repo-server -- sh
$ getent hosts github.com
20.205.243.166 github.com.localdomain
```

Test from a netshoot pod within argocd NS

```
# k run netshoot -n argocd --rm -it --image=nicolaka/netshoot -- bash
If you don't see a command prompt, try pressing enter.
netshoot:~# curl -I https://github.com
HTTP/2 200
date: Wed, 24 Sep 2025 12:36:54 GMT
content-type: text/html; charset=utf-8
vary: X-PJAX, X-PJAX-Container, Turbo-Visit, Turbo-Frame, X-Requested-With, Accept-Language,Accept-Encoding, Accept, X-Requested-With
content-language: en-US
etag: W/"bd1f03eeece7fbf502d53ea53e83bbee"
cache-control: max-age=0, private, must-revalidate
```

strict-transport-security: max-age=31536000; includeSubdomains; preload
x-frame-options: deny
x-content-type-options: nosniff
x-xss-protection: 0
referrer-policy: origin-when-cross-origin, strict-origin-when-cross-origin
content-security-policy: default-src 'none'; base-uri 'self'; child-src github.githubassets.com
github.com/assets-cdn/worker/ github.com/assets/ gist.github.com/assets-cdn/worker/;
connect-src 'self' uploads.github.com www.githubstatus.com collector.github.com
raw.githubusercontent.com api.github.com github-cloud.s3.amazonaws.com github-production-
repository-file-5c1aeb.s3.amazonaws.com github-production-upload-manifest-
file-7fdce7.s3.amazonaws.com github-production-user-asset-6210df.s3.amazonaws.com
.rel.tunnels.api.visualstudio.com wss://.rel.tunnels.api.visualstudio.com
github.githubassets.com objects-origin.githubusercontent.com copilot-
proxy.githubusercontent.com proxy.individual.githubcopilot.com
proxy.business.githubcopilot.com proxy.enterprise.githubcopilot.com
.actions.githubusercontent.com wss://.actions.githubusercontent.com
productionresultssa0.blob.core.windows.net/ productionresultssa1.blob.core.windows.net/
productionresultssa2.blob.core.windows.net/ productionresultssa3.blob.core.windows.net/
productionresultssa4.blob.core.windows.net/ productionresultssa5.blob.core.windows.net/
productionresultssa6.blob.core.windows.net/ productionresultssa7.blob.core.windows.net/
productionresultssa8.blob.core.windows.net/ productionresultssa9.blob.core.windows.net/
productionresultssa10.blob.core.windows.net/ productionresultssa11.blob.core.windows.net/
productionresultssa12.blob.core.windows.net/ productionresultssa13.blob.core.windows.net/
productionresultssa14.blob.core.windows.net/ productionresultssa15.blob.core.windows.net/
productionresultssa16.blob.core.windows.net/ productionresultssa17.blob.core.windows.net/
productionresultssa18.blob.core.windows.net/ productionresultssa19.blob.core.windows.net/
github-production-repository-image-32fea6.s3.amazonaws.com github-production-release-
asset-2e65be.s3.amazonaws.com insights.github.com wss://alive.github.com wss://alive-
staging.github.com api.githubcopilot.com api.individual.githubcopilot.com
api.business.githubcopilot.com api.enterprise.githubcopilot.com edge.fullstory.com
rs.fullstory.com; font-src github.githubassets.com; form-action 'self' github.com gist.github.com
copilot-workspace.githubnext.com objects-origin.githubusercontent.com; frame-ancestors
'none'; frame-src viewscreen.githubusercontent.com notebooks.githubusercontent.com
www.youtube-nocookie.com; img-src 'self' data: blob: github.githubassets.com
media.githubusercontent.com camo.githubusercontent.com identicons.github.com
avatars.githubusercontent.com private-avatars.githubusercontent.com github-
cloud.s3.amazonaws.com objects.githubusercontent.com release-assets.githubusercontent.com
secured-user-images.githubusercontent.com/ user-images.githubusercontent.com/ private-user-
images.githubusercontent.com opengraph.githubassets.com marketplace-
screenshots.githubusercontent.com/ copilotprodattachments.blob.core.windows.net/github-
production-copilot-attachments/ github-production-user-asset-6210df.s3.amazonaws.com
customer-stories-feed.github.com spotlights-feed.github.com objects-
origin.githubusercontent.com *.githubusercontent.com images.ctfassets.net/8aevphvgewt8/;
manifest-src 'self'; media-src github.com user-images.githubusercontent.com/ secured-user-
images.githubusercontent.com/ private-user-images.githubusercontent.com github-production-
user-asset-6210df.s3.amazonaws.com gist.github.com github.githubassets.com
assets.ctfassets.net/8aevphvgewt8/ videos.ctfassets.net/8aevphvgewt8/; script-src
github.githubassets.com; style-src 'unsafe-inline' github.githubassets.com; upgrade-insecure-
requests; worker-src github.githubassets.com github.com/assets-cdn/worker/
github.com/assets/ gist.github.com/assets-cdn/worker/
server: github.com
accept-ranges: bytes
set-cookie: _gh_sess=TwwwcmStrg21ZdQNQmGKXhYqf%2Fi%2BbWGEP%2Bo%2B%

```

2F4w7vkmUVWTu%2FgNInziOsQzqwuDgJH0aUxrkpa4iD7xgDEvX%2FsAMXmC%2FyXVzuu%
2BG6t1hO6mlXow%2BIN7ou4qNdAEMnrrWSYUUvrW59hzHb2H3%
2FHkbpFmbUZZ6EmZZlr8tHMCg72EII7U3e8TtzSzA0nLxJNHLHV4IXPUc%2Bql4EJT9%
2F11fMtKvOW7e5%2BhBwqYzQH%
2FPWI8sLiFm6TIEQy0TgmtJOGfnTF913Bp3ce9Geby31sm8AA%3D%3D--CR7j3l4eHdkAD8%2B--
UI3I5S3DWRSzU72mlAEJsA%3D%3D; Path=/; HttpOnly; Secure; SameSite=Lax
set-cookie: _octo=GH1.1.613555136.1758717416; Path=/; Domain=github.com; Expires=Thu, 24
Sep 2026 12:36:56 GMT; Secure; SameSite=Lax
set-cookie: logged_in=no; Path=/; Domain=github.com; Expires=Thu, 24 Sep 2026 12:36:56 GMT;
HttpOnly; Secure; SameSite=Lax
x-github-request-id: 9C9D:2CCE73:1BC0738:20A6B22:68D3E5E7

```

```

netshoot:~# ping -c 4 github.com
PING github.com (20.205.243.166) 56(84) bytes of data.
64 bytes from 20.205.243.166: icmp_seq=1 ttl=127 time=105 ms
64 bytes from 20.205.243.166: icmp_seq=2 ttl=127 time=330 ms
64 bytes from 20.205.243.166: icmp_seq=3 ttl=127 time=44.9 ms
^C
--- github.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 44.913/159.731/329.677/122.615 ms

```

I adjusted git path and branch name and git URL.

Still Unknown sync status.

When I ran below:

```
# k logs -n argocd argocd-application-controller-0
{"level":"warning","msg":"Failed to save cluster info: dial tcp 10.43.186.128:6379: i/o
timeout","time":"2025-09-24T12:20:00Z"}
{"level":"warning","msg":"Failed to save cluster info: dial tcp 10.43.186.128:6379: i/o
timeout","time":"2025-09-24T12:20:15Z"}
{"level":"warning","msg":"Failed to save cluster info: dial tcp 10.43.186.128:6379: i/o
timeout","time":"2025-09-24T12:20:25Z"}
```

6379 is redis bro. There shld be an argocd redis pod which application controller uses to persist or fetch state. If not present, it will be an endless Unknown sync status.

Below we see there is a argocd-redis pod and deploy

```
# k get po -n argocd
NAME                               READY   STATUS    RESTARTS   AGE
argocd-application-controller-0      1/1     Running   1 (47m ago)  3d4h
argocd-applicationset-controller-54f96997f8-m2bjb 1/1     Running   1 (47m ago)  3d4h
argocd-dex-server-798cbff4c7-pg447  1/1     Running   1 (47m ago)  3d4h
argocd-notifications-controller-644f66f7df-qw6tc 1/1     Running   1 (47m ago)  3d4h
argocd-redis-6684c6947f-2lkwq       1/1     Running   1 (47m ago)  3d4h
argocd-repo-server-8664486b9d-zknpr 1/1     Running   0          13m
argocd-server-64d5fcbd58-8qbdt     1/1     Running   1 (47m ago)  3d4h

# k get svc -n argocd
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)           AGE
argocd-applicationset-controller   ClusterIP  10.43.125.26  <none>        7000/TCP,8080/TCP
3d4h
```

```

argocd-dex-server           ClusterIP 10.43.151.178 <none>
5556/TCP,5557/TCP,5558/TCP 3d4h
argocd-metrics              ClusterIP 10.43.246.181 <none>     8082/TCP      3d4h
argocd-notifications-controller-metrics ClusterIP 10.43.179.58 <none>     9001/TCP
3d4h
argocd-redis                ClusterIP 10.43.186.128 <none>     6379/TCP      3d4h
argocd-repo-server           ClusterIP 10.43.25.69 <none>     8081/TCP,8084/TCP
3d4h
argocd-server               NodePort 10.43.157.58 <none>
80:31805/TCP,443:32619/TCP 3d4h
argocd-server-metrics        ClusterIP 10.43.22.170 <none>     8083/TCP
3d4h

# k get deploy,sts -n argocd | grep redis
deployment.apps/argocd-redis          1/1   1       1       3d4h

```

I checked logs of app

```

# k describe app -n argocd sampapp

Status:
Conditions:
Last Transition Time: 2025-09-24T12:48:16Z
Message: Failed to load target state: failed to generate manifest for source 1 of 1: rpc
error: code = Unknown desc = failed to list refs: dial tcp 10.43.186.128:6379: i/o timeout
Type: ComparisonError

```

Checked if endpoints of svc matches pod and if redis pod listening on 6379

```

# k exec -n argocd -it argocd-redis-6684c6947f-2lkwq -- netstat -tlnp | grep 6379
Defaulted container "redis" out of: redis, secret-init (init)
tcp    0  0 0.0.0.0:6379      0.0.0.0:*      LISTEN  1/redis-server *:63
tcp    0  0 :::6379         ...*      LISTEN  1/redis-server *:63

$ k -n argocd get endpoints argocd-redis
Warning: v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice
NAME      ENDPOINTS      AGE
argocd-redis  172.18.137.119:6379  3d4h

$ k get po -n argocd -owide | grep argocd-redis -i
argocd-redis-6684c6947f-2lkwq          1/1   Running  1 (52m ago)  3d4h  172.18.137.119
master1 <none>      <none>

```

Tried to apply netpol to allow connection

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-argocd-controller-to-redis
  namespace: argocd
spec:
  podSelector:
    matchLabels:
      app.kubernetes.io/name: argocd-redis

```

```

ingress:
  - from:
    - podSelector:
      matchLabels:
        app.kubernetes.io/name: argocd-application-controller
  ports:
    - protocol: TCP
      port: 6379
  policyTypes:
    - Ingress

```

Reapplied argocd manifest

```
# kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

Restart controller and repo server

```
# kubectl rollout restart deployment -n argocd argocd-repo-server
# kubectl rollout restart statefulset -n argocd argocd-application-controller
```

Status still Unknown for sync status. Health status Healthy

Checked connectivity to redis pod as above logs always complained about redis connectivity

```
# k run netshoot --rm -it --image=nicolaka/netshoot -n argocd -- sh
If you don't see a command prompt, try pressing enter.
~ # nc -zv argocd-redis 6379
```

It hung

Edited the netpol from above to include all allowed in argocd ns

```
$ k get netpol -n argocd allow-argocd-all-to-redis -oyaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"networking.k8s.io/v1","kind":"NetworkPolicy","metadata":{"annotations":{},"name":"allow-argocd-all-to-redis","namespace":"argocd"},"spec":{"ingress":[{"from":[{"namespaceSelector":{"matchLabels":{"kubernetes.io/metadata.name":"argocd"} } }],"ports":[{"port":6379,"protocol":"TCP"}]},"podSelector":{"matchLabels":{"app.kubernetes.io/name":"argocd-redis"}}, "policyTypes":["Ingress"]}}
  creationTimestamp: "2025-09-25T06:21:08Z"
  generation: 1
  name: allow-argocd-all-to-redis
  namespace: argocd
  resourceVersion: "102689"
  uid: a858e622-2006-4051-a651-80095396e8b6
spec:
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
```

```

    kubernetes.io/metadata.name: argocd
  ports:
    - port: 6379
      protocol: TCP
  podSelector:
    matchLabels:
      app.kubernetes.io/name: argocd-redis
  policyTypes:
    - Ingress

```

Tried to test connectivity to redis pod with netshoot pod again but still hung.

Found this netpol:

```

$ k get netpol -n argocd argocd-redis-network-policy -oyaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"networking.k8s.io/v1","kind":"NetworkPolicy","metadata":{"annotations":{},"labels":{"app.kubernetes.io/component":"redis","app.kubernetes.io/name":"argocd-redis","app.kubernetes.io/part-of":"argocd"},"name":"argocd-redis-network-policy","namespace":"argocd"},"spec":{"ingress":[{"from":[{"podSelector":{"matchLabels":{"app.kubernetes.io/name":"argocd-server"}}, {"podSelector":{"matchLabels":{"app.kubernetes.io/name":"argocd-repo-server"}}, {"podSelector":{"matchLabels":{"app.kubernetes.io/name":"argocd-application-controller"}}}], "ports":[{"port":6379,"protocol":"TCP"}]}],"podSelector":{"matchLabels":{"app.kubernetes.io/name":"argocd-redis"}}, "policyTypes":["Ingress"]}}
  creationTimestamp: "2025-09-21T08:19:13Z"
  generation: 1
  labels:
    app.kubernetes.io/component: redis
    app.kubernetes.io/name: argocd-redis
    app.kubernetes.io/part-of: argocd
  name: argocd-redis-network-policy
  namespace: argocd
  resourceVersion: "94067"
  uid: f7560754-bf94-460a-82d4-42a149122597
spec:
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app.kubernetes.io/name: argocd-server
        - podSelector:
            matchLabels:
              app.kubernetes.io/name: argocd-repo-server
        - podSelector:
            matchLabels:
              app.kubernetes.io/name: argocd-application-controller
  ports:
    - port: 6379

```

```
protocol: TCP
podSelector:
  matchLabels:
    app.kubernetes.io/name: argocd-redis
policyTypes:
- Ingress
```

No wonder it hung - it only explicitly allows traffic from argocd-server and repo server and app controller.

But after applying the all to redis netpol, now the argocd-repo-server does not complain to timeout errors to redis IP.

Now I check the application sampapp it is still Unknown, with below error:

```
Status:
Conditions:
Last Transition Time: 2025-09-25T06:09:24Z
Message: Failed to load target state: failed to generate manifest for source 1 of 1: rpc error: code = Unavailable desc = connection error; desc = "transport: Error while dialing: dial tcp 10.43.25.69:8081: i/o timeout"
```

Previously, on ArgocdUI it said 2 errors. Now after fixing the redis thing it is 1 error.

The above log states the argoCD app controller cannot reach the repo server on port 8081

Got some issues with the netpols!!! There are multiple default netpols (when applying the origin argocd manifest)

I deleted ALL netpols in argocd - will reapply some later on.

NOW my cluster sync status is Synced.

```
# k get app sampapp -n argocd
NAME   SYNC STATUS  HEALTH STATUS
sampapp Synced     Healthy
```

ISSUE 4: BUILD FAILED 1

Current deployment.yaml and image tag in DockerHub are same - v1

Changed deployment.yaml's image tag to oldtag instead of v1

Triggered Build and it failed because current deployment.yaml and the image tag in DockerHub are the same - v1

Changed deployment.yaml's image tag to oldtag instead of v1.

Triggered build again.

ISSUE 5: BUILD FAILED 2

Indentation error in Jenkinsfile

Added 4 tabs in Jenkinsfile before image to handle indentation

Build failed again

```
#2 (Sep 25, 2025, 10:01:32 AM) Add description Keep this build forever Started 2 min 11 sec ago Took
25 sec Started by user Jenkins Admin This run spent: 12 sec waiting; 25 sec build duration; 25 sec total
from scheduled to completion. Revision: f8bf15ec15040f22b06a4f90a63c59253e509269 Repository:
https://github.com/mikiangelica/sampapp.git refs/remotes/origin/master Revision:
f8bf15ec15040f22b06a4f90a63c59253e509269 Repository: https://github.com/mikiangelica/k8s-manifests.git refs/remotes/origin/master The following steps that have been detected may have
insecure interpolation of sensitive variables (click here for an explanation): sh: [DOCKER_PASS] No
changes.
```

Why?

The `sed -i 's|image: .*|image: $IMAGE_NAME:$IMAGE_TAG|'` command in `manifests/base/deployment.yaml` assumes your manifest has a line like:

yaml

Copy code

```
image: something
```

But in real-world YAMLs, it's usually indented:

yaml

Copy code

```
image: mikiangelica/sampapp:oldtag
```

Because of the spaces, `sed` didn't match it.

Edited Jenkinsfile to have 4 tabs before image so it can handle indentation:

```
sed -i 's|\s*image: .*|    image: $IMAGE_NAME:$IMAGE_TAG|' manifests/base/deployment.yaml
```

ISSUE 6: BUILD FAILED 3

Docker not found

Made jenkins-agent file for jenkins to use

Failed again. Checked pipeline logs and saw that pipeline failed at this line

```
docker: not found
```

By default, Jenkins agent pod (`jenkins/inbound-agent`) does not include Docker. That's why can't build or push

Make Dockerfile.jenkins-agent file under Dockerfiles/jenkins/

```
FROM jenkins/inbound-agent:latest
```

```
USER root
```

```
# Install Docker CLI and buildx
RUN apt-get update && \
    apt-get install -y curl apt-transport-https ca-certificates gnupg lsb-release && \
    curl -fsSL https://download.docker.com/linux/debian/gpg | apt-key add - && \
    echo "deb [arch=amd64] https://download.docker.com/linux/debian $(lsb_release -cs) stable" \
        > /etc/apt/sources.list.d/docker.list && \
    apt-get update && \
    apt-get install -y docker-ce-cli && \
    mkdir -p ~/.docker/cli-plugins && \
    curl -L https://github.com/docker/buildx/releases/download/v0.16.2/buildx-v0.16.2.linux-amd64 \
        -o ~/.docker/cli-plugins/docker-buildx && \
    chmod +x ~/.docker/cli-plugins/docker-buildx
```

```
USER admin
```

```
# docker login
# docker buildx build -t mikiangelica/jenkins-agent-buildx:latest -f Dockerfile.jenkins-agent .
```

```
# docker push mikiangelica/jenkins-agent-buildx:latest
```

Registering custom agent in Jenkins

Add Globally in Jenkins UI

Go to Jenkins > Manage Jenkins > Clouds > Kubernetes

Under Pod Templates, add new template:

- Name: docker-build-agent
- Namespace: jenkins
- Labels: docker-build
- Usage: Only build jobs with label expressions matching this node
- Pod template to inherit from: (empty)
- Name of the container that will run the Jenkins agent: jnlp
- "Inject Jenkins agent in agent container"
- Containers
 - Name: jnlp
 - Docker image: mikiangelica/jenkins-agent-buildx:latest
 - "Always pull image"
 - Working dir: /home/jenkins/agent
 - Command to run: sleep
 - Args: 9999999
- Service Account: jenkins
- Workspace Volume: Empty Dir Workspace Volume

Click Create button

This will register the pod template globally so that any pipeline can use the label docker-build given.

Edit Jenkinsfile to accept the jenkins agent

Original:

```
pipeline {  
    agent any  
    environment {  
        IMAGE_NAME = "mikiangelica/sampapp"  
        IMAGE_TAG = "v1"  
    }  
}
```

Adjusted:

```
pipeline {  
    agent { label 'docker-build' }  
    environment {  
        IMAGE_NAME = "mikiangelica/sampapp"  
        IMAGE_TAG = "v1"  
    }  
}
```

ISSUE 7: BUILD FAILED 4

Rerun pipeline... failed. I realise I forgot to push Jenkinsfile to GitHub. Pushed to GitHub and went for round 5 of pipeline run.

ISSUE 8: BUILD FAILED 5

This run FAILED again!!

Docker-build-agent pod could not find user

Amended user in Dockerfile.jenkins-agent

This time with different error... The docker-build-agent pod is coming up but with below error:

| | | | | | |
|---------|--------------------------|-----|----------------------|---|------|
| jenkins | docker-build-agent-xrxdp | 0/1 | CreateContainerError | 0 | 2m6s |
|---------|--------------------------|-----|----------------------|---|------|

Events:

| Type | Reason | Age | From | Message |
|---------|-----------|------------------|-------------------|---|
| Normal | Scheduled | 2m17s | default-scheduler | Successfully assigned jenkins/docker-build-agent-xrxdp to master1 |
| Normal | Pulling | 2m17s | kubelet | Pulling image "jenkins/inbound-agent:3327.v868139a_d00e0-6" |
| Normal | Pulled | 87s | kubelet | Successfully pulled image "jenkins/inbound-agent:3327.v868139a_d00e0-6" in 49.376s (49.376s including waiting). Image size: 136833139 bytes. |
| Normal | Created | 87s | kubelet | Created container: set-up-jenkins-agent |
| Normal | Started | 87s | kubelet | Started container set-up-jenkins-agent |
| Normal | Pulled | 41s | kubelet | Successfully pulled image "mikiangelica/jenkins-agent-buildx:latest" in 46.266s (46.266s including waiting). Image size: 239373498 bytes. |
| Warning | Failed | 41s | kubelet | Error: failed to create containerd container: mount callback failed on /var/lib/rancher/k3s/agent/containerd/tmpmounts/containerd-mount126234113: no users found |
| Normal | Pulled | 37s | kubelet | Successfully pulled image "mikiangelica/jenkins-agent-buildx:latest" in 2.549s (2.549s including waiting). Image size: 239373498 bytes. |
| Warning | Failed | 37s | kubelet | Error: failed to create containerd container: mount callback failed on /var/lib/rancher/k3s/agent/containerd/tmpmounts/containerd-mount1069881075: no users found |
| Normal | Pulled | 20s | kubelet | Successfully pulled image "mikiangelica/jenkins-agent-buildx:latest" in 3.061s (3.061s including waiting). Image size: 239373498 bytes. |
| Warning | Failed | 20s | kubelet | Error: failed to create containerd container: mount callback failed on /var/lib/rancher/k3s/agent/containerd/tmpmounts/containerd-mount3507521232: no users found |
| Normal | Pulling | 9s (x4 over 87s) | kubelet | Pulling image "mikiangelica/jenkins-agent-buildx:latest" |
| Normal | Pulled | 6s | kubelet | Successfully pulled image "mikiangelica/jenkins-agent-buildx:latest" in 2.734s (2.734s including waiting). Image size: 239373498 bytes. |
| Warning | Failed | 6s | kubelet | Error: failed to create containerd container: mount callback failed on /var/lib/rancher/k3s/agent/containerd/tmpmounts/containerd-mount1962973324: no users found |

Seeing this error:

| |
|--|
| Error: failed to create containerd container: mount callback failed on /var/lib/rancher/k3s/agent/containerd/tmpmounts/containerd-mount1962973324: no users found |
|--|

Need to amend user in Dockerfile.jenkins-agent

Original:

| |
|-----------------------------------|
| FROM jenkins/inbound-agent:latest |
|-----------------------------------|

| |
|-----------|
| USER root |
|-----------|

| |
|---------------------------------|
| # Install Docker CLI and buildx |
| RUN apt-get update && \ |

```
apt-get install -y curl apt-transport-https ca-certificates gnupg lsb-release && \
curl -fsSL https://download.docker.com/linux/debian/gpg | apt-key add - && \
echo "deb [arch=amd64] https://download.docker.com/linux/debian $(lsb_release -cs) stable" \
> /etc/apt/sources.list.d/docker.list && \
apt-get update && \
apt-get install -y docker-ce-cli && \
mkdir -p ~/.docker/cli-plugins && \
curl -L https://github.com/docker/buildx/releases/download/v0.16.2/buildx-v0.16.2.linux-amd64 \
-o ~/.docker/cli-plugins/docker-buildx && \
chmod +x ~/.docker/cli-plugins/docker-buildx
```

USER admin

Amended to:

```
FROM jenkins/inbound-agent:latest
```

```
USER root
```

```
# Install Docker CLI and buildx
RUN apt-get update && \
    apt-get install -y curl apt-transport-https ca-certificates gnupg lsb-release && \
    curl -fsSL https://download.docker.com/linux/debian/gpg | apt-key add - && \
    echo "deb [arch=amd64] https://download.docker.com/linux/debian $(lsb_release -cs) stable" \
> /etc/apt/sources.list.d/docker.list && \
    apt-get update && \
    apt-get install -y docker-ce-cli && \
    mkdir -p ~/.docker/cli-plugins && \
    curl -L https://github.com/docker/buildx/releases/download/v0.16.2/buildx-v0.16.2.linux-amd64 \
-o ~/.docker/cli-plugins/docker-buildx && \
    chmod +x ~/.docker/cli-plugins/docker-buildx
```

USER jenkins

The repushed Dockerfile to DockerHub:

```
# docker login
# docker buildx build -t mikiangelica/jenkins-agent-buildx:latest -f Dockerfile.jenkins-agent .
# docker push mikiangelica/jenkins-agent-buildx:latest
```

Deleted the docker-build-agent pod and reran pipeline...

ISSUE 9: BUILD FAILED 6

FAILED AGAIN

No Docker daemon running inside the Jenkins agent container

Added hostpath volume to jenkins pod template for jenkins agent (hosts' docker daemon socket file)

```
ERROR: Cannot connect to the Docker daemon at unix:///var/run/docker.sock. Is the docker
daemon running?
```

So the Jenkins agent pod is now running with the custom image but the Jenkins agent container doesn't have a Docker daemon running.

We shall do this dirtily because we want to just test the pipeline. Edit the docker-build-agent pod

template under Jenkins > Clouds > kubernetes > pod template

Volumes ?
List of volumes to mount in agent pod

Host Path Volume

Host path ?
/var/run/docker.sock

Mount path ?
/var/run/docker.sock

Read Only ?

Add Volume ▾

So on the host machine (master node), the Docker daemon (dockerd) is running in the background... That daemon exposes a **Unix socket file** at /var/run/docker.sock.

- This is a doorway into the Docker engine basically
- Any process that can read/write to this socket can issue Docker commands (build, push, run, remove containers etc)

Without the above mount, the Jenkins agent container has the docker CLI installed but no Docker daemon running inside... so when we run docker buildx build ..., the CLI looks for the socket file at /var/run/docker.sock.

Inside the pod itself, the path doesn't exist (or is just an empty file so we get the above error. (cannot connect to Docker daemon at /var/run/docker.sock)

So with the mount, which in YAML looks like:

```
volumes:  
- hostPath:  
  path: /var/run/docker.sock  
  name: docker-sock  
containers:  
- name: jnlp  
  volumeMounts:  
  - mountPath: /var/run/docker.sock  
    name: docker-sock
```

It mounts the hosts' socket file /var/run/docker.sock into the container at the same path,, so that when the docker CLI inside the Jenkins agent tries to talk to /var/run/docker.sock - it is actually talking to the host's Docker daemon.

So the agent container provides the tools (docker, buildx etc)

Host machine does the real building, running, pushing

The container is just a remote control for Docker running outside.

Note: This setup works fine for testing, but in production it's risky (the agent gets root-level access to the host Docker). For prod, **Kaniko/BuildKit** can be used for a safer approach.

Security Implications here:

- Any process in the container can control all Docker containers on the host (root-level access)
- In shared clusters, this is considered insecure
- That's why in production, people use **Kaniko, BuildKit, or img** (which run builds without needing

the host Docker socket).

ISSUE 10: BUILD FAILED 7

Permission denied error

Removed USER jenkins line from Dockerfile.jenkins-agent-build

Reran pipeline after adding the hostpath volume above and it failed again. With permission denied error:

```
ERROR: permission denied while trying to connect to the Docker daemon socket at
unix:///var/run/docker.sock: Head "http://%2Fvar%2Frun%2Fdocker.sock/\_ping": dial unix
/var/run/docker.sock: connect: permission denied
```

So it seems we have mounted it correctly but permission denied.

Usually that socket file is owned by root and group docker on the host.

But in our container, Jenkins run as user jenkins (not root, and not in docker group)

So when docker CLI tries to talk to the socket, the kernel blocks it > permission denied.

Ok so now lets just adjust our Dockerfile so it stays as root. Previously we amended admin to jenkins. Now we just completely remove it

```
FROM jenkins/inbound-agent:latest
```

```
USER root
```

```
# Install Docker CLI and buildx
RUN apt-get update && \
    apt-get install -y curl apt-transport-https ca-certificates gnupg lsb-release && \
    curl -fsSL https://download.docker.com/linux/debian/gpg | apt-key add - && \
    echo "deb [arch=amd64] https://download.docker.com/linux/debian $(lsb_release -cs) stable" \
        > /etc/apt/sources.list.d/docker.list && \
    apt-get update && \
    apt-get install -y docker-ce-cli && \
    mkdir -p ~/.docker/cli-plugins && \
    curl -L https://github.com/docker/buildx/releases/download/v0.16.2/buildx-v0.16.2.linux-amd64 \
        -o ~/.docker/cli-plugins/docker-buildx && \
    chmod +x ~/.docker/cli-plugins/docker-buildx
```

Push to DockerHub again

```
# docker login
# docker buildx build -t mikangelica/jenkins-agent-buildx:latest -f Dockerfile.jenkins-agent .
# docker push mikangelica/jenkins-agent-buildx:latest
```

ISSUE 11: BUILD FAILED 8

Failed to read Dockerfile

Amended path to my sampapp Dockerfile

RUN PIPELINE ATTEMPT #8

Which FAILED again... lol

```
ERROR: failed to solve: failed to read dockerfile: open Dockerfile: no such file or directory
```

So this means we got past the docker not found & permish denied errors. But now it is looking for a file named Dockerfile in the workspace root... So we need to recheck where our sampapp Dockerfile is..

```
# /c/Users/mikia/OneDrive/Documents/SampApp/sampapp/Dockerfiles/sampapp
```

And in my Jenkins file it is

```
stage('Build & Push Docker Image') {
    steps {
        withCredentials([usernamePassword(credentialsId: 'dockerhub-creds',
                                         usernameVariable: 'DOCKER_USER',
                                         passwordVariable: 'DOCKER_PASS')]) {
            sh """
                docker build -t $IMAGE_NAME:$IMAGE_TAG -f Dockerfiles/Dockerfile .
                echo $DOCKER_PASS | docker login -u $DOCKER_USER --password-stdin
                docker push $IMAGE_NAME:$IMAGE_TAG
            """
        }
    }
}
```

Fixed to:

```
stage('Build & Push Docker Image') {
    steps {
        withCredentials([usernamePassword(credentialsId: 'dockerhub-creds',
                                         usernameVariable: 'DOCKER_USER',
                                         passwordVariable: 'DOCKER_PASS')]) {
            sh """
                docker build -t $IMAGE_NAME:$IMAGE_TAG -f Dockerfiles/sampapp/Dockerfile .
                echo $DOCKER_PASS | docker login -u $DOCKER_USER --password-stdin
                docker push $IMAGE_NAME:$IMAGE_TAG
            """
        }
    }
}
```

ISSUE 12: BUILD FAILED 9

No github credentials in Jenkinsfile

Added github credentials in Jenkinsfile (github PAT)

RUN PIPELINE ATTEMPT #9

Failed due to this error:

```
+ sed -i s|image: .*|      image: mikiangelica/sampapp:v1| manifests/base/deployment.yaml
+ git config user.email mikiangelica@gmail.com
+ git config user.name admin
+ git add manifests/base/deployment.yaml
+ git commit -m Update image to mikiangelica/sampapp:v1
[master 8ee8443] Update image to mikiangelica/sampapp:v1
 1 file changed, 1 insertion(+), 1 deletion(-)
+ git push origin master
fatal: could not read Username for 'https://github.com
```

BTW our image built and pushed successfully but the error is about pushing the updated manifest back to Github Repo.

This happens because:

- Jenkins is running inside your custom agent pod.
- That pod doesn't have your GitHub credentials configured globally.
- So when git push runs, it can't authenticate to GitHub.

Need to add withCredentials block to Jenkinfile:

```
stage('Update Manifests') {
    steps {
        withCredentials([usernamePassword(credentialsId: 'github-pat',
                                         usernameVariable: 'GIT_USER',
                                         passwordVariable: 'GIT_PASS')]) {
            sh """
                sed -i 's|image: .*|    image: $IMAGE_NAME:$IMAGE_TAG|'
manifests/base/deployment.yaml
                git config user.email "mikiangelica@gmail.com"
                git config user.name "admin"
                git add manifests/base/deployment.yaml
                git commit -m "Update image to $IMAGE_NAME:$IMAGE_TAG" || echo "No changes to
commit"
                git push https://\$GIT\_USER:\$GIT\_PASS@github.com/mikiangelica/k8s-manifests.git master
"""
        }
    }
}
```

Repushed Jenkinsfile to Github. Had to push twice more (because 10th try I missed out 1 } in the Jenkinsfile which caused it to fail)

Finally 11th time's the charm! It succeeded....

Jenkins / sampapp-pipeline

Status Changes Build Now Configure Delete Pipeline Favorite Open Blue Ocean Rename Pipeline Syntax Credentials

sampapp-pipeline

Edit description

Permalinks

- Last build (#11), 53 sec ago
- Last stable build (#11), 53 sec ago
- Last successful build (#11), 53 sec ago
- Last failed build (#10), 2 min 20 sec ago
- Last unsuccessful build (#10), 2 min 20 sec ago
- Last completed build (#11), 53 sec ago

Builds

Filter

Today

- #11 11:54 AM
- #10 11:53 AM
- #9 11:45 AM
- #8 11:32 AM
- #7 11:17 AM
- #6 11:06 AM
- #5 10:52 AM
- #4 10:45 AM

September 25, 2025

- #3 10:07 AM
- #2 10:01 AM
- #1 9:56 AM